



CS213

Principles of Database Systems(H)

Chapter 6

Shiqi YU 于仕琪

yusq@sustech.edu.cn

Most contents are from Stéphane Faroult's slides



6.1 Ordering

Shiqi Yu 于仕琪

yusq@sustech.edu.cn

Most contents are from Stéphane Faroult's slides

order by

There is one simple expression in SQL to order a result set, which is **ORDER BY**. It comes at the end of a query (although you can have it in subqueries, as you'll see). It is followed by the list of columns used as sort columns.

This will return all movies, starting with the oldest one.

```
select title, year_released  
from movies  
order by year_released
```

Sorts the result of the query

table unchanged

```
select title, year_released  
from movies  
where country = 'us'  
order by year_released
```

We can apply it to any result set ...

```
select m.title,  
       m.year_released  
from movies m  
where m.movieid in  
      (select distinct c.movieid  
       from credits c  
        inner join people p  
        on p.peopleid = c.peopleid  
       where c.credited_as = 'A'  
        and p.birth_year >= 1970)  
order by m.year_released
```

... no matter how complicated the query.

```
select c.country_name,  
       m.title,  
       m.year_released  
from movies m  
       inner join countries c  
       on c.country_code = m.country  
where m.movieid in  
      (select distinct c.movieid  
       from credits c  
        inner join people p  
        on p.peopleid = c.peopleid  
       where c.credited_as = 'A'  
        and p.birth_year >= 1970)  
order by m.year_released
```

and with joins you can
sort by any column of
any table in the join
(remember the super
wide table with all
the columns from all
tables involved)

```
order by col1 desc, col2 asc, ...
```


You can specify that a sort is descending by following the column name with DESC. You can also use ASC to say ascending, but as it's the default nobody uses it.


```
select c.country_name,      1
       m.title,             3
       m.year_released      2
from movies m
     inner join countries c
       on c.country_code = m.country
where m.movieid in
      (select distinct c.movieid
       from credits c
        inner join people p
          on p.peopleid = c.peopleid
       where c.credited_as = 'A'
          and p.birth_year >= 1970)
order by c.country_name,
       m.year_released desc, m.title
```

ordering depends on the data type

Remember that strings are sorted alphabetically,
numbers numerically and dates and times
chronologically. What happens when data is missing?

NULLs ?



Microsoft®
SQL Server®

MySQL®

SQLite

IBM® DB2®

ORACLE®

PostgreSQL

It depends on the DBMS. SQL Server, MySQL and SQLite consider by default that nothing is smaller than everything, and DB2, Oracle and PostgreSQL that it's greater than anything.

Don't believe that things are simple with text, either. They are relatively simple in English, as long as you don't use a foreign word with an accent such as attaché.

In this case, you would probably think that é should sort with e (so do I), but that's not necessarily what internal encoding says. Besides, local habits may vary. Swedes think that ö should come after z. German speakers rather see it with o (Swedish is the default language for MySQL)



ö

a
b

...

y
z



ö

Local text sorting rules are known as "collations". Some products allow you to specify how data in a column should be sorted when you create the table. It's also sometimes possible to specify how you want data to be sorted when you do it.

Collation

```
create table ... (  
    some_text_column varchar(100)  
    collate <collation name> not null,  
    ...)  
  
order by nls_sort(some_text_column,  
    '<collation name>')
```

PostgreSQL



Microsoft
SQL Server



ORACLE®

How are
Chinese text
strings sorted?

How many
collations can
you choose for
Chinese?



I've told you that usually dates are converted to a user-friendly format when returned, for instance with TO_CHAR() available in several products.

```
select to_char(a_date_column, 'MM/DD/YYYY')  
        as event_date, ...  
  
from ...  
where ...  
  
order by event_date
```



But if you sort by this column (text) the sort will be alphabetical!
You should sort by the original, date column:

```
order by a_date_column
```

You can sort by a column that isn't returned.

movies

movied	title	country	year_released
1832	Gone With The Wind	us	1939

For instance suppose that we add producers to the movie database (credited_as = 'P')

credits


movied	peopleid	credited_as
1832	237	A
1832	312	A
1832	742	P
1832	128	D

people


peopleid	first_name	surname	born	died
237	Clark	Gable	1901	1960
742	David	Selznick	1902	1965
312	Vivien	Leigh	1913	1967
128	Victor	Fleming	1889	1949

If we want to sort people by function first, with the director first, producer second and actors last ...

Director



Producer



Actors

Actors


Director

Producer


... no matter whether CREDITED_AS is ascending or descending, sorting by it won't work.

order by credited_as desc

Director



Producer



Actors


```
order by  
  case credited_as  
    when 'D' then 1  
    when 'P' then 2  
    when 'A' then 3  
  end
```

The solution is to use CASE ... END to replace each code with a value that sorts as intended. This is frequently used for "custom sorts".

Three oldest ?

order by year_released

Another problem that isn't so easy is displaying only a limited number of oldest values, or successive "slices" in a long sorted list.



Top 10

title ▼	country	year_released
Annie Hall	us	1977
Blade Runner	us	1982
Bronenosets Potyomkin	ru	1925
Casablanca	us	1942
Citizen Kane	us	1941
Das Boot	de	1985
Det sjunde inseglet	se	1957
Doctor Zhivago	us	1965
Goodfellas	us	1990
Il buono, il brutto, il cattivo	it	1966

1

2

3

4

Successive pages are common on websites. Here titles are sorted.

Skip 10, Top 10

title ▼	country	year_released
Inglourious Basterds	us	2009
Jaws	us	1975
La Belle et la Bête	fr	1946
Ladri di biciclette	it	1948
Lawrence of Arabia	gb	1962
Le cinquième élément	fr	1997
Les Visiteurs du Soir	fr	1942
Mary Poppins	us	1964
On The Waterfront	us	1954
Pather Panchali	in	1955

1

2

3

4

First Page

```
select title,  
        country,  
        year_released  
from movies  
order by title  
limit 10
```

PostgreSQL



MySQL



Several products implement a LIMIT clause that is executed AFTER the sort; this syntax seems to be gaining in popularity.

First Page

```
select title,  
        country,  
        year_released  
from movies  
order by title  
fetch first 10 rows only
```



ORACLE®



DB2 has something slightly different, which was also (more recently) adopted by Oracle and Postgres.

First Page

```
select top 10  
    title,  
    country,  
    year_released  
from movies  
order by title
```



SQL Server is frankly different, but the logic is the same: you sort, then discard everything but what you want.

Third Page

```
select title,  
        country,  
        year_released  
from movies  
order by title  
limit 10 offset 20
```

Retrieving rows 20 to 30 in a sorted result is easy with PostgreSQL, MySQL and SQLite.

PostgreSQL



MySQL®



Third Page

```
select title,  
        country,  
        year_released  
from movies  
order by title  
offset 20  
fetch first 10 rows only
```



ORACLE®

PostgreSQL



What when order is a bit more subtle?

There are many cases when plain ordering isn't satisfying.

10:23



Jennifer

What do you think of 2001 A Space Odyssey?

10:29

1723




Holly

Kubrick's best movie

Reply

10:31

1727



Lorelei

I didn't understand anything

10:35

1732



Darth Vader

Nothing beats Star Wars – [reply to 1723](#)

Such a case is a forum. Somebody posts a topic, then people post their comments in sequence. Things turn ugly when somebody starts posting an answer to a comment rather than to the original topic. Some forums always keep a sequential order and force users to add say @Holly or @1723 (the post id) to help others understand what they are reacting at.

10:23



Jennifer

What do you think of 2001 A Space Odyssey?

10:29

1723



Holly

Kubrick's best movie

Reply

10:35

1732



Darth Vader

Nothing beats Star Wars – [reply to 1723](#)

10:31

1727



Lorelei

I didn't understand anything

10:38

1743



Strangelove

I prefer another one 😊 ... – [reply to 1723](#)

A better solution (for visitors, not developers) is to maintain "threads"

10:23



Jennifer

What do you think of 2001 A Space Odyssey?

10:29

1723



Holly

Kubrick's best movie

10:35

1732



Darth Vader

Nothing beats Star Wars – [reply to 1723](#)

Reply

10:38

1743




Strangelove

I prefer another one 😊 ... – [reply to 1723](#)

10:31

1727



Lorelei

I didn't understand anything

10:36

1733



Harry Lime

Are you kidding? – [reply to 1732](#)

But threads can develop into complicated hierarchies.

10:23



Jennifer

What do you think of 2001 A Space Odyssey?

10:29

1723



Holly

Kubrick's best movie

10:35

1732



Darth Vader

Nothing beats Star Wars – [reply to 1723](#)

Reply

10:36

1733



Harry Lime

Are you kidding? – [reply to 1732](#)

10:38

1743




Strangelove

I prefer another one 😊 ... – [reply to 1723](#)

10:31

1727




Lorelei

I didn't understand anything

10:40

1747



Vito

Darth, you'll stop trolling if I ask you gently.
– [reply to 1732](#)

Nobody's perfect, and the area where SQL database management systems struggle a bit is the management of hierarchies (sometimes referred to as the BOM problem – Bill Of Materials). This is something you encounter everywhere you have to deal with items that can be divided in subitems that can also be subdivided and indefinite number of times. A few example:

- * Cars, made of components that can themselves have subcomponents
- * Chemistry. Ingredients rarely are "pure" ingredients but already the result of chemical processes
- * Financial participations. You can have parts in two companies, one of which also has parts in the other (also known as "financial exposure")

10:23



Jennifer

What do you think of 2001 A Space Odyssey?

10:29

1723



Holly

NULL

order by concat(coalesce(path, ''),
<formatted id>)

10:35

1732



Darth Vader

000001723

10:36

1733




Harry Lime

000001723000001732

10:40

1747



Vito

000001723000001732

10:38

1743




Strangelove

000001723

10:31

1727




Lorelei

NULL

One way to try to solve the problem is the "materialized path", turning the "ancestry" into an attribute.

```
select message, ....  
from forum_posts ...  
connect by answered_postid = prior postid  
start with answered_postid is null  
       and topicid = ...  
order siblings by postid
```

postid	answered_postid
X	
	X



← prior

Oracle has long (since the first half of the 1980s) implemented a way to refer to a 'prior row' in a kind of "dynamic ordering"

Most products handle hierarchies through recursive queries, that we'll see in some detail next time.

6.2 Window Functions

Shiqi Yu 于仕琪

yusq@sustech.edu.cn

Most contents are from Stéphane Faroult's slides

Another very important (but not available in MySQL before 8 or SQLite) set of functions for ordering/reporting are window functions. They bear different names, Oracle calls them analytic functions, DB2 calls them OLAP (OnLine Analytical Processing) functions. They are of two kinds, we'll start with non-ranking functions.

Non-ranking Window Functions

Ranking

We have seen so far two categories of functions: functions that operate on values in the current row (called **scalar functions**), and **aggregate functions**, that operate on sets of rows.

Year of oldest movie per country?

```
select country,  
       min(year_released) earliest_year  
from movies  
group by country
```

The problem with aggregate functions is that details just vanish. If I ask for the year of the oldest movie per country, I get a country, a year, and nothing else.

TITLE and year of the earliest movie per country

If I want some detail, for instance which was the title of this oldest movie, the only option with aggregate functions is to join their output to the very same table that has been aggregated to retrieve the lost detail.

For instance, by joining with movies I can retrieve the title(s) of the movie(s) released in this country that year. Intuitively, we feel that we visited MOVIES twice and that perhaps we could have done better.

```
select a.country,  
       a.title,  
       a.year_released  
from movies a  
  inner join  
    (select country,  
     min(year_released) minyear  
     from movies  
     group by country) b  
on b.country = a.country  
and b.minyear = a.year_released
```


Window functions hold the middle-ground between scalar and aggregate functions. Like scalar functions, they return a result for a single row; but like aggregate functions, this result is computed out of several rows.

The syntax is as follows

func(parameters) over (magic clause)

With DBMS products that support window functions, every aggregate function can be used as a window function. Instead of specifying with GROUP BY the subset on which the result is computed, you say OVER (PARTITION BY ...)

`min(year_released)
over (partition by country)`



```
select country,  
       title,  
       year_released,  
       min(year_released)  
         over (partition by country)  
              earliest_year  
from movies
```

Thus, this query returns two years for every movie: the one when this particular movie was released, and the one when the earliest movie for the same country was released. You get both detail and an aggregate value on the same row.

TITLE and year of the earliest movie per country

country	title	year_released	earliest_year
am	Sayat Nova	1969	1969
ar	La Ciénaga	2001	1945
ar	La bestia debe morir	1952	1945
ar	Truman	2015	1945
ar	Waiting for the Hearse	1985	1945
ar	El hombre de al lado	2010	1945
ar	Derecho de familia	2006	1945
ar	Carancho	2010	1945
ar	Savage Pampas	1966	1945
ar	Cama adentro	2004	1945
ar	Un cuento chino	2011	1945
ar	El hijo de la novia	2001	1945
ar	Delirium	2014	1945
ar	Madame Bovary	1947	1945
ar	La hora de los hornos	1968	1945
ar	El abrazo partido	2004	1945
ar	Hombre mirando al sudeste	1986	1945
ar	Crónica de una fuga	2006	1945
ar	Las aventuras del Capitán Piluso	1963	1945
ar	Albéniz	1947	1945

```
select m.country,  
       m.title,  
       m.year_released  
from (select country,  
            title,  
            year_released,  
            min(year_released)  
              over (partition by country)  
                  earliest_year  
            from movies) m  
where m.year_released = m.earliest_year
```

You just need to limit output to those movies for which the year of release happens to be the same as the earliest one for their country.

country	title	year_released
am	Sayat Nova	1969
ar	Pampa bárbara	1945
at	The Curse	1925
au	The Story of the Kelly Gang	1906
ba	Grbavica	2006
bd	Titāsa Êkaꣳ Nadīra Nāma	1973
be	Miss Mend	1926
bf	Sankofa	1993
bg	Otklonenie	1967
bo	Sangre de cóndor	1969
br	Limite	1931
ca	The Wizard of Oz	1933
ch	Die letzte Chance	1945
cl	El huérfano	1926
cn	Nànfū Nàmqī	1913
co	El inmigrante latino	1980
cu	Soy Cuba	1964
cz	Císařův slavík	1949
de	Nerves	1919
de	Die Austernprinzessin	1919
de	Harakiri	1919
dk	The Picture of Dorian Gray	1910
dz	The Battle of Algiers	1966
ec	Ratas, ratones, rateros	1999
ee	Autumn Ball	2007
eg	Ghazal al-banat	1949
fi	Tuntematon sotilas	1955
fr	L'arrivée d'un train en gare de La Ciotat	1896
gb	Lady Windermere's Fan	1916
ge	Im Shvante	1930

Oldest movie
you like
least?

(country with
several movies)

```
select m.country, m.title,  
       m.year_released  
from  
  (select country,  
         title,  
         year_released,  
         min(year_released)  
          over (partition by country)  
              earliest_year  
        from movies  
       where title <> 'A title here') m  
where m.year_released = m.earliest_year
```

If you filter out, with a WHERE condition, one movie, it will be excluded from the window function computation. The earliest year may become the second earliest.

Window functions always operate against rows that belong to a result set. One related characteristics is that they can only appear after the SELECT, not in the WHERE clause, and there is nothing with them similar to HAVING with aggregate functions (it's not a real limitation; you can always work around it by wrapping the query into another one that applies conditions to its output, as shown previously)

Reporting function

SELECTED rows

```
select a.country, a.title, a.year_released
from movies a
    inner join
    (select country,
        min(year_released) earliest_year
    from movies
    where title <> 'A title here'
    group by country) b
    on b.country = a.country
    and b.earliest_year = a.year_released
```

We have seen the functional equivalence with GROUP BY + join, the previous example works like what is above, with the minimum computed on everything but one movie.


```
select m.country, m.title,  
       m.year_released  
from  
    (select country,  
         title,  
         year_released,  
         min(year_released)  
         over (partition by country)  
              earliest_year  
        from movies) m  
where m.year_released = m.earliest_year  
     and title <> 'A title here'
```

If the query is nested, then the minimum is computed over everything, then filtered out. One country may disappear out of the picture.

```
select a.country, a.title, a.year_released
from movies a
    inner join
    (select country,
        min(year_released) earliest_year
    from movies
    group by country) b
    on b.country = a.country
    and b.earliest_year = a.year_released
where a.title <> 'A title here'
```

This query is functionally equivalent to the previous one.

`min(year_released) over()`

In the same way that you can have an aggregate function without a **GROUP BY** when you want **ONE** result for the whole table, you can have an empty **OVER** clause to indicate that you want the result computed over all rows selected. Note that **OVER ()** is still mandatory otherwise the function would be interpreted as a regular aggregate function, not as a window function.

```
select country_name,  
       cnt as number_of_movies,  
       round(100 * cnt / sum(cnt) over (), 0)  
         as percentage  
from (select c.country_name,  
            coalesce(m.cnt, 0) cnt  
      from countries c  
      left outer join (select country,  
                             count(*) cnt  
                      from movies  
                      group by country) m  
        on m.country = c.country_code  
      ) q  
order by country_name
```

This is frequently used
in operations such as
computing a value as a
percentage of the total.


Side note: when there is an ORDER BY you cannot start returning rows
before you have seen all of them - so you may count them too when sorting,
and the marginal cost of the window function is near zero.

```
select country_name,  
       cnt as number_of_movies,  
       round(100 * cnt / t.movie_count, 0) percentage  
from (select c.country_name,  
            coalesce(m.cnt, 0) cnt  
      from countries c  
      left outer join (select country,  
                           count(*) cnt  
                      from movies  
                      group by country) m  
      on m.country = c.country_code  
     ) q  
  cross join (select count(*) movie_count  
              from movies) t  
 order by country_name
```

The same thing can be obtained with a type of join we haven't seen yet,
a **CROSS JOIN** (without any join condition, also called a Cartesian join)

If all aggregate functions can be used as window functions, there are also some window functions that provide ranking capabilities. These functions are original functions and unrelated to either aggregate functions or scalar functions. There are a few of them, we'll only discuss the most important ones.

ranking window function



aggregate function

func() over (...)


When we talk about "ranking", of course, we implicitly talk about "ordering". In the same way as we can put into the OVER clause how we group, we can also say there how we order.

There are three main ranking functions. In many cases, they return identical values. Differences are interesting.

`row_number()`

`rank()`

`dense_rank()`

With a ranking window function you **MUST** have an **ORDER BY** clause in the **OVER()** (you cannot have an empty **OVER()** clause). You can combine it with a **PARTITION BY** to order with groups.

over (order by ...)

**over (partition by ...
order by ...)**

id	title	country	year_released
1	Casablanca	us	1942
2	Blade Runner	us	1982
3	On The Waterfront	us	1954
4	Lawrence Of Arabia	gb	1962
5	Annie Hall	us	1977
6	Goodfellas	us	1990
7	The Third Man	gb	1949
8	Citizen Kane	us	1941
9	Bicycle Thieves	it	1948
10	The Battleship Potemkin	ru	1925
11	Sholay	in	1975
12	A Better Tomorrow	hk	1986

id	title	country	year_released	
1	Casablanca	us	1942	5
2	Blade Runner	us	1982	2
3	On The Waterfront	us	1954	4
4	Lawrence Of Arabia	gb	1962	1
5	Annie Hall	us	1977	3
6	Goodfellas	us	1990	1
7	The Third Man	gb	1949	2
8	Citizen Kane	us	1941	6
9	Bicycle Thieves	it	1948	1
10	The Battleship Potemkin	ru	1925	1
11	Sholay	in	1975	1
12	A Better Tomorrow	hk	1986	1

```
select title,  
       country,  
       year_released ,  
       row_number()  
         over ( partition by country  
                order by year_released desc ) rn  
  
from movies
```

In this example, movies are grouped by country, and a sequential number is assigned by country to each movie, starting with the most recent movie.

title	country	year_released
Casablanca	us	1942
Blade Runner	us	1982
On The Waterfront	us	1954
Lawrence Of Arabia	gb	1962
Annie Hall	us	1977
Goodfellas	us	1990
The Third Man	gb	1949
Citizen Kane	us	1941
Bicycle Thieves	it	1948
The Battleship Potemkin	ru	1925
Sholay	in	1975
A Better Tomorrow	hk	1986

over (partition by *col1, col2, ...*
order by *col3, col4, ...*)

As with plain GROUP BY and plain ORDER BY, both partitioning and ordering can be applied to several columns (and "columns" can of course be expressions).

What happens to ranks when we have ties?

Tie



What will be the rank of the one who didn't do it to the podium?

`row_number()` assigns distinct,
sequential numbers
to everyone

`row_number()`



rank() assigns the same number to ties,
but there is a gap in ranks.

rank()



`dense_rank()` also assigns the same number toties, with no gap



`dense_rank()`

Which are the **two most recent** movies for each country?

As an aside, a condition on ROW_NUMBER() works a bit like LIMIT applied to ORDER BY, except that ordering can be by group.

Ranking window functions allow answering easily some really tough questions which are almost impossible to answer efficiently otherwise.

```
select x.country,  
       x.title,  
       x.year_released  
from  
(select country,  
         title,  
         year_released,  
         row_number()  
           over (partition by country  
                 order by year_released desc) rn  
  from movies ) x  
where x.rn <= 2
```