



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

09 Randomized Algorithms

CS216 Algorithm Design and Analysis (H)

Instructor: Shan Chen

chens3@sustech.edu.cn



Randomization

- **Algorithm design patterns:**

- Greedy
- Divide and conquer
- Dynamic programming
- Duality (e.g., network flow)
- Reductions

- **Randomization**

in practice, access to a pseudorandom number generator

- **Randomization.** Allow **fair** coin flip in unit time.

- **Why randomize?** Can lead to **simplest, fastest, or only known algorithm** for a particular problem.

- E.g., symmetry-breaking protocols, graph algorithms, quicksort, hashing, load balancing, closest pair, Monte Carlo integration, cryptography, etc.




1. Content Resolution



Contention Resolution in Distributed System

- **Contention resolution.** Given n processes P_1, \dots, P_n , each competing for access to a shared database. If **two or more** processes access the database **simultaneously**, all processes are locked out. Devise protocol to ensure all processes get through on a regular basis.
- **Restriction.** Processes can't communicate.
- **Challenge.** Need **symmetry-breaking** paradigm.





Contention Resolution: Randomized Protocol

- **Randomized protocol.** Each process requests access to the database at any round t with probability $p = 1/n$.
- **Lemma 1.** Let $S[i, t]$ = event that process i succeeds in accessing the database at round t . Then $1/(2n) \geq \Pr[S(i, t)] \geq 1/(e \cdot n)$.
- **Useful facts from calculus.** As n increases from 2, the function:
 - $(1 - 1/n)^n$ converges monotonically from $1/4$ up to $1/e$.
 - $(1 - 1/n)^{n-1}$ converges monotonically from $1/2$ down to $1/e$.
- **Pf.** By independence, $\Pr[S(i, t)] = p(1 - p)^{n-1}$.
← process i requests access, none of remaining processes requests access

Setting $p = 1/n$, we have $\Pr[S(i, t)] = 1/n \underbrace{(1 - 1/n)^{n-1}}_{\text{between } 1/e \text{ and } 1/2}$. ■

↑
value that maximizes $\Pr[S(i, t)]$

between $1/e$ and $1/2$



Contention Resolution: Randomized Protocol

- **Randomized protocol.** Each process requests access to the database at any round t with probability $p = 1/n$.
- **Lemma 2.** The probability that process i fails to access the database in $e \cdot n$ rounds is at most $1/e$. After $e \cdot n (c \ln n)$ rounds, the probability $\leq n^{-c}$.
- **Pf.** Let $F[i, t]$ = event that process i fails to access database between rounds $1 \sim t$. By independence and **Lemma 1**, $\Pr[F(i, t)] \leq (1 - 1/(en))^t$.

$$\text{Choose } t = \lceil e \cdot n \rceil: \Pr[F(i, t)] \leq \left(1 - \frac{1}{en}\right)^{\lceil en \rceil} \leq \left(1 - \frac{1}{en}\right)^{en} \leq \frac{1}{e}$$

$$\text{Choose } t = \lceil e \cdot n \rceil \lceil c \cdot \ln n \rceil: \Pr[F(i, t)] \leq \left(\frac{1}{e}\right)^{c \ln n} = n^{-c}$$



Contention Resolution: Randomized Protocol

- **Theorem.** The probability that **all** processes succeed within $2en \ln n$ rounds is $\geq 1 - 1/n$.
- **Pf.** Let $F[t]$ = event that **at least one** of the n processes fails to access database in any rounds $1 \sim t$.

$$\Pr[F[t]] = \Pr\left[\bigcup_{i=1}^n F[i, t]\right] \leq \sum_{i=1}^n \Pr[F[i, t]]$$

union bound Lemma 2

Choosing $t = \lceil e \cdot n \rceil \lceil 2 \ln n \rceil$ yields $\Pr[F[t]] \leq n \cdot n^{-2} = 1/n$. ■

Union bound. Given events E_1, \dots, E_n ,
$$\Pr\left[\bigcup_{i=1}^n E_i\right] \leq \sum_{i=1}^n \Pr[E_i]$$



2. Global Min Cut



Global Minimum Cut

- **Global min cut.** Given a connected, undirected graph $G = (V, E)$, find a cut (A, B) of minimum cardinality.
- **Applications.** Partitioning items in a database, identify clusters of related documents, network reliability, circuit design, TSP solvers, etc.
- **Network flow solution:**
 - Replace every edge (u, v) with two antiparallel edges (u, v) and (v, u) .
 - Pick any vertex $s \in V$: for every other node $v \in V$, compute min s - v cut.
- **False intuition.** Global min-cut is harder than min s - t cut.



Global Min Cut: Contraction Algorithm

- **Contraction algorithm:** [Karger 1995]

- Pick an edge $e = (u, v)$ uniformly at random.
- **Contract** edge e .
 - ✓ replace u and v by single new supernode w
 - ✓ preserve edges, updating endpoints of u and v to w
 - ✓ keep parallel edges, but delete self-loops






Global Min Cut: Contraction Algorithm

- **Contraction algorithm:** [Karger 1995]

- Pick an edge $e = (u, v)$ uniformly at random. Contract edge e .
- Repeat until graph has just **two** supernodes v_1 and v_2 .
- Return the cut $(S(v_1), S(v_2))$ (where $S(v_i)$ denote all nodes contracted to v_i).




Reference: Thore Husfeldt



Contraction Algorithm: Analysis

- **Theorem.** The contraction algorithm returns a **min cut** with prob $\geq 2/n^2$.
- **Pf.** Consider a global min cut (A^*, B^*) of G . Let F^* be edges in this min cut and let $k = |F^*|$ = size of **min cut**.
 - In **first step**, algorithm contracts an edge in F^* with probability $k / |E|$.
 - Every node has degree $\geq k$ since otherwise (A^*, B^*) would not be a min-cut. Therefore, we have $2|E| \geq kn \Leftrightarrow k / |E| \leq 2/n$.
 - Thus, the algorithm contracts an edge in F^* with probability $\leq 2/n$.





Contraction Algorithm: Analysis

- **Theorem.** The contraction algorithm returns a **min cut** with prob $\geq 2/n^2$.
- **Pf.** Consider a global min cut (A^*, B^*) of G . Let F^* be edges in this min cut and let $k = |F^*|$ = size of **min cut**.
 - Let $G' = (V', E')$ be graph after j iterations, then G' has $n' = n - j$ (super)nodes.
 - If no edge in F^* has been contracted, the min-cut in G' is still k . Then, as before, $k / |E'| \leq 2/n'$. Thus, algorithm contracts an edge in F^* with probability $\leq 2/n'$.
 - Let E_j = event that **no edge** in F^* is contracted in iteration j .

$$\begin{aligned}\Pr[E_1 \cap E_2 \cdots \cap E_{n-2}] &= \Pr[E_1] \times \Pr[E_2 \mid E_1] \times \cdots \times \Pr[E_{n-2} \mid E_1 \cap E_2 \cdots \cap E_{n-3}] \\ &\geq \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \cdots \left(1 - \frac{2}{4}\right) \left(1 - \frac{2}{3}\right) \\ &= \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \cdots \left(\frac{2}{4}\right) \left(\frac{1}{3}\right) \\ &= \frac{2}{n(n-1)} \\ &\geq \frac{2}{n^2}\end{aligned}$$



Contraction Algorithm: Amplification

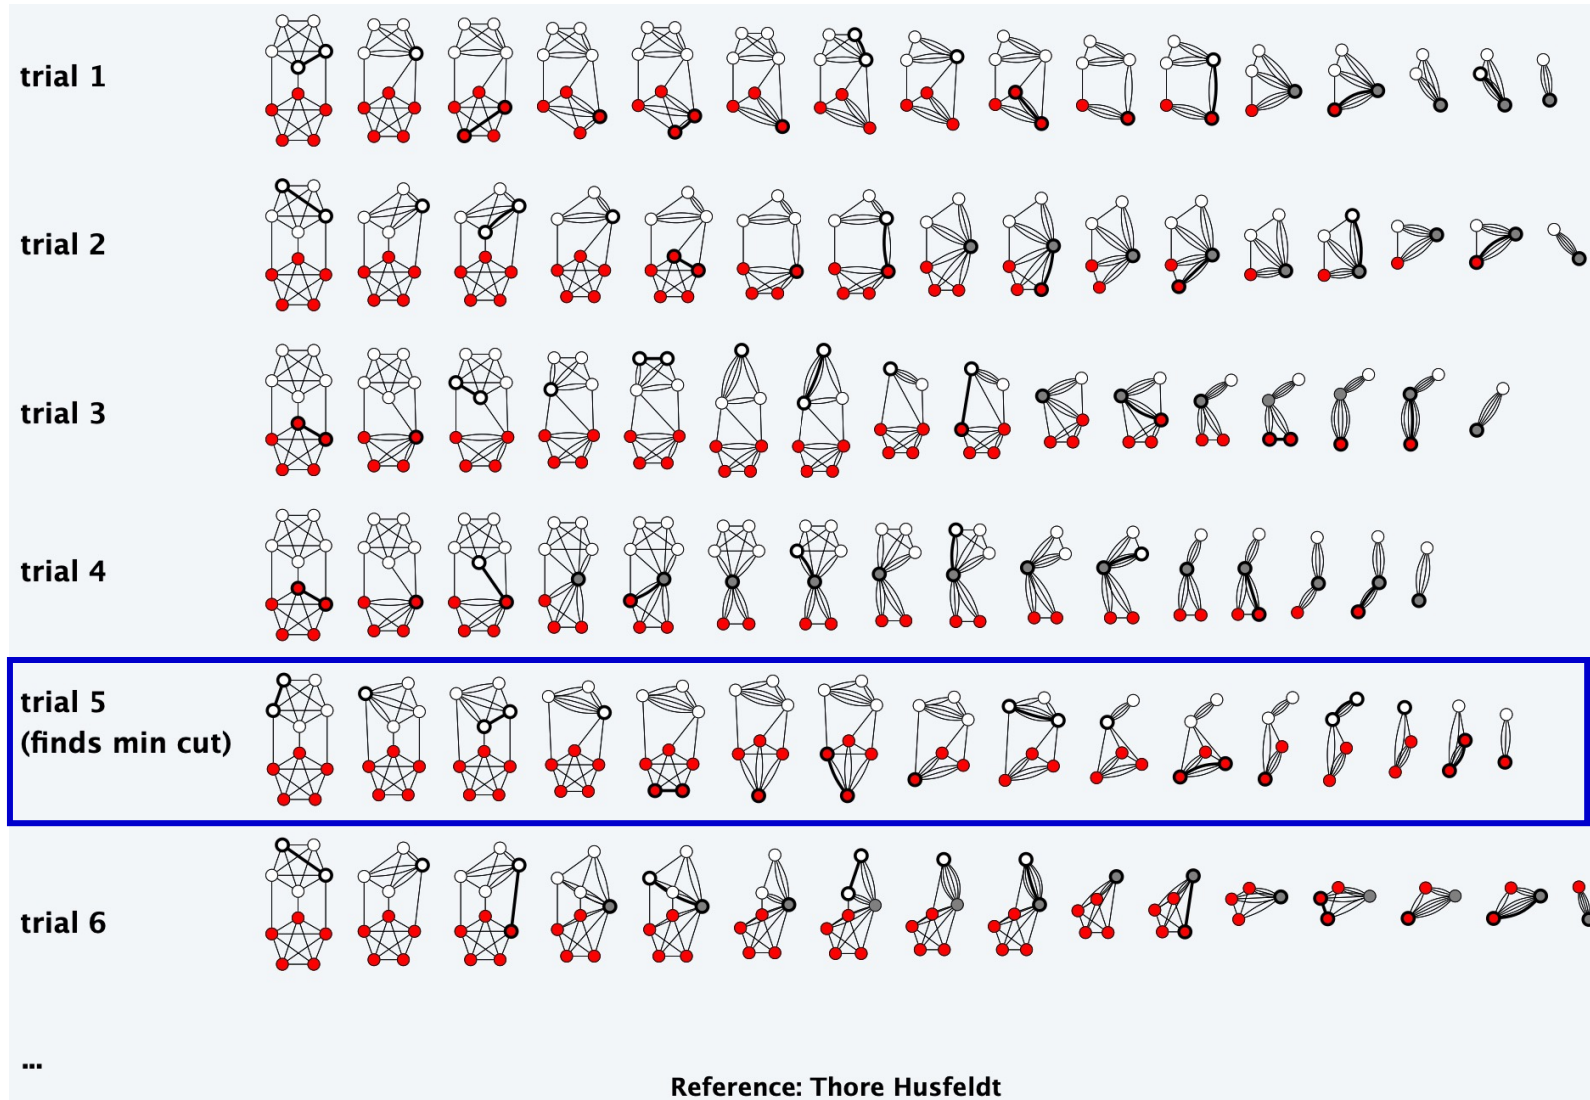
- **Amplification.** To **amplify** the probability of success, run the contraction algorithm **many times** with **independent randomness**.
- **Claim.** If we repeat the contraction algorithm $n^2 \ln n$ times, then the probability of failing to find the global min cut is $\leq 1/n^2$.
- **Pf.** By independence, the probability of failure is at most

$$\left(1 - \frac{2}{n^2}\right)^{n^2 \ln n} = \left[\left(1 - \frac{2}{n^2}\right)^{\frac{1}{2}n^2}\right]^{2 \ln n} \leq \left(e^{-1}\right)^{2 \ln n} = \frac{1}{n^2}$$

\uparrow
 $(1 - 1/x)^x \leq 1/e$



Contraction Algorithm: Demo





More on Global Minimum Cut

- **Remark.** Overall running time $\Theta(n^2 m \log n)$ is slow since we perform $\Theta(n^2 \log n)$ iterations and each takes $\Omega(m)$ time.
- **Improvement:** [Karger-Stein 1996] $O(n^2 \log^3 n)$
 - Early iterations are less risky than later ones: probability of contracting an edge in min cut hits **50%** when $n/\sqrt{2}$ nodes remain.
 - Run contraction algorithm until $n/\sqrt{2}$ nodes remain.
 - Run contraction algorithm **twice** on resulting graph and return **best** of two cuts.
- **Extensions.** Naturally generalizes to handle positive weights.
- **Best known.** [Karger 2000] $O(m \log^3 n)$. ← faster than best known max flow algorithm or deterministic global min cut algorithm



3. Load Balancing



Load Balancing

- **Load balancing.** System in which m jobs arrive in a stream and need to be processed immediately on n identical processors. Find an assignment that balances the workload across processors.
- **Centralized controller.** Assign jobs in round-robin manner. Each processor receives at most $\lceil m/n \rceil$ jobs.
- **Decentralized controller.** Assign jobs to processors **uniformly at random**. How likely is it that some processor is assigned **“too many”** jobs?



Chernoff Bounds

- **Setting:**

- X_1, \dots, X_n : independent random variables on $\{0, 1\}$
- $X = X_1 + \dots + X_n$
- $E(X) = E(X_1) + \dots + E(X_n)$

- **Theorem. (above mean)** For any $\delta > 0$ and $\mu \geq E(X)$, we have

$$\Pr[X > (1 + \delta)\mu] < \left(\frac{e^\delta}{(1 + \delta)^{1+\delta}} \right)^\mu$$

typically choose $\mu = E(X)$

- **Theorem. (below mean)** For any $\delta > 0$ and $\mu \leq E(X)$, we have

$$\Pr[X < (1 - \delta)\mu] < e^{-\delta^2 \mu / 2}$$

- **Takeaway.** Provide exponentially decreasing bounds on the probabilities of large deviations from the expected value.



Load Balancing: # Jobs = # Processors

- **Analysis:** (number of jobs m = number of processors n)

- Let X_i = number of jobs assigned to processor i .
- Let $Y_{ij} = 1$ if job j is assigned to processor i , and $Y_{ij} = 0$ otherwise.
- Thus, $X_i = \sum_j Y_{ij}$. We have $\mathbf{E}[Y_{ij}] = 1/n$ and $\mathbf{E}[X_i] = 1$.
- Chernoff bounds with $\mu = \mathbf{E}[X_i] = 1$ and $\delta = c - 1 > 0 \Rightarrow \Pr[X_i > c] < e^{c-1} / c^c$.
- Let $\gamma(n)$ be number x such that $x^x = n$, and choose $c = e \gamma(n)$.

$$\Pr[X_i > c] < \frac{e^{c-1}}{c^c} < \left(\frac{e}{c}\right)^c = \left(\frac{1}{\gamma(n)}\right)^{e\gamma(n)} \leq \left(\frac{1}{\gamma(n)}\right)^{2\gamma(n)} = \frac{1}{n^2}$$

- **Union bound** \Rightarrow with probability $\leq n \cdot 1/n^2 = 1/n$ exists some processor receives more than c jobs \Rightarrow with probability $\geq 1 - 1/n$ no processor receives more than $c = e \gamma(n) = \Theta(\log n / \log \log n)$ jobs.
 - ✓ Do \ln and $\ln \ln$ on both sides of $\gamma(n)^{\gamma(n)} = n \Rightarrow \gamma(n)/2 \leq \log n / \log \log n \leq \gamma(n)$.



Load Balancing: # Jobs > # Processors

- **Theorem.** Suppose the number of jobs $m = 16 n \ln n$. Then **on average**, each of the n processors handles $16 \ln n$ jobs. With high probability, every processor will have **between half and twice** the average load.

- **Pf.** (number of jobs $m >$ number of processors n)

- Let X_i = number of jobs assigned to processor i .
- Applying **Chernoff bounds** with $\delta = 1$ and $\mu = \mathbf{E}(X_i) = 16 \ln n$ yields

$$\Pr[X_i > 2\mu] < \left(\frac{e}{4}\right)^{16 \ln n} < \left(\frac{1}{e}\right)^{2 \ln n} = \frac{1}{n^2}$$

$$\Pr\left[X_i < \frac{1}{2}\mu\right] < e^{-\frac{1}{2}\left(\frac{1}{2}\right)^2 16 \ln n} = \frac{1}{n^2}$$

- **Union bound** \Rightarrow every processor has load between half and twice the average with probability $\geq 1 - 2/n$. ▀



4. MAX 3-SAT



Maximum 3-Satisfiability

- **MAX 3-SAT.** Given a **3-SAT** formula, find a truth assignment that **satisfies as many clauses as possible**.

$$C_1 = x_2 \vee \overline{x_3} \vee \overline{x_4}$$

$$C_2 = x_2 \vee x_3 \vee \overline{x_4}$$

$$C_3 = \overline{x_1} \vee x_2 \vee x_4$$

$$C_4 = \overline{x_1} \vee \overline{x_2} \vee x_3$$

$$C_5 = x_1 \vee \overline{x_2} \vee \overline{x_4}$$

- **Remark.** **NP-hard** optimization problem.
- **Simple idea.** Flip a coin, and set each variable true with probability $\frac{1}{2}$, **independently** for each variable.



Maximum 3-Satisfiability: Analysis

- **Theorem.** Given a **3-SAT** formula with k clauses, the **expected number** of clauses satisfied by a random assignment is $7k/8$.
- **Pf.** Consider random variables $Z_j = \begin{cases} 1 & \text{if clause } C_j \text{ is satisfied} \\ 0 & \text{otherwise.} \end{cases}$

Let $Z = \sum_j Z_j$ be **number of clauses satisfied** by random assignment.


$$\begin{aligned} E[Z] &= \sum_{j=1}^k E[Z_j] \\ \text{linearity of expectation} &\nearrow \\ &= \sum_{j=1}^k \Pr[\text{clause } C_j \text{ is satisfied}] \\ &= \frac{7}{8}k \end{aligned}$$

\nwarrow **disjunction of 3 literals**
each literal corresponds to a different variable



The Probabilistic Method

- **Corollary.** For any instance of **3-SAT**, there **exists** a truth assignment that satisfies **at least** a **$7/8$** fraction of all clauses.
- **Pf.** Random variable is **at least its expectation** some of the time. ■
- **Probabilistic method.** [Paul Erdős] Prove the **existence** of a non-obvious property by showing that a **random** construction produces it with **positive probability**!





Maximum 3-Satisfiability: Further Analysis

- **Q.** Can we turn this idea into a **7/8-approximation** algorithm?
- **A.** Yes (but a random variable can almost always be below its mean).
- **Lemma.** The probability that a random assignment satisfies $\geq 7k/8$ clauses is at least $1/(8k)$.
- **Pf.** Let p_j be probability that exactly j clauses are satisfied; let p be probability that $\geq 7k/8$ clauses are satisfied.

$$\begin{aligned}\frac{7}{8}k = E[Z] &= \sum_{j \geq 0} j p_j = \sum_{j < 7k/8} j p_j + \sum_{j \geq 7k/8} j p_j \\ &\leq \underbrace{\left(\frac{7k}{8} - \frac{1}{8}\right)}_{\substack{\text{j is integer} \\ \nearrow}} \sum_{j < 7k/8} p_j + k \sum_{j \geq 7k/8} p_j \leq \left(\frac{7}{8}k - \frac{1}{8}\right) \cdot 1 + k p\end{aligned}$$

Rearranging terms yields $p \geq 1/(8k)$. ■



Maximum 3-Satisfiability: Analysis

- **Johnson's algorithm.** Repeatedly generate random truth assignments until one of them satisfies $\geq 7k/8$ clauses.
- **Theorem.** Johnson's algorithm is a $7k/8$ -approximation algorithm.
- **Pf. (direct proof)**
 - **Lemma** \Rightarrow each iteration succeeds with probability $p \geq 1/(8k)$
 - The expected number of trials to find the satisfying assignment is
$$\sum_{j=1}^{\infty} j \Pr[j \text{ trials}] = \sum_{j=1}^{\infty} j(1-p)^{j-1}p = \frac{1}{(1-(1-p))^2}p = \frac{1}{p} \leq 8k \quad \blacksquare$$

calculus fact \nearrow
- **Takeaway.** NP-hard problems may have good approximation algorithms.



Maximum Satisfiability

- **Extensions:**

- **MAX-SAT:** Allow one, two, or more literals per clause.
- **Weighted MAX-SAT:** Find max **weighted** set of satisfied clauses.

- **Theorem.** [Asano-Williamson 2000] There exists a **0.784-approximation** algorithm for **MAX-SAT**.
- **Theorem.** [Karloff-Zwick 1997, Zwick+computer 2002] There exists a **deterministic $7/8$ -approximation** algorithm for version of **MAX 3-SAT** in which each clause has ≤ 3 literals.
- **Theorem.** [Håstad 1997] Unless **P = NP**, no **ρ -approximation** algorithm for **MAX 3-SAT** (and hence **MAX SAT**) for any **$\rho > 7/8$** .

↑
very unlikely to improve over
simple randomized algorithm for **MAX 3-SAT**



Randomized Algorithms: Closing Remarks

- **Monte Carlo.** Guaranteed to **run poly-time**, likely to find correct answer.
- **Example.** Contraction algorithm for global min cut.
- **Las Vegas.** Guaranteed to **find correct answer**, likely to run in poly-time.
- **Example.** Randomized quicksort, Johnson's MAX 3-SAT algorithm.

stop algorithm after a certain point

- **Remark.** Can always convert a Las Vegas algorithm into Monte Carlo, but **no known method (in general)** to convert the other way.

e.g., don't know when to stop



Randomized Algorithms: Closing Remarks

- **RP.** (Randomized Poly-Time) [Monte Carlo] Decision problems solvable with **one-sided error** in poly-time.
- **One-sided error:**
 - If the correct answer is **no**, **always** return **no**.
 - If the correct answer is **yes**, return **yes** with probability $\geq 1/2$.

can decrease probability of false negative to 2^{-100} by 100 independent repetitions
- **ZPP.** [Las Vegas] Decision problems solvable in **expected** poly-time.

running time can be unbounded, but fast on average
- **Theorem.** $P \subseteq ZPP \subseteq RP \subseteq NP$
- **Fundamental open questions.** To what extent does **randomization** help?
 - Does $P = ZPP$? Does $ZPP = RP$? Does $RP = NP$?