

CS217 - Data Structures & Algorithm Analysis (DSAA)

Lecture #8

► Elementary Data Structures

Prof. Pietro S. Oliveto

Department of Computer Science and Engineering

Southern University of Science and Technology (SUSTech)

`olivetop@sustech.edu.cn`

<https://faculty.sustech.edu.cn/olivetop>

Reading: Part III Introduction & Chapter 10

► Aims of this lecture

- To introduce **data structures** and their typical operations.
- **Stacks, queues, priority queues** and **linked lists**.
- To work out the **running time** for operations on these data structures.
- To identify pros and cons for data structures in terms of efficiency.

► Data Structures

- **Dynamic sets** that can store and retrieve elements.
- Data structures are techniques for representing finite dynamic sets of elements
- Each element can contain:
 - a **key**, used to identify the element
 - **Satellite data**, carried around but unused by the data structure
 - **Attributes**, that are manipulated by the data structure eg., pointers to other objects
- Often keys stem from a **totally ordered set** (e. g. numbers)
 - Allows to define the minimum, successor and predecessor

► Data Structure Operations

- Operations on a dynamic sets S can be grouped into **queries** and **modifying operations**:
- Typical operations:
 - **Search(S, k)**: returns element x with **key k** , or NIL
 - **Insert(S, x)**: adds **element x** to S
 - **Delete(S, x)**: removes element x from S
 - **Minimum(S), Maximum(S)**: return x resp. with smallest or largest key
 - **Successor(S, x), Predecessor(S, x)**: next larger (smaller) than $\text{Key}(x)$
- **Time** often measured using n as the number of elements in S .

► Data Structure Operations

- What's the runtime of each operation on an **array**?
- **Search(S, k)**: returns element x with **key k**, or NIL $\Theta(n)$
- **Insert(S, x)**: adds **element x** to S $\Theta(1)$
- **Delete(S, x)**: removes element x from S $\Theta(1)$
- **Minimum(S), Maximum(S)**: return x resp. with smallest or largest key $\Theta(n)$
- **Successor(S, x), Predecessor(S, x)**: next larger (smaller) than Key(x) $\Theta(n)$

► Data Structure Operations

- What's the runtime of each operation on a **sorted array**?
- **Search(S, k)**: returns element x with **key** k , or NIL $\Theta(\log n)$
- **Insert(S, x)**: adds **element** x to S $\Theta(n)$
- **Delete(S, x)**: removes element x from S $\Theta(n)$
- **Minimum(S), Maximum(S)**: return x resp. with smallest or largest key $\Theta(1)$
- **Successor(S, x), Predecessor(S, x)**: next larger (smaller) than $\text{Key}(x)$ $\Theta(1)$


We'll now see some data structures that improve on the array implementation for many of the dynamic-set operations.

► Roadmap for the next lectures

- Simple data structures
 - Stacks
 - Queues
 - Linked lists
 - Binary search trees
 - Graphs
- Advanced data structures
 - Balanced trees
 - Priority queues

► Stacks

3
6
8



- Only the **top element** is accessible in a stack.
 - Last-in, first-out policy (LIFO)
- Insert is usually called **Push**, and Delete is called **Pop**.



► Stacks implemented using arrays

- Stacks can be implemented as an array S with attribute $S.top$.

PUSH(S, x)


```
1  if  $S.top == S.size$ 
2      error "overflow"
3  else  $S.top = S.top + 1$ 
4       $S[S.top] = x$ 
```

STACK-EMPTY(S)


```
1  if  $S.top == 0$ 
2      return TRUE
3  else return FALSE
```

POP(S)


```
1  if STACK-EMPTY( $S$ )
2      error "underflow"
3  else  $S.top = S.top - 1$ 
4      return  $S[S.top + 1]$ 
```



(a)



(b)



(c)

- All stack operations take time $O(1)$.

► Stacks Application (1): Bracket Balance Checking

- $1 + \{2 * [x + (4y - z)] * [5x - (5y + z)] - 5t\}$
- $\{[()] [()]\}$
- Are the brackets correctly balanced or not?
- Read the expression: **Push** each opening bracket and **pop** for each closing bracket
- If the type of popped bracket always matches return **true**, else return **false**
- What's the runtime of the algorithm?

► Stacks Application (2): Postfix expression

- $5 * ((9 + 3) * (4 * 2) + 7)$ (infix expression)
- $5\ 9\ 3\ +\ 4\ 2\ *\ * \ 7\ +\ *$ (postfix expression)
- Parsing postfix expressions is somewhat easier than infix expressions. Why?
- Read the tokens one at a time:
 - If it is an operand, **push** it on the stack
 - If it is a binary operator **pop** twice, apply the operator, and **push** the result back on the stack
- What is the runtime of the algorithm?

► Stacks Application (2): Postfix expression

- $5 * ((9 + 3) * (4 * 2) + 7)$ (infix expression)
- $5\ 9\ 3\ +\ 4\ 2\ *\ * \ 7\ +\ *$ (postfix expression)

Stack operations

- ◆ push(5)
- ◆ push(9)
- ◆ push(3)
- ◆ push(pop() + pop())
- ◆ push(4)
- ◆ push(2)
- ◆ push(pop() * pop())
- ◆ push(pop() * pop())
- ◆ push(7)
- ◆ push(pop() + pop())
- ◆ push(pop() * pop())

Stack elements

5
5 9
5 9 3
5 12
5 12 4
5 12 4 2
5 12 8
5 96
5 96 7
5 103
515

► Queues



head

3	6	8
---	---	---

 tail

- The British love them 😊
- The first element in a queue is accessible.
 - First-in, first-out policy (FIFO)
- Insert is called **Enqueue**, Delete is called **Dequeue**.
- Queues have a **head** and a **tail**, like in a supermarket
 - Elements are added to the tail
 - Elements are extracted from the head

► Queues implemented using arrays




- Queues can be stored in an array “**wrapped around**”.

ENQUEUE(Q, x)

```
1   $Q[Q.tail] = x$ 
2  if  $Q.tail == Q.size$ 
3       $Q.tail = 1$ 
4  else  $Q.tail = Q.tail + 1$ 
```

DEQUEUE(Q)

```
1   $x = Q[Q.head]$ 
2  if  $Q.head == Q.size$ 
3       $Q.head = 1$ 
4  else  $Q.head = Q.head + 1$ 
5  return  $x$ 
```



- All queue operations take time $O(1)$.

► Queues: Applications

- Playlists (eg., iTunes)
- Dispensing requests on a shared resource (eg., a printer, a server, a processor etc.,)
- Data buffers (eg., streaming services)
- What if I have priorities on the use of the resource?

➤ Priority Queues: Motivation

- Schedule jobs on a computer shared among multiple users
- A max-priority queue keeps track of the jobs to be performed and their relative priorities
- When a job is finished the scheduler selects the job with highest priority from those pending
- Jobs can be added to the scheduler at any time


Job	Owner	Priority (key)
Job 1	Yao Xin	35
Job 12	Oliveto Pietro	2
Job 24	Hao Qi	22
Job 25	Yu Shiqi	18
Job 72	Yao Xin	30

- **Use a heap!**


➤ Heap Properties

- **Max-heap property:** for every node other than the root, the parent is no smaller than the node, $A[\textit{Parent}(i)] \geq A[i]$.
- In a max-heap, the **root** always stores a **largest** element.

→ this is what we want!



(a)



(b)

- **Min-heap property:** for every node other than the root, the parent is no larger than the node, $A[\textit{Parent}(i)] \leq A[i]$.

➤ Priority Queue based on max-heap

- A data structure for maintaining a set S of elements with an associated element called key (the priority).

Operation	Time
Insert(S, x, k) – inserts x with key k into S	
Maximum (S) – returns the element in S with the largest key	
Extract-Max(S) – removes and returns element in S with the largest key	
Increase-Key(S, x, k) – increases the key of x to a larger value k (element may float up in the heap)	

➤ Priority Queue based on max-heap

- A data structure for maintaining a set S of elements with an associated element called key (the priority).

Operation	Time
Insert(S, x, k) – inserts x with key k into S	$O(\log n)$
Maximum (S) – returns the element in S with the largest key	$O(1)$
Extract-Max(S) – removes and returns element in S with the largest key	$O(\log n)$
Increase-Key(S, x, k) – increases the key of x to a larger value k (element may float up in the heap)	$O(\log n)$

Job x : x .satellite_data; x .job_address x .priority (key)

(We need a way to map the position of job x in the heap (and update it as it moves in the heap) as well as the pointer to the job to execute it)

Min-priority queue based on min-heap also exist: we will use them in graph algorithms (eg., Dijkstra, Prim)

➤ Find and extract next job

MAX-HEAP-MAXIMUM(A)

```
1  if  $A.heap-size < 1$ 
2      error “heap underflow”
3  return  $A[1]$ 
```

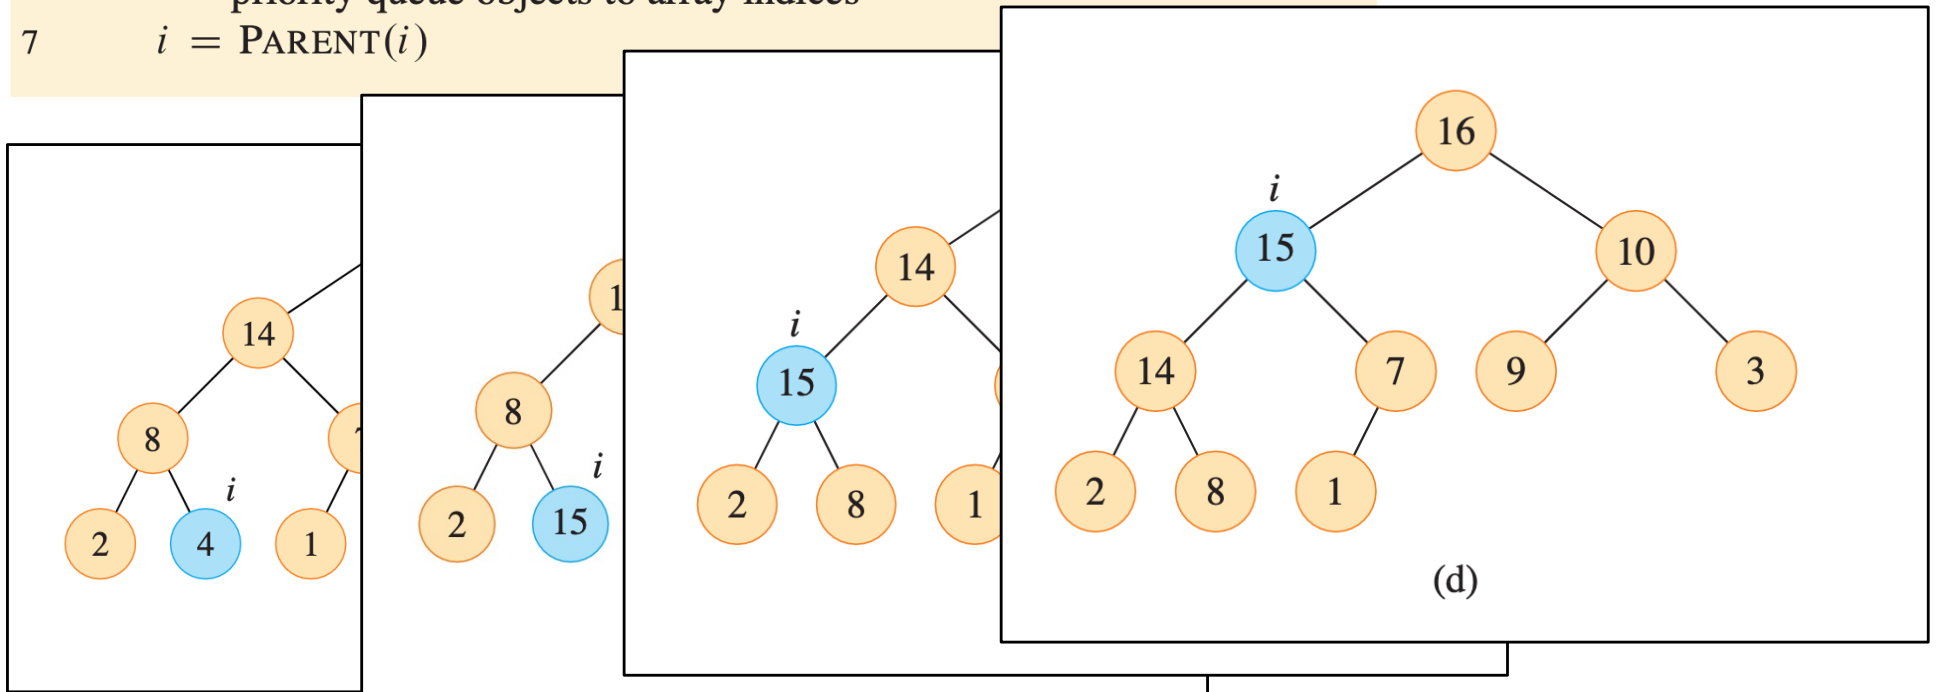
MAX-HEAP-EXTRACT-MAX(A)

```
1   $max = \text{MAX-HEAP-MAXIMUM}(A)$ 
2   $A[1] = A[A.heap-size]$ 
3   $A.heap-size = A.heap-size - 1$ 
4   $\text{MAX-HEAPIFY}(A, 1)$ 
5  return  $max$ 
```

➤ Increase job priority

MAX-HEAP-INCREASE-KEY(A, x, k)

```
1  if  $k < x.key$ 
2      error "new key is smaller than current key"
3   $x.key = k$ 
4  find the index  $i$  in array  $A$  where object  $x$  occurs
5  while  $i > 1$  and  $A[\text{PARENT}(i)].key < A[i].key$ 
6      exchange  $A[i]$  with  $A[\text{PARENT}(i)]$ , updating the information that maps
        priority queue objects to array indices
7       $i = \text{PARENT}(i)$ 
```



➤ Insert new job

```
MAX-HEAP-INSERT( $A, x, n$ )
1  if  $A.heap\text{-}size == n$ 
2      error “heap overflow”
3   $A.heap\text{-}size = A.heap\text{-}size + 1$ 
4   $k = x.key$ 
5   $x.key = -\infty$ 
6   $A[A.heap\text{-}size] = x$ 
7  map  $x$  to index  $heap\text{-}size$  in the array
8  MAX-HEAP-INCREASE-KEY( $A, x, k$ )
```

► Linked Lists: Array Disadvantages


- You need to specify an initial size
- Changing the size of an array is troublesome
- Inserting and deleting elements in specific positions is difficult
- Let's say we want to delete 10 and keep the order of the rest:

A	5	8	10	13	16	19	27	46	51	86
A	5	8		13	16	19	27	46	51	86
A	5	8	13	16	19	27	46	51	86	


- What's the time complexity?

► Linked Lists

- Objects are linked using **pointers to the next element**.
- Linked lists can be **singly linked** or **doubly linked**: pointers to next and previous elements.
- Each element x has attributes
 - $x.key$ – the key used to identify the element
 - $x.next$ – a pointer to the next element
 - $x.prev$ – a pointer to the previous element
 - Optional: further satellite data



► Linked Lists: Searching



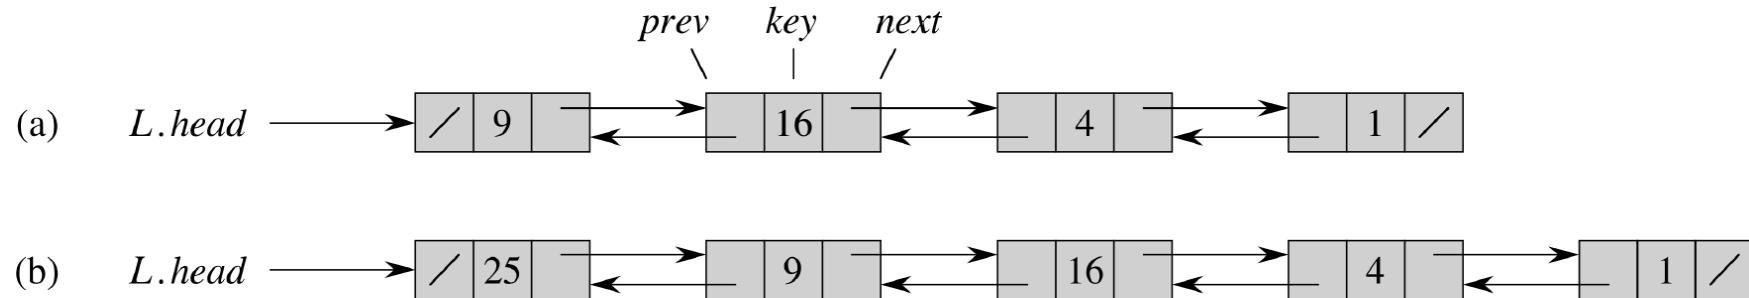
- Search inspects all elements in sequence and stops when the key has been found or the end of the list is reached.

LIST-SEARCH(L, k)

```
1:  $x = L.head$ 
2: while  $x \neq \text{NIL}$  and  $x.key \neq k$  do
3:    $x = x.next$ 
4: return  $x$ 
```

- The worst-case time is $\Theta(n)$, since it may have to search the entire list.

► Linked Lists: Inserting at the front




- New elements are added to the front of the list.

LIST-PREPEND(*L*, *x*)

```
1  x.next = L.head
2  x.prev = NIL
3  if L.head ≠ NIL
4      L.head.prev = x
5  L.head = x
```

- The time for an insertion is $O(1)$.

► Linked Lists: Inserting after element x




- New element added after element y.

```
LIST-INSERT( $x, y$ )
1   $x.next = y.next$ 
2   $x.prev = y$ 
3  if  $y.next \neq \text{NIL}$ 
4       $y.next.prev = x$ 
5   $y.next = x$ 
```

- The time for an insertion is $O(1)$ if you know the pointer to y

► Linked Lists: Deleting



- If element x is known, update pointers to take it out.

LIST-DELETE(L, x)

```

1: if  $x.prev \neq \text{NIL}$  then
2:    $x.prev.next = x.next$ 
3: else
4:    $L.head = x.next$ 
5: if  $x.next \neq \text{NIL}$  then
6:    $x.next.prev = x.prev$ 

```

- The time for a deletion is $O(1)$.
But if we only have the key and need to search the element x , it's time $\Theta(n)$ in the worst case.

► Summary

- **Stacks** and **Queues** are simple data structures that can
 - be implemented efficiently in arrays (modulo space issues)
 - Have a restricted set of operations, but these run in time $O(1)$.
- **Priority Queues**: all operations in at most $O(\log n)$ time
- Linked lists form an **unordered list** of elements
 - **Insertion** is fast if not important where it occurs: time $O(1)$.
 - **Searching** takes worst-case time $\Theta(n)$.
 - **Deletion** runs in time $O(1)$ if the element is known, otherwise we need to run a search beforehand and incur time $\Theta(n)$.