

# CS 305: Computer Networks

## Fall 2024

### **Lecture 5: Application Layer**

**Ming Tang**

Department of Computer Science and Engineering  
Southern University of Science and Technology (SUSTech)

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

# DNS: domain name system

People: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

- hostname, e.g.,  
www.yahoo.com -  
used by humans
- IP address (32 bit) -  
used for addressing  
datagrams

Q: how to map between IP  
address and name, and  
vice versa ?

Domain Name System (DNS):

- distributed database  
implemented in hierarchy of  
many *name servers*
- application-layer protocol: hosts  
and name servers communicate  
to *resolve* names (address/name  
translation)

# DNS Overview

- **DNS Services**
- DNS Structure
  - Hierarchical structure
  - Iterated and recursive query
- DNS protocol
  - DNS Records
  - Query and reply messages
- Inserting records into DNS

# DNS Services

---

- **hostname to IP address translation**
- **host aliasing**
  - canonical, alias hostnames
  - **www.ibm.com** (alias) is really **servereast.backup2.ibm.com** (canonical)
  - From supplied alias hostname to canonical hostname
- **mail server aliasing**
- **load distribution**
  - replicated Web servers: many IP addresses correspond to one name
  - rotation distributes the traffic (rotate the ordering of IP addresses)



# DNS Services

---

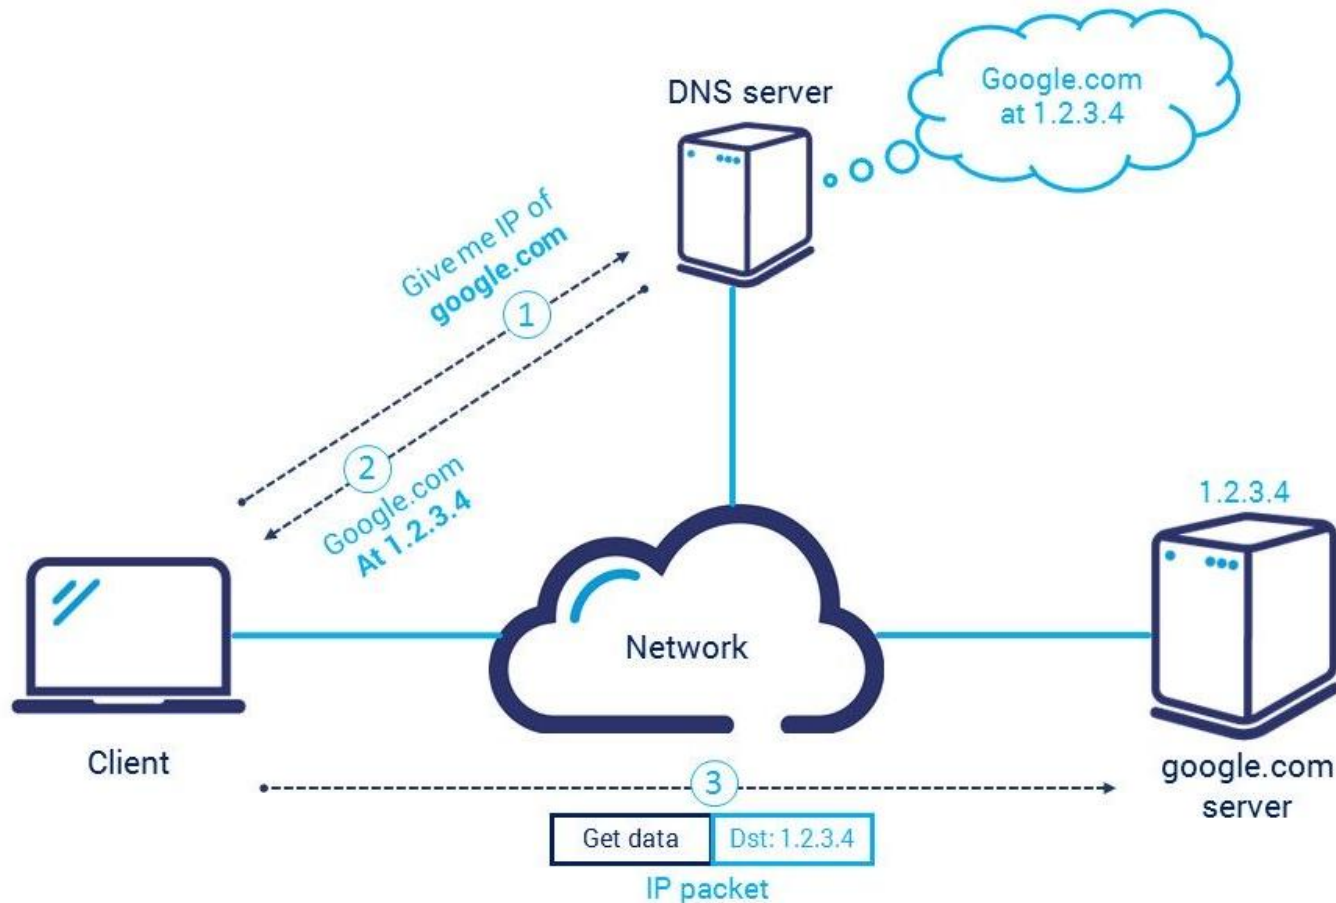
1. An application invokes the client side of DNS
  - specifying the **hostname** that needs to be translated
2. DNS in the user's host takes over, sending a query message into the network.
  - DNS query and reply messages
  - **UDP datagrams** to port 53.
3. After a delay, ranging from milliseconds to seconds, DNS in the user's host receives a DNS reply message that provides the desired mapping.
4. The **mapping (hostname - IP)** is then passed to the invoking application.

## Why UDP?

- fast speed
- smaller data packets

# DNS Services

From the perspective of the invoking application in the user's host, DNS is a **black box** providing a simple, straightforward translation service.



# DNS Overview

- DNS Services
- **DNS Structure**
  - Hierarchical structure
  - Iterated and recursive query
- DNS protocol
  - DNS Records
  - Query and reply messages
- Inserting records into DNS



# DNS Structure

---

## Centralized DNS:


Clients simply direct all queries to the single DNS server, and the DNS server responds directly to the querying clients.

### *Why not centralize DNS?*

- Single point of failure
- Traffic volume
- Distant centralized database
- Maintenance: huge database, update frequently

A: **doesn't scale!**

# DNS: a distributed, hierarchical database

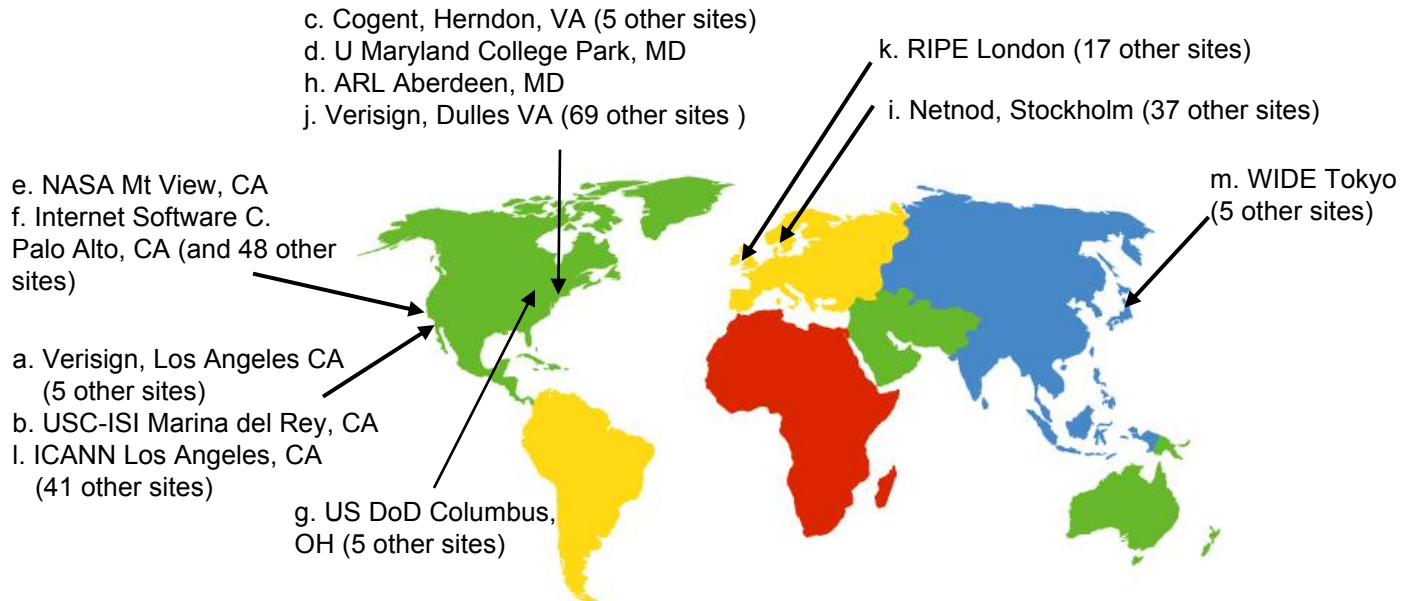


Client wants IP for [www.google.com](http://www.google.com) / [scholar.google.com](http://scholar.google.com):

- **Root DNS Servers:** find IP address of the [.com](#) TLD DNS server
- **Top-Level Domain (TLD) DNS:** client queries [.com](#) DNS server to get [google.com](http://google.com) authoritative DNS server
- **Authoritative DNS servers:** client queries [google.com](http://google.com) DNS server to get IP address for [www.google.com](http://www.google.com) / [scholar.google.com](http://scholar.google.com)

# DNS: root servers

- Root name server:
  - Provide the IP addresses of the TLD servers



13 logical root name  
“servers” worldwide

# TLD, authoritative servers

## Top-level domain (TLD) servers:

- Top-level domains: com, org, net, edu, aero, jobs, museums; top-level country domains: uk, fr, ca, jp
- *Network Solutions* maintains servers for .com TLD
- *Educause* for .edu TLD

## Authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

# Local DNS server

- Does not strictly belong to hierarchy
- Each ISP (residential ISP, company, university) has one
  - also called “default name server”

When a host connects to an ISP, the ISP provides the **IP addresses** of one or more of local DNS servers

- A host’s local DNS server may be typically “close to” the host

When host makes DNS query, query is sent to local DNS server

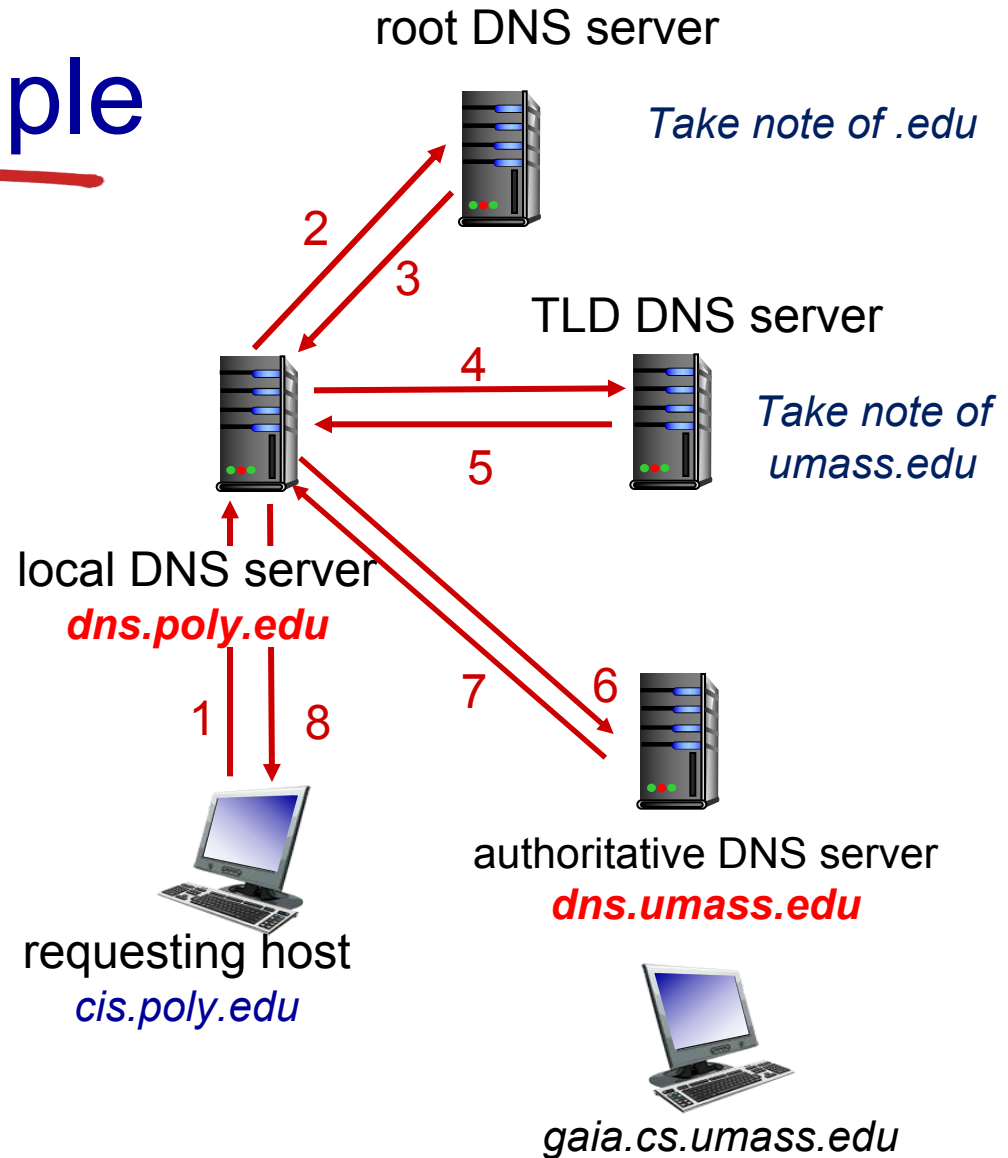
- **acts as proxy**, forwards query into hierarchy
- has **local cache** of recent name-to-address translation pairs (but may be out of date!)

# DNS name resolution example

- host at cis.poly.edu wants IP address for **gaia.cs.umass.edu**

## Iterated query:

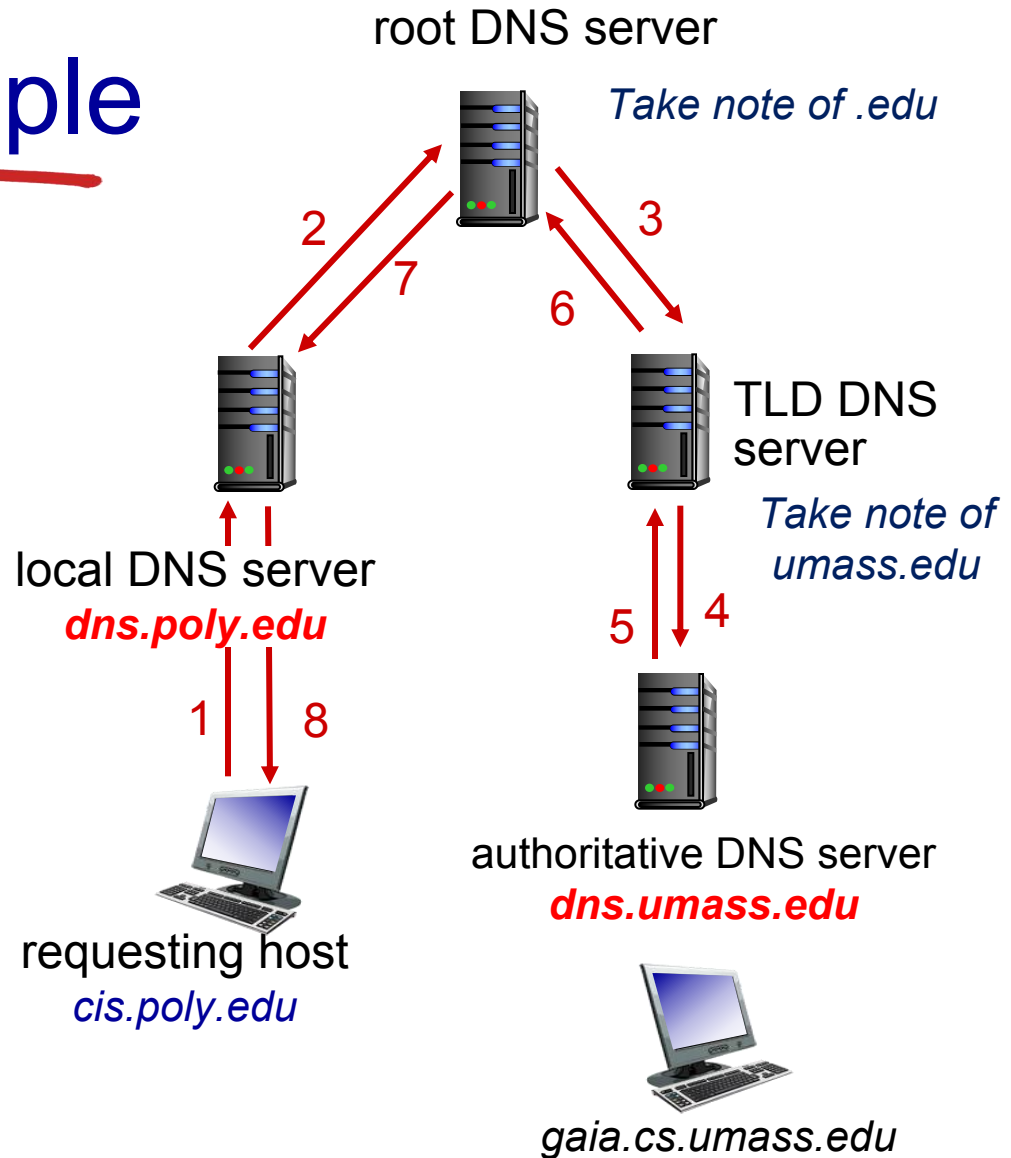
- contacted server replies with the name of another server to contact
- “I don’t know this name, but ask this server”



# DNS name resolution example

## Recursive query:

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?



# DNS: caching, updating records

- Once (any) name server learns mapping, it *caches* mapping
  - TLD servers typically cached in local name servers
  - thus root DNS servers not often visited
- Cached entries may be *out-of-date*
  - cache entries timeout (disappear) after some time (e.g., two days)
- Update/notify mechanisms proposed IETF standard
  - RFC 2136



# DNS Overview

- DNS Services
- DNS Structure
  - Hierarchical structure
  - Iterated and recursive query
- **DNS protocol**
  - DNS Records
  - Query and reply messages
- Inserting records into DNS

# DNS records

**DNS**: distributed database storing resource records (**RR**)

RR format: (name, value, type, ttl)

## type=A

- **name** is hostname
- **value** is IP address

## type=NS

- **name** is domain  
(e.g., foo.com)
- **value** is hostname of  
authoritative server for  
this domain  
(e.g., dns.foo.com)

## type=CNAME

- **name** is alias name for some  
“canonical” (the real) name
- **www.ibm.com** is really  
**servereast.backup2.ibm.com**
- **value** is canonical name

## type=MX

- **value** is canonical name of the  
**mailserver** with **name** (alias name)

# DNS records

If a DNS server is **authoritative** for a particular hostname

- the DNS server will contain a Type A record for the hostname
- (Even if the DNS server is not authoritative, it may contain a Type A record in its cache.)

If a server is **not authoritative for** a hostname

- the server will contain a Type NS record for the domain that includes the hostname
- it will also contain a Type A record that provides the IP address of the DNS server in the **Value** field of the NS record.

Example: an .edu TLD server is not authoritative for gaia.cs.umass.edu

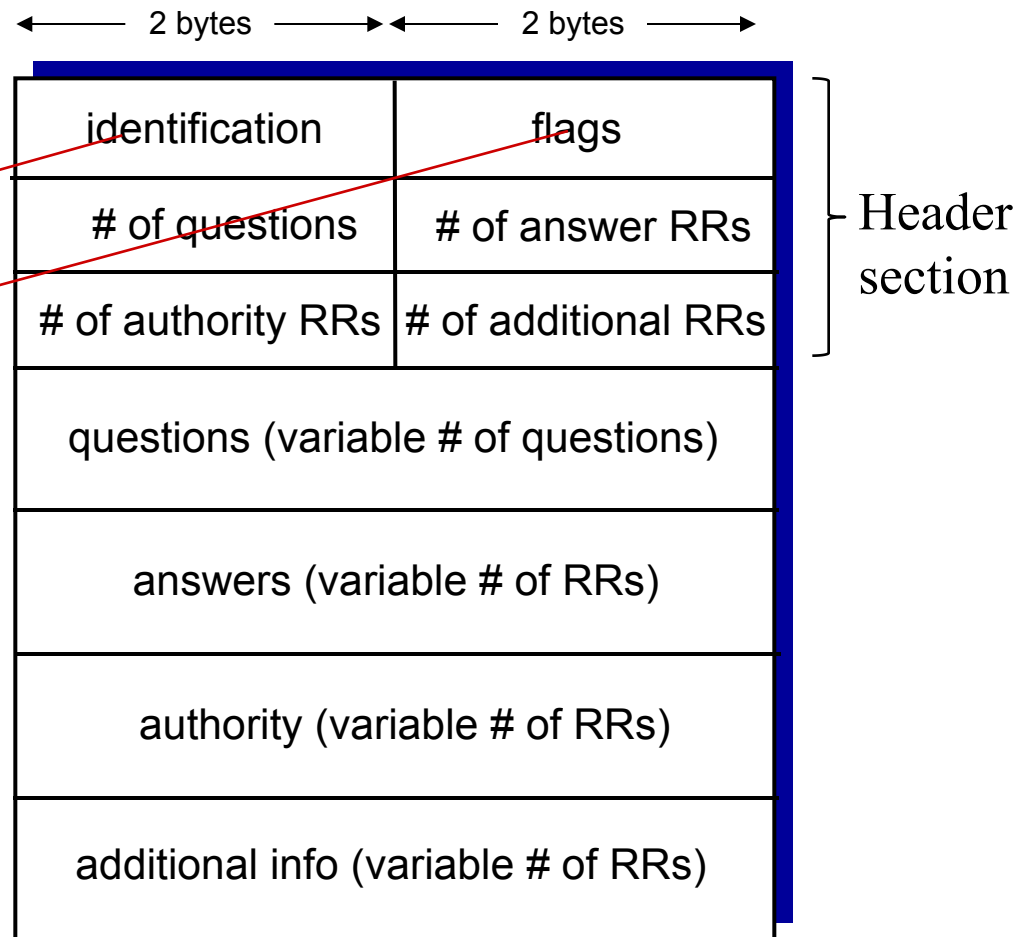
- (umass.edu, dns.umass.edu, NS) .
- (dns.umass.edu, 128.119.40.111, A)

# DNS protocol, messages

Query and reply messages, both with same message format

message header

- **identification**: 16 bit number for query, reply to query uses same number
- **flags**:
  - query or reply
  - recursion desired
  - recursion available
  - reply is authoritative



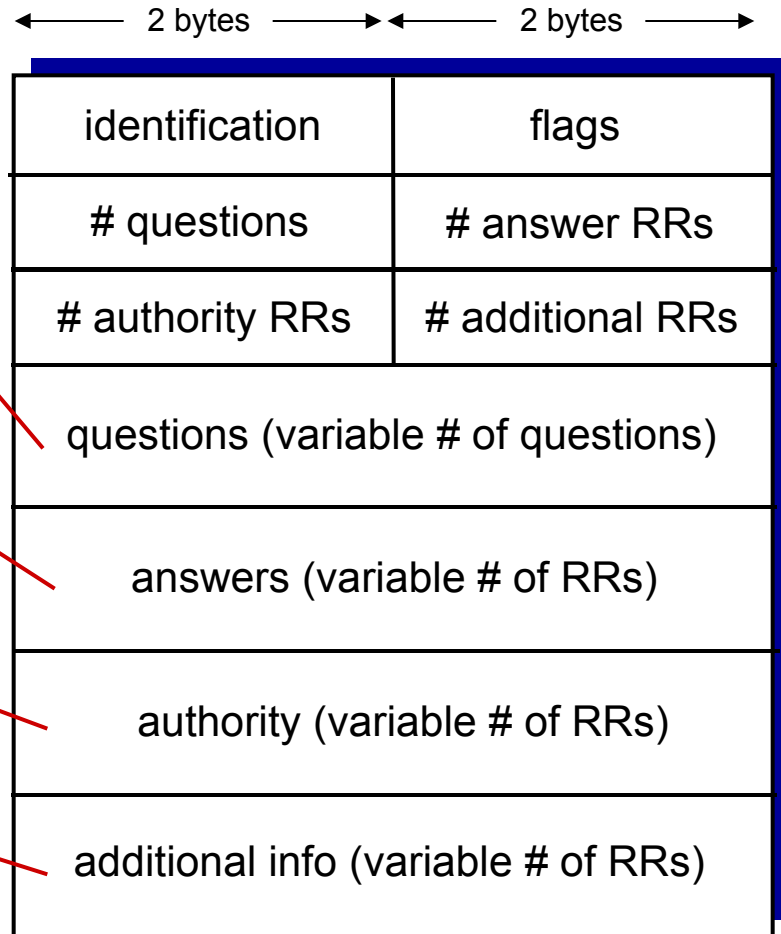
# DNS protocol, messages

Name & type fields  
(e.g., Type A or Type MX)

RRs in response  
to query  
(a reply can return  
multiple RRs)

records of other  
authoritative servers

additional “helpful”  
info that may be used



# DNS protocol, messages

For example, a reply to **an MX query**

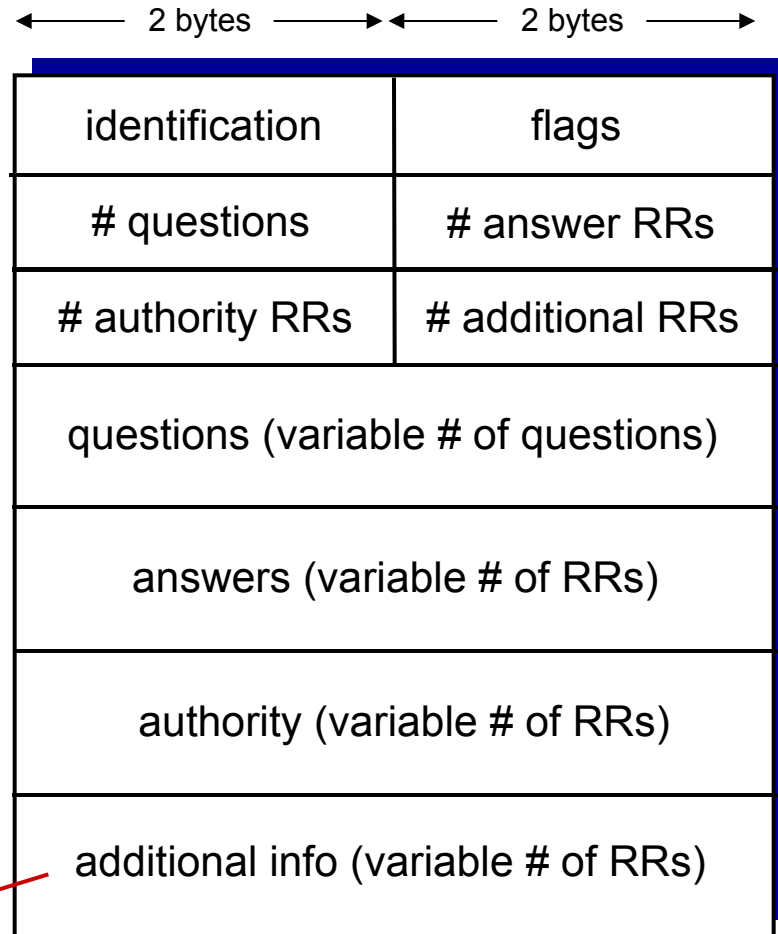
Answer section: Type MX

- an RR providing the canonical hostname of a mail server.

Additional section: Type A

- the IP address for the canonical hostname of the mail server.

additional “helpful”  
info that may be used



# DNS Overview

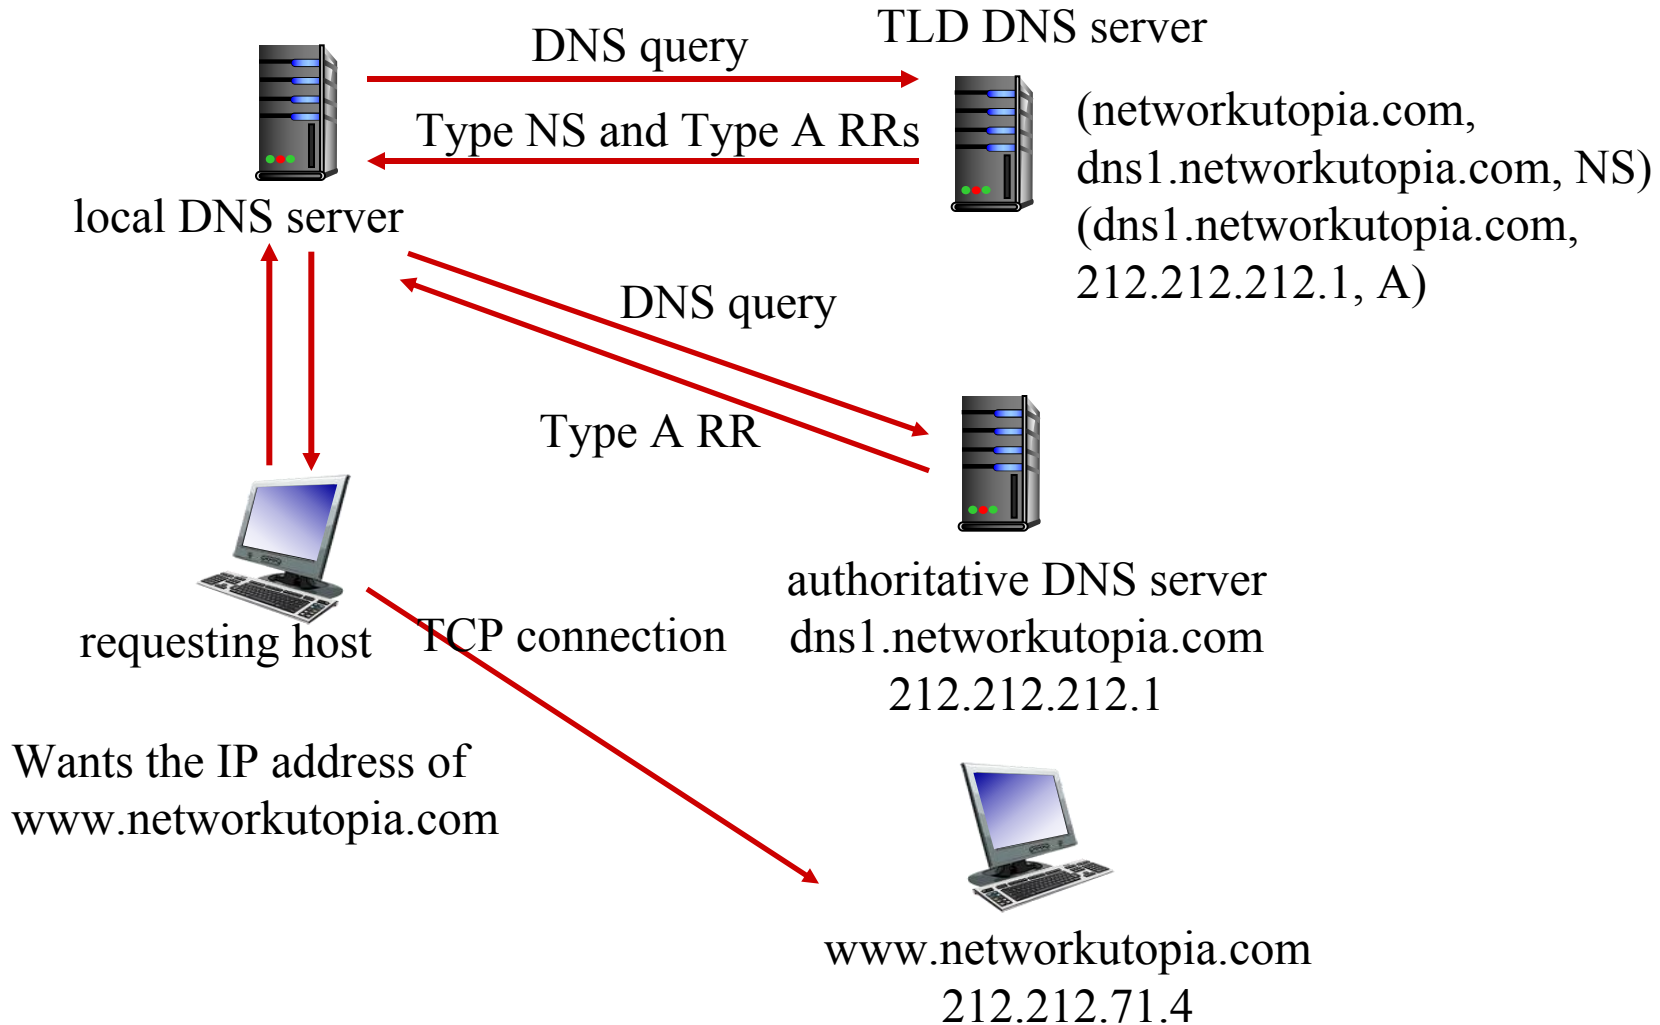
- DNS Services
- DNS Structure
  - Hierarchical structure
  - Iterated and recursive query
- DNS protocol
  - DNS Records
  - Query and reply messages
- Inserting records into DNS server

# Inserting records into DNS

- Example: new startup “Network Utopia”
- Register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
  - provide names, IP addresses of authoritative DNS server (primary and secondary)
  - registrar inserts two RRs into .com TLD server:  
(networkutopia.com, dns1.networkutopia.com, NS)  
(dns1.networkutopia.com, 212.212.212.1, A)



# Inserting records into DNS



# Attacking DNS

## Distributed denial-of-service (DDoS) attacks

- bombard root servers with traffic
  - not successful to date
  - traffic filtering
  - local DNS servers cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers
  - potentially more dangerous

## Redirect attacks

- man-in-middle
  - Intercept queries; bogus reply
- DNS poisoning
  - Send bogus replies to DNS server

## Exploit DNS for DDoS

- target IP
- Redirect an unsuspecting Web user to attack Web site

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks


2.7 socket programming with UDP and TCP

# Pure P2P architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

## Examples:

- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)




# DNS Overview

- P2P vs Client Server
- BitTorrent

# File distribution: client-server vs P2P

**Question:** How much time to distribute file (size  $F$ ) from one server to  $N$  peers?


- peer upload/download capacity is limited resource
- **Distribution time:** the time it takes to get a copy of the file to all  $N$  peers.



# File distribution time: client-server

- **Server transmission:** must sequentially send (upload)  $N$  file copies:

- time to send one copy:  $F/u_s$
- time to send  $N$  copies:  $NF/u_s$



- **Client:** each client must download file copy

- $d_{min}$  = min client download rate
- maximum client download time:  $F/d_{min}$

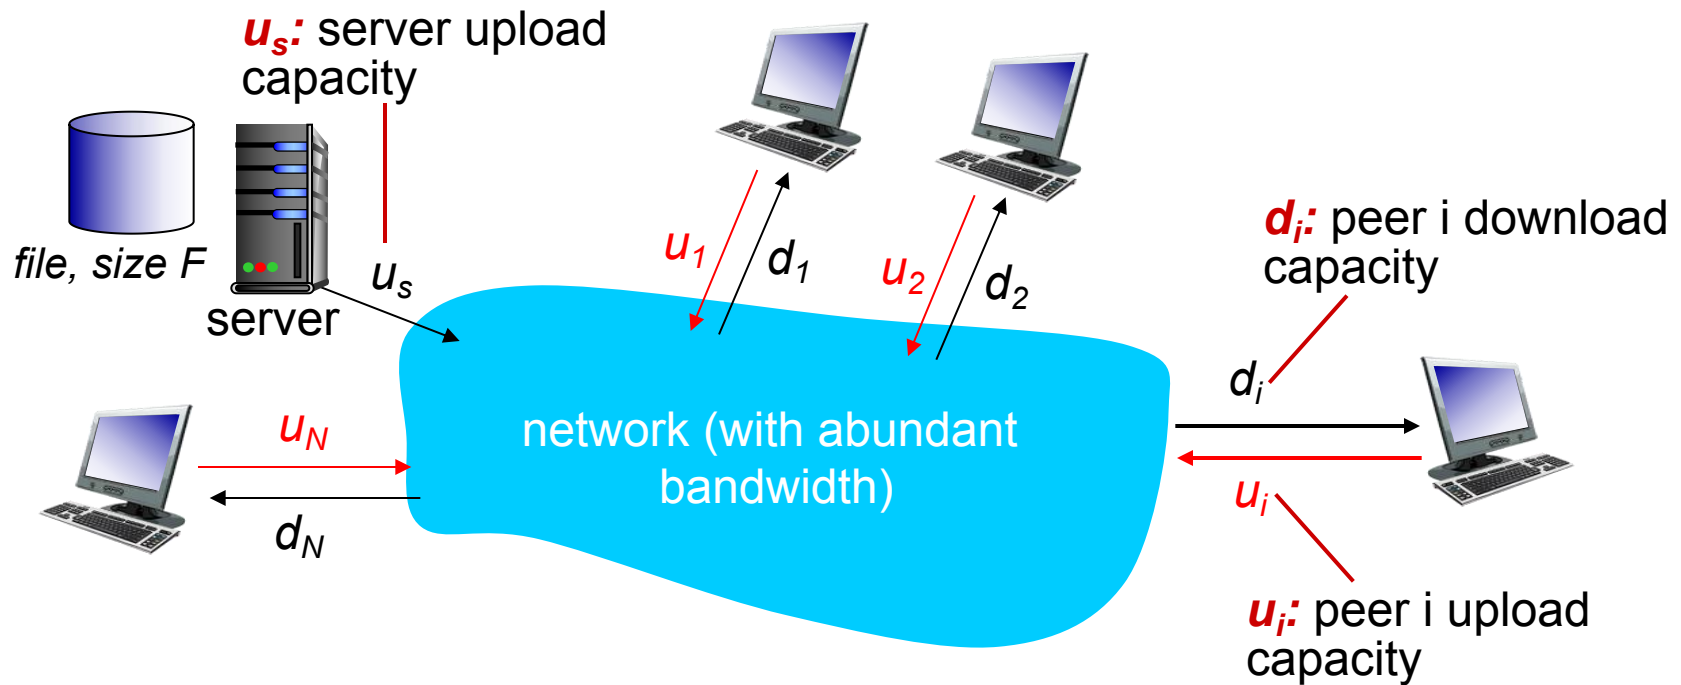
time to distribute  $F$   
to  $N$  clients using  
client-server approach

$$D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

increases linearly in  $N$

# File distribution time: P2P


In P2P model, clients are both downloaders and uploaders.





# File distribution time: P2P

- **Server transmission:** must upload at least one copy
  - time to send one copy:  $F/u_s$
- **Client downloading:** each client must download file copy
  - maximum client download time:  $F/d_{\min}$
- **Clients and server:** delivering a total of  $NF$  bits
  - max upload rate (limiting max download rate) is  $u_s + \sum u_i$



time to distribute  $F$   
to  $N$  clients using  
P2P approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$


If each peer can redistribute a bit as soon as it receives the bit, then there is a scheme that actually achieves this lower bound

increases linearly in  $N$  ...

... but so does this, as each peer brings service capacity

# Client-server vs. P2P: example

client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$



# DNS Overview


- P2P vs Client Server
- BitTorrent

# P2P file distribution: BitTorrent

- File divided into 256Kb chunks
- Peers in BitTorrent send/receive file chunks


*tracker*: tracks peers participating in BitTorrent

*torrent*: group of peers exchanging chunks of a file



# P2P file distribution: BitTorrent

- Peer **joining** BitTorrent:
  - has no chunks, but will accumulate them over time from other peers
  - registers with tracker to get **list of peers**
  - TCP connections with subset of peers ( **“neighbors”** )
- While **downloading**, peer uploads chunks to other peers
  - Peers may leave
  - Peers may come, initiating connections with Alice
- Once peer has entire file, it may (selfishly) leave or (altruistically) remain in BitTorrent



# BitTorrent: requesting, sending file chunks

Q1: which chunks should she request first from her neighbors?

Q2: to which of her neighbors should she send requested chunks?

## requesting chunks:


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each “neighbor” for list of chunks that they have
- Alice requests missing chunks from peers, **rarest first**

## sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks **at highest rate**
  - other peers are choked by Alice (do not receive chunks from her)
  - re-evaluate every 10 secs
- every 30 secs: **randomly** select one **additional** peer, starts sending chunks
  - “optimistically unchoke” this peer
  - newly chosen peer may join top 4

# BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP




# Video Streaming and CDNs: ~~context~~

- Video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- Challenge: scale - how to reach ~1B users?
  - single mega-video server won't work (why?)
- Challenge: heterogeneity
  - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- **Solution:** distributed, application-level infrastructure



# Multimedia: video

- Video: sequence of images displayed at constant rate
  - e.g., 24 images/sec
- Digital image: array of pixels
  - each pixel represented by bits




# Multimedia: video

Coding (Compression): use redundancy *within* and *between* images to decrease # bits used to encode image

- spatial (within image)
- temporal (from one image to next)

**Spatial coding example:** instead of sending N values of same color (all purple), send only two values: color value (purple) and number of repeated values (N)



frame  $i$



frame  $i+1$


**temporal coding example:** instead of sending complete frame at  $i+1$ , send only differences from frame  $i$

# Multimedia: video

Type	Video Bitrate, Standard Frame Rate (24, 25, 30)	Video Bitrate, High Frame Rate (48, 50, 60)
2160p (4k)	35-45 Mbps	53-68 Mbps
1440p (2k)	16 Mbps	24 Mbps
1080p	8 Mbps	12 Mbps
720p	5 Mbps	7.5 Mbps
480p	2.5 Mbps	4 Mbps
360p	1 Mbps	1.5 Mbps

- **CBR: (constant bit rate):** video encoding rate fixed
- **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes

# HTTP Streaming



All clients receive the same encoding of the video:

- Human users may have different requirements
- Clients may have different available bandwidth, which may be time-varying

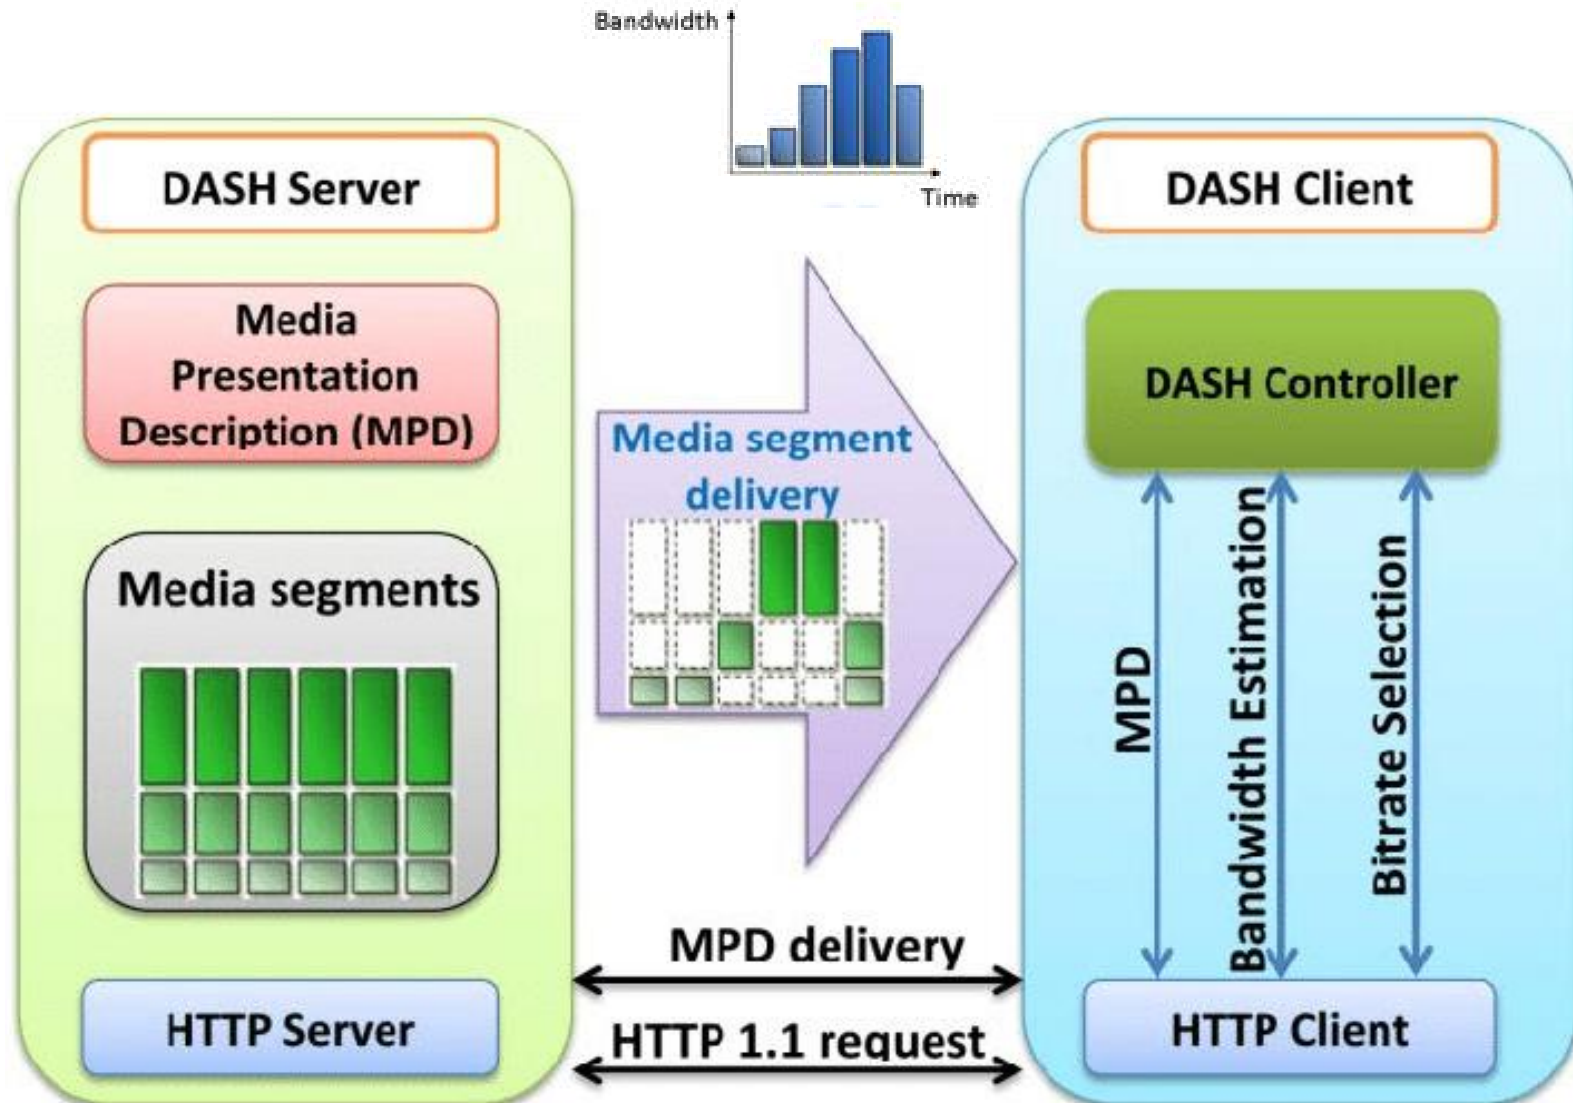
How to deal with this?

# Streaming multimedia: DASH

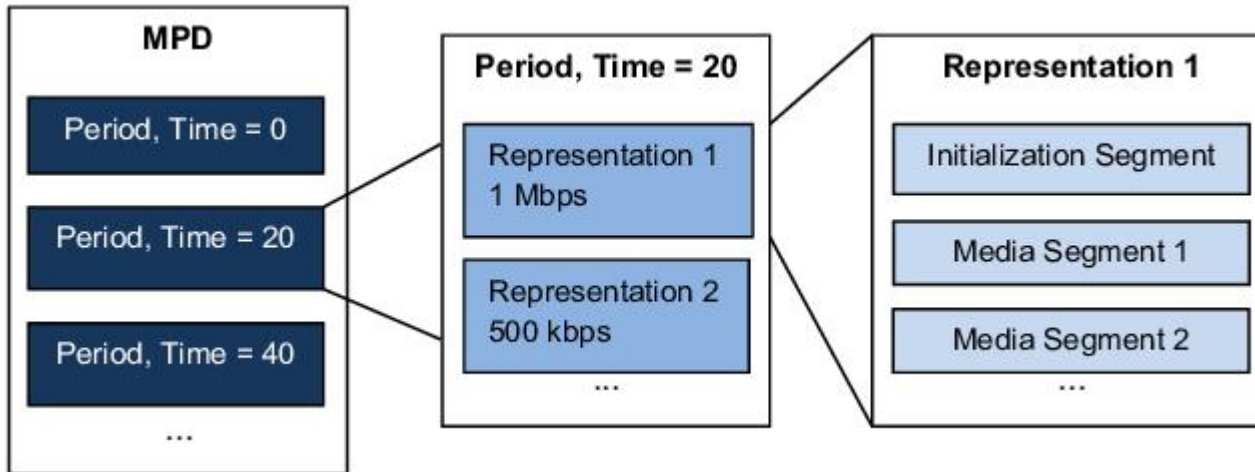
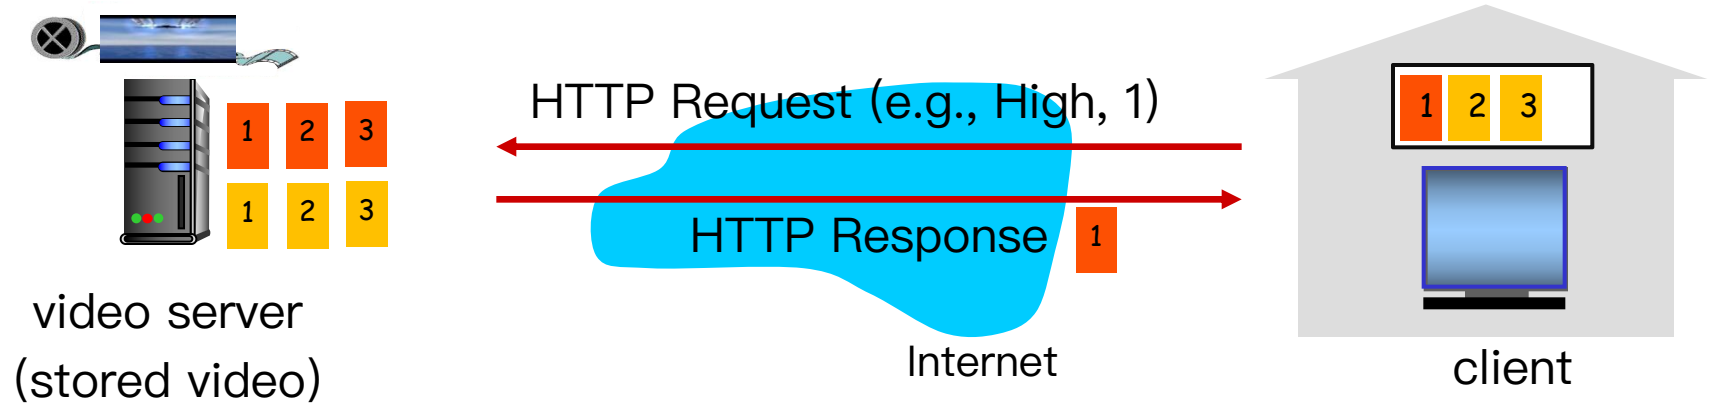
---

- **DASH: D**ynamic, **A**daptive **S**treaming over **H**TTP
- **Server:**
  - divides video file into **multiple chunks**
  - each chunk stored, encoded at different rates
  - **manifest file:** provides URLs for different chunks encoded at different rates
- **Client:**
  - periodically measures server-to-client bandwidth
  - consulting manifest, requests one chunk at a time
    - chooses maximum coding rate sustainable given current bandwidth
    - can choose different coding rates at different points in time (depending on available bandwidth at time)

# Streaming multimedia: DASH



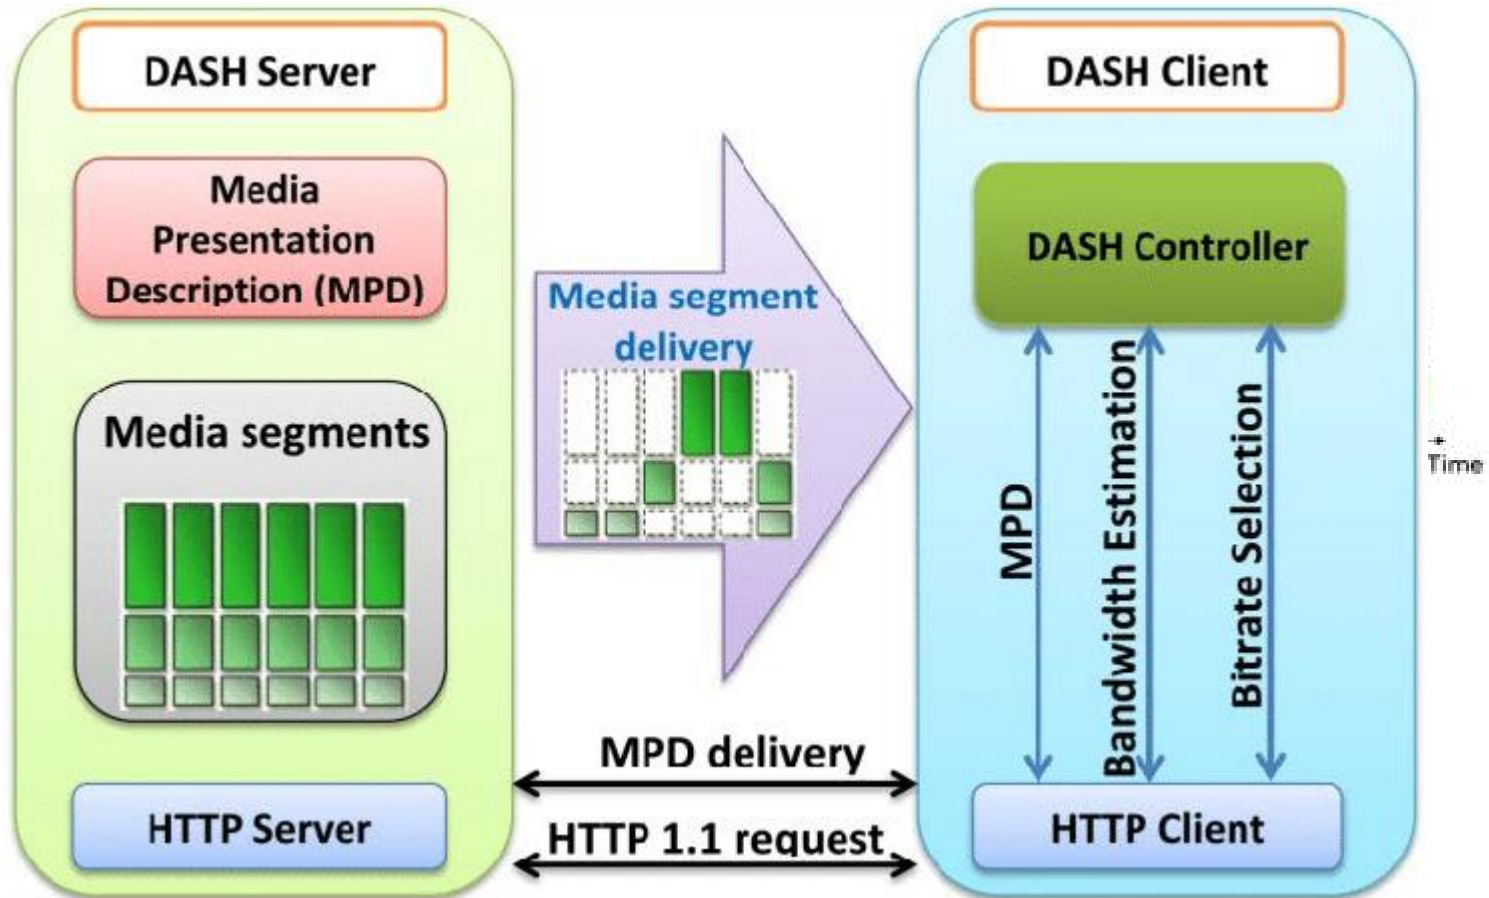
# Streaming multimedia: DASH



Media presentation description (MPD), also known as manifest



# Streaming multimedia: DASH



# Streaming multimedia: DASH

---

“intelligence” at client: client determines

- **when** to request chunk (so that buffer starvation, or overflow does not occur)
- **what encoding rate** to request (higher quality when more bandwidth available)
- **where** to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications


2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

# Content distribution networks





来源: 小黑羊JoinWings  
[https://baijia.baidu.com/s?old\\_id=126615&wfr=pc&fr=app\\_list](https://baijia.baidu.com/s?old_id=126615&wfr=pc&fr=app_list)






# Content distribution networks





# Content distribution networks



# Content distribution networks

---

- **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- **Option 1:** single, large “mega-server”
  - single point of failure
  - huge traffic
  - long path to distant clients
  - multiple copies of video sent over outgoing link

....quite simply: this solution **doesn't scale**

# Content distribution networks


---

- **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- **Option 2: Content distribution networks (CDN)** store/serve multiple copies of videos at multiple geographically distributed sites
  - **Enter deep:** push CDN servers deep into many access networks; inside ISPs
    - close to users
    - used by Akamai, 1700 locations
  - **Bring home:** smaller number (10's) of larger clusters in Internet Exchange Point (IXP); outside ISPs
    - used by Limelight



# Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
  - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
  - directed to nearby copy, retrieves content
  - may choose different copy if network path congested



# Content Distribution Networks (CDNs)

---

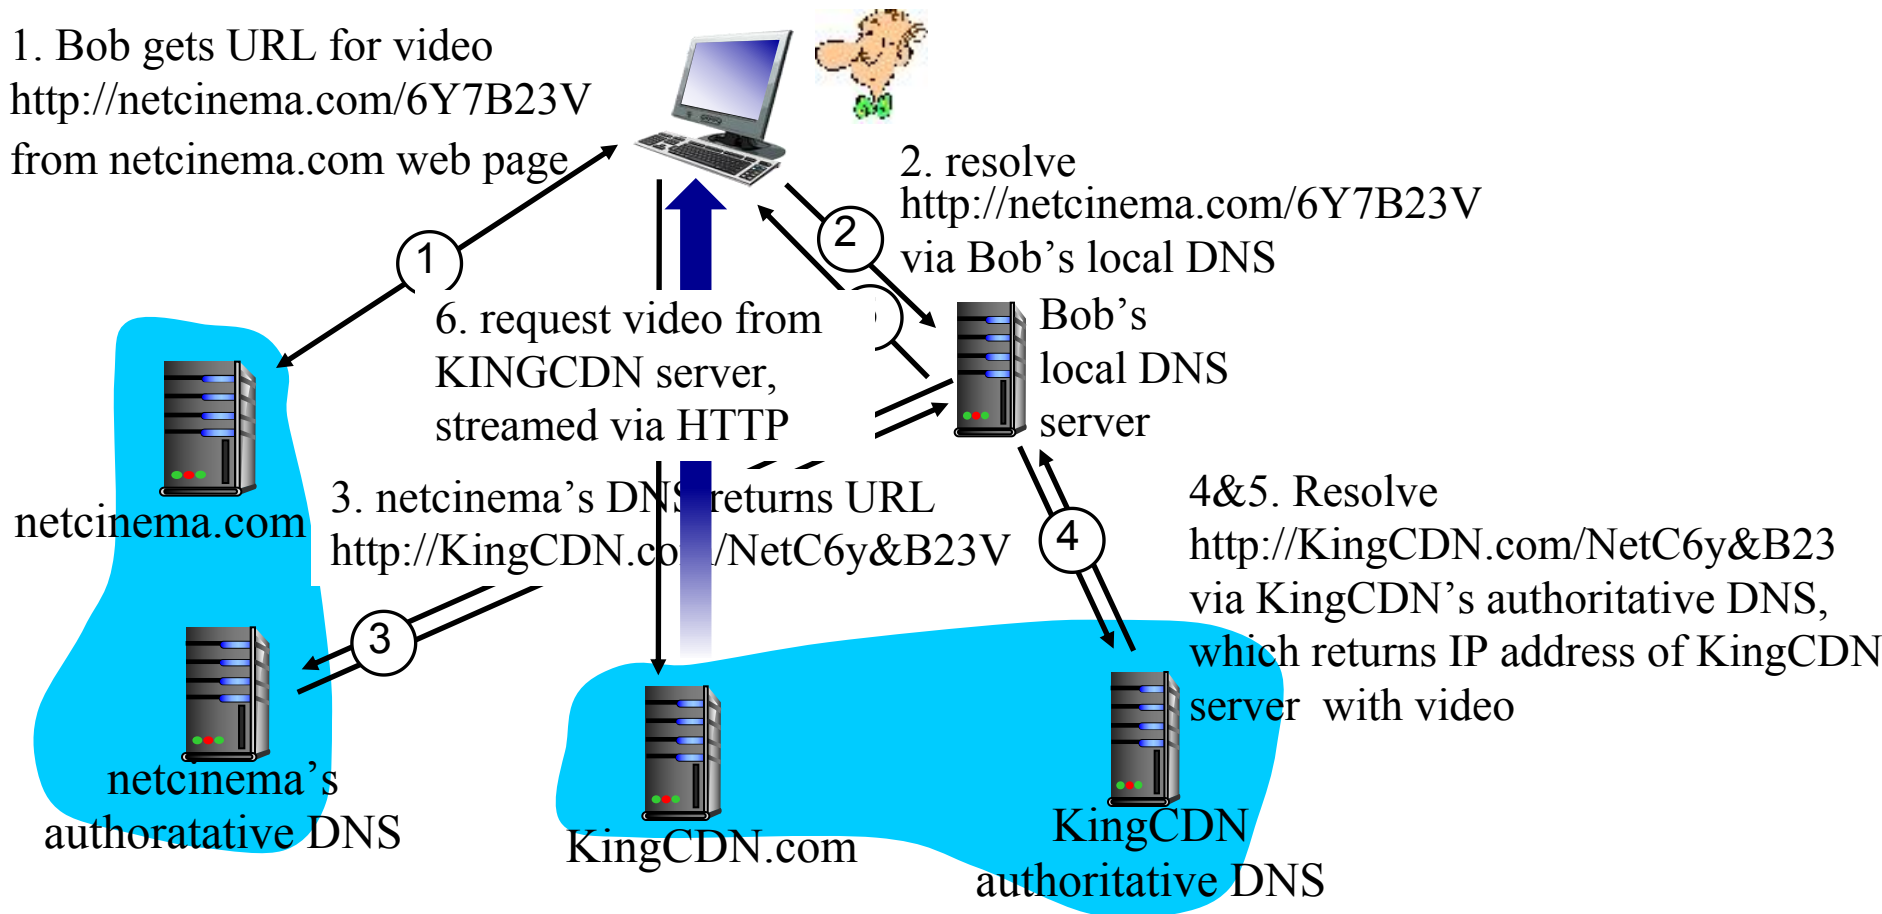
## Challenges: Coping with a congested Internet

- what content to place in CDN node?
  - Simple pull strategy: request, then store
- from which CDN node to retrieve content?
  - Cluster selection strategy
- the operation for retrieving content?
  - CDN operation

# CDN Operation

Bob (client) requests video <http://netcinema.com/6Y7B23V>

- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



# CDN: Cluster Selection Strategy

---

One simple strategy is to assign the client to the cluster that is

**geographically closest:**

- When a DNS request is received from a particular LDNS, the CDN chooses the geographically closest cluster
- may not be the closest cluster in terms of the length or number of hops
- ignore the variation in delay and available bandwidth over time

Periodic **real-time measurements** of delay and loss performance between their clusters and clients:

- a CDN can have each of its clusters periodically send probes to all of the LDNSs around the world.
- many LDNSs are configured to not respond to such probes.

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP