

CS 305: Computer Networks

Fall 2024

Lecture 6: Transport Layer

Ming Tang

Department of Computer Science and Engineering
Southern University of Science and Technology (SUSTech)

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

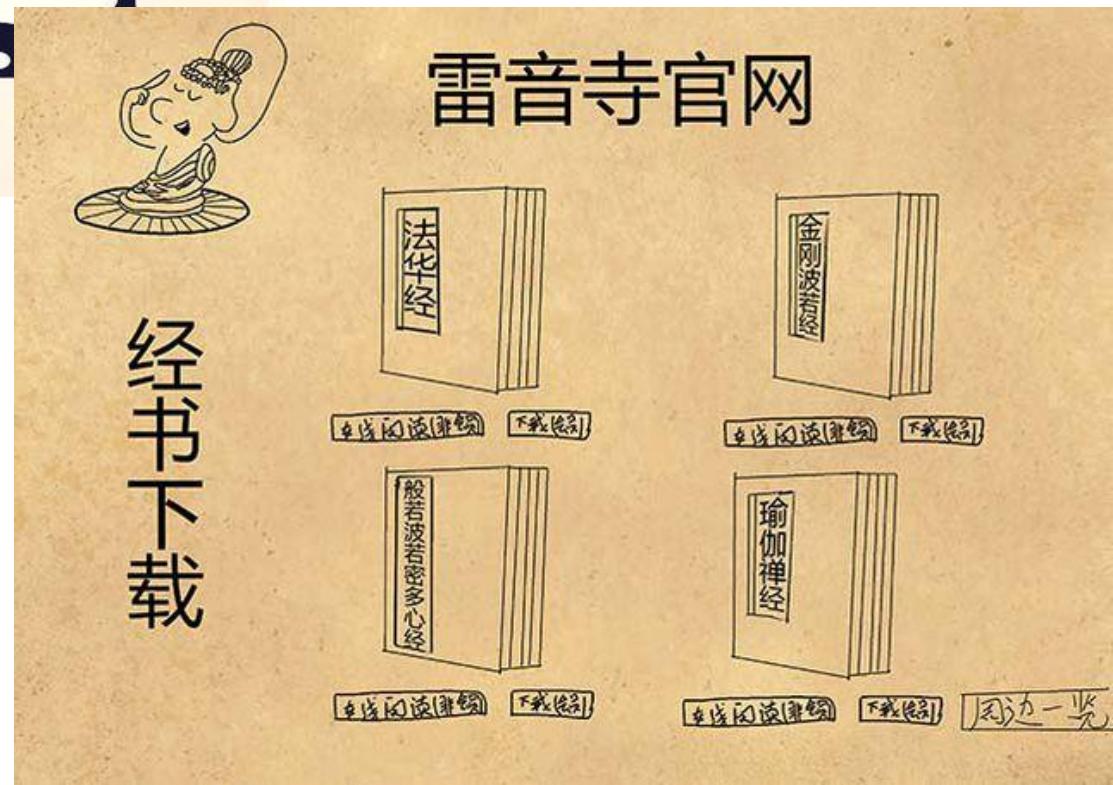
2.6 video streaming and **content distribution networks**

2.7 socket programming with UDP and TCP

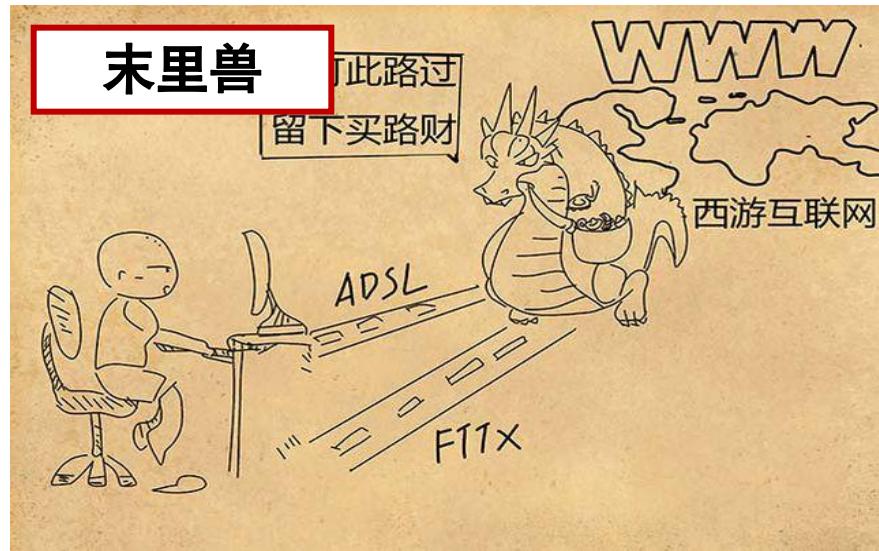
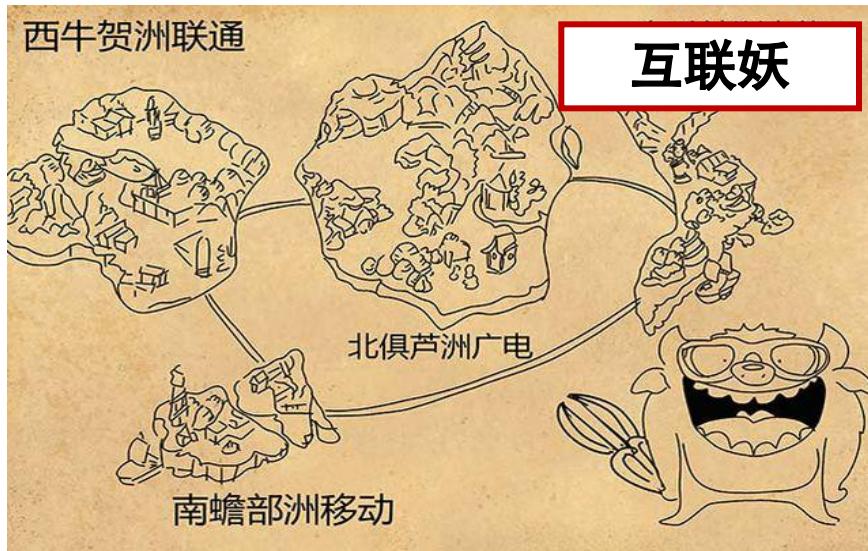
Content distribution networks



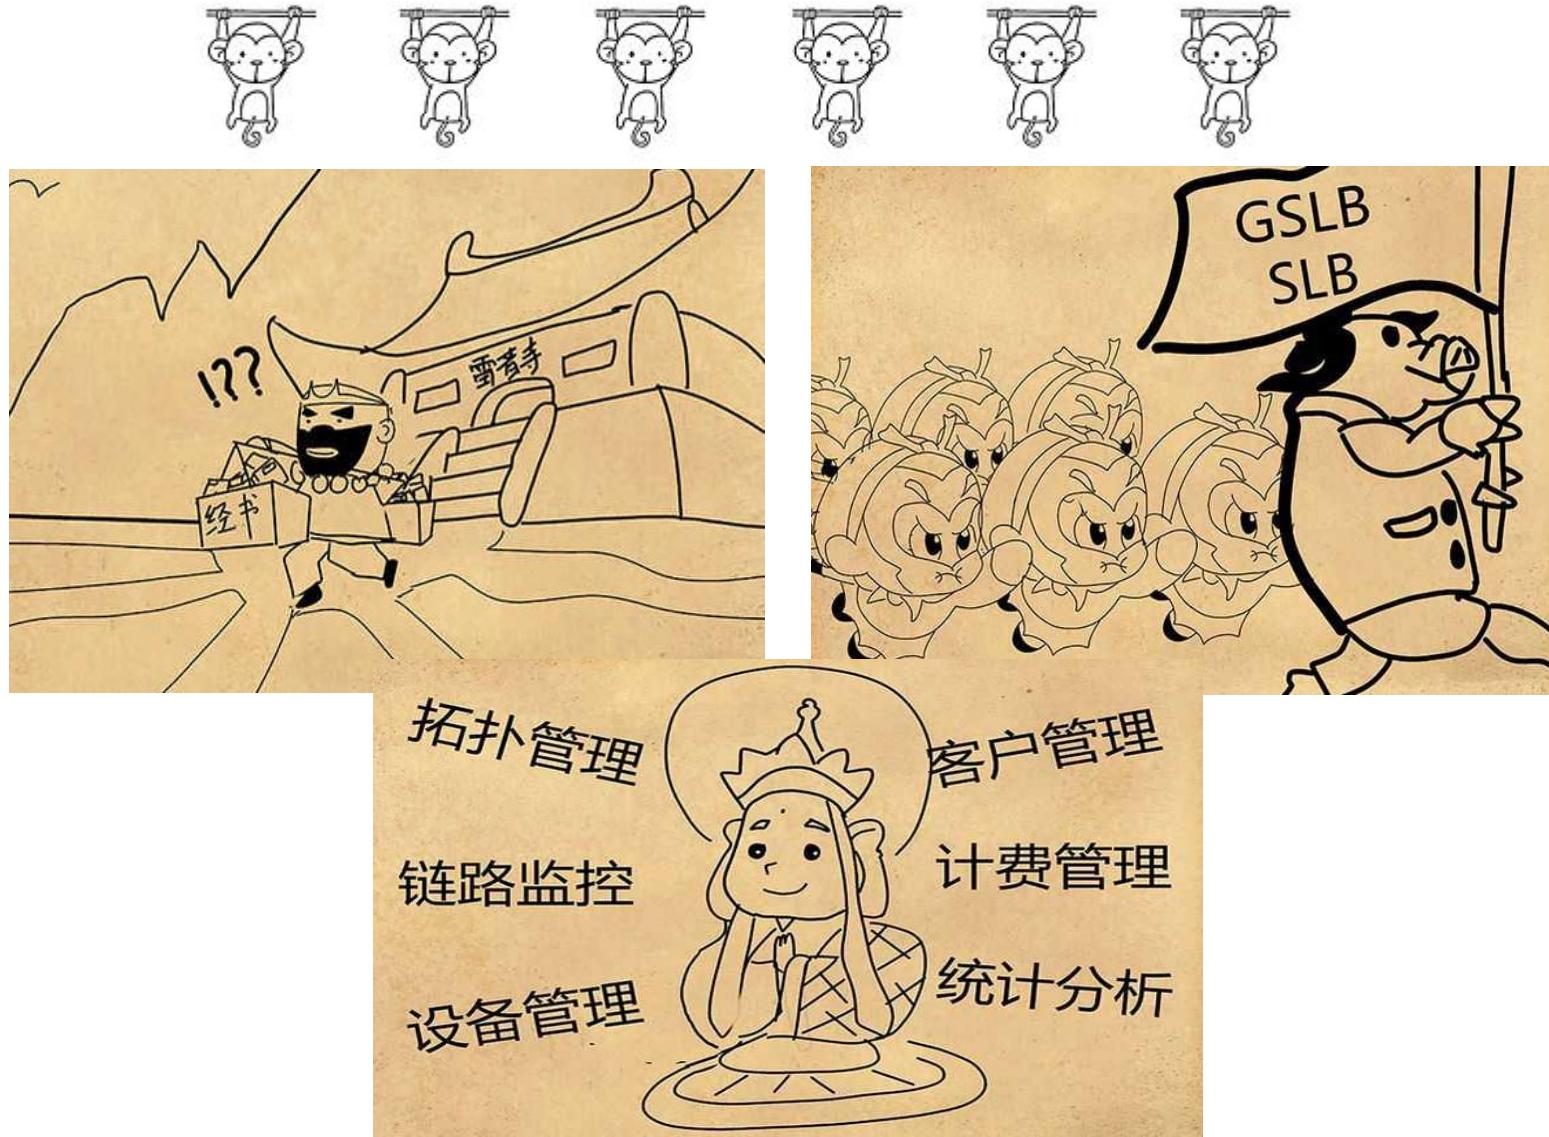
来源：小黑羊JoinWings
[https://baijia.baidu.com/
s?old_id=126615&wfr=p
c&fr=app_list](https://baijia.baidu.com/s?old_id=126615&wfr=p&fr=app_list)



Content distribution networks



Content distribution networks



Content distribution networks

- ❖ **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- ❖ **Option 1:** single, large “mega-server”
 - single point of failure
 - huge traffic
 - long path to distant clients
 - multiple copies of video sent over outgoing link

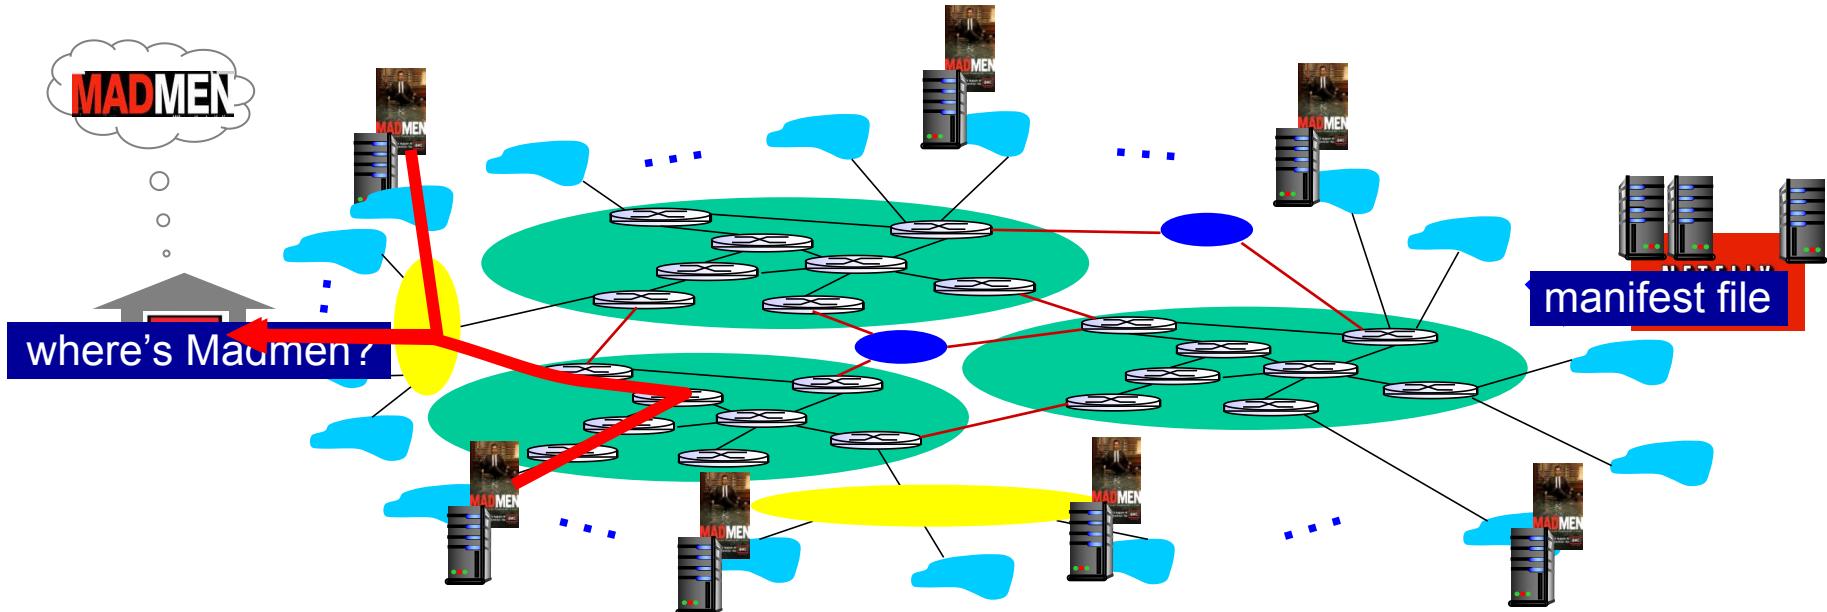
....quite simply: this solution **doesn't scale**

Content distribution networks

- ❖ **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- ❖ **Option 2: Content distribution networks (CDN)** store/serve multiple copies of videos at multiple geographically distributed sites
 - **Enter deep:** push CDN servers deep into many access networks; inside ISPs
 - close to users
 - used by Akamai, 1700 locations
 - **Bring home:** smaller number (10's) of larger clusters in Internet Exchange Point (IXP); outside ISPs
 - used by Limelight

Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
 - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
 - directed to nearby copy, retrieves content
 - may choose different copy if network path congested



Content Distribution Networks (CDNs)

Challenges: Coping with a congested Internet

- what content to place in CDN node?
 - Simple pull strategy: request, then store
- from which CDN node to retrieve content?
 - Cluster selection strategy
- the operation for retrieving content?
 - CDN operation

CDN Operation

Bob (client) requests video <http://netcinema.com/6Y7B23V>

- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>

1. Bob gets URL for video
<http://netcinema.com/6Y7B23V>
from netcinema.com web page

netcinema.com

netcinema's
authoritative DNS



1
6. request video from
KINGCDN server,
streamed via HTTP

2. resolve
<http://netcinema.com/6Y7B23V>
via Bob's local DNS

Bob's
local DNS
server

3. netcinema's DNS returns URL
<http://KingCDN.com/NetC6y&B23V>

KingCDN.com

KingCDN
authoritative DNS

4&5. Resolve
<http://KingCDN.com/NetC6y&B23V>
via KingCDN's authoritative DNS,
which returns IP address of KingCDN
server with video

CDN: Cluster Selection Strategy

One simple strategy is to assign the client to the cluster that is **geographically closest**:

- When a DNS request is received from a particular LDNS, the CDN chooses the geographically closest cluster
- may not be the closest cluster in terms of the length or number of hops
- ignore the variation in delay and available bandwidth over time

Periodic **real-time measurements** of delay and loss performance between their clusters and clients:

- a CDN can have each of its clusters periodically send probes to all of the LDNSs around the world.
- many LDNSs are configured to not respond to such probes.

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

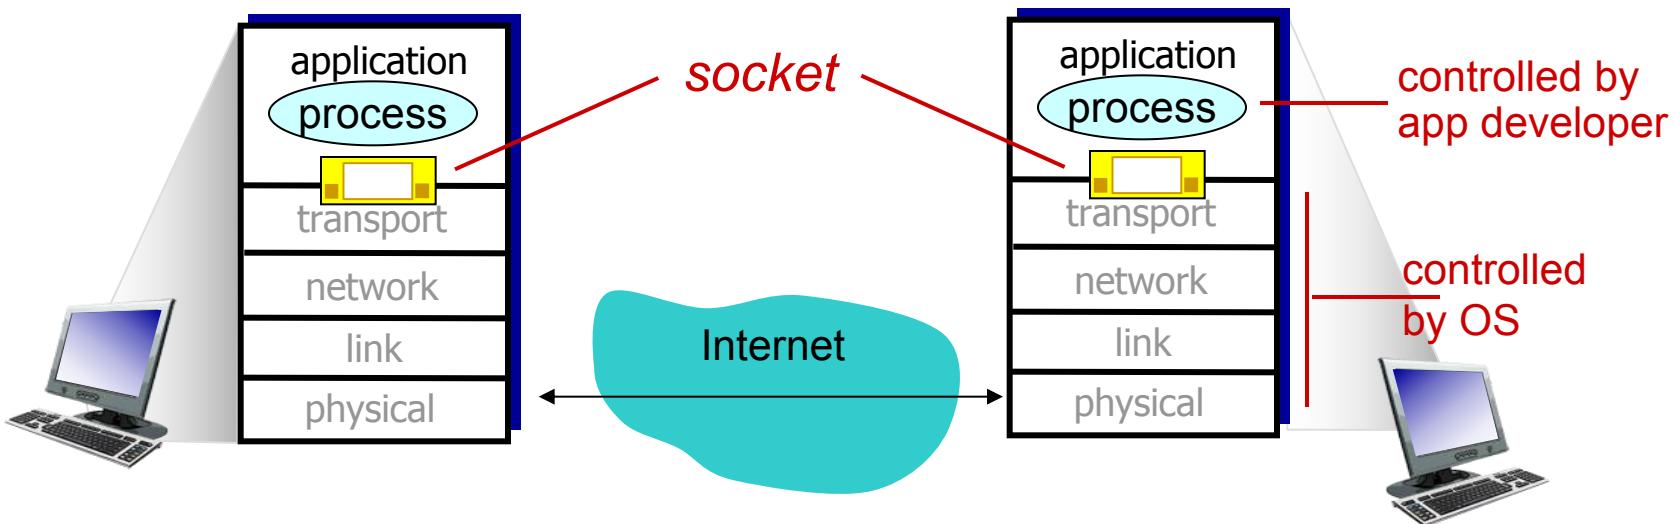
2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Socket programming

Goal: learn how to build client/server applications that communicate using sockets

Socket: door between application process and end-end-transport protocol



Socket programming

Socket programming: how we can use socket API for creating communication between client and server processes.

Two socket types for two transport services:

- **UDP**: unreliable datagram
- **TCP**: reliable, connection-oriented

Application Example:

1. client reads a line of characters (data) from its keyboard and sends data to server
2. server receives the data and converts characters to uppercase
3. server sends modified data to client
4. client receives modified data and displays line on its screen

Socket programming with UDP

UDP: no “connection” between client & server

- ❖ no handshaking before sending data
- ❖ sender explicitly attaches **destination IP address and port number** to each packet
- ❖ receiver extracts sender IP address and port number from received packet

UDP: transmitted data may be lost or received out-of-order

Application viewpoint:

- ❖ UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server

Client/server socket interaction: ~~UDP~~

server (running on serverIP)

create socket, port= x:

```
serverSocket =  
socket(AF_INET,SOCK_DGRAM)
```

read UDP datagram from
serverSocket

write reply to **serverSocket**
specifying client address,
port number

client

create socket:

```
clientSocket =  
socket(AF_INET,SOCK_DGRAM)
```

Create datagram with server IP and
port=x; send datagram via
clientSocket

read datagram from
clientSocket

close
clientSocket

Segment: the transport-layer packet for TCP

Datagram: the packet for UDP

Example app: UDP server

Python UDPServer

```
include Python's socket  
library → from socket import *  
  
create UDP socket → serverSocket = socket(AF_INET, SOCK_DGRAM)  
  
bind socket to local port  
number 12000 → serverSocket.bind(("", serverPort))  
print ( "The server is ready to receive")  
  
loop forever → while True:  
  
Read from UDP socket into  
message, getting client's  
address (client IP and port) → message, clientAddress = serverSocket.recvfrom(2048)  
modifiedMessage = message.decode().upper()  
  
send upper case string → serverSocket.sendto(modifiedMessage.encode(),  
clientAddress)
```

IPv4
UDP socket
UDP socket is identified by destination IP address and port number
IP + portnumber

Example app: UDP client

Python UDPClient

We did not
specify the client
port number

Create the client's socket

get user keyboard
input

Attach server name, port to
message; send into socket

read reply characters from
socket into string

print out received string
and close socket

```
from socket import *
serverName = 'hostname'
serverPort = 12000
clientSocket = socket(AF_INET,
                      SOCK_DGRAM)
message = raw_input('Input lowercase sentence:')
clientSocket.sendto(message.encode(),
                     (serverName, serverPort))
modifiedMessage, serverAddress =
clientSocket.recvfrom(2048)
print modifiedMessage.decode()
clientSocket.close()
```

either the IP address (e.g.,
“128.138.32.126”) or the hostname
(e.g., “cis.poly.edu”)

Socket programming with TCP

Client must contact server

- ❖ server process must first be running
- ❖ server must have created socket (door) that welcomes client's contact

Client contacts server by:

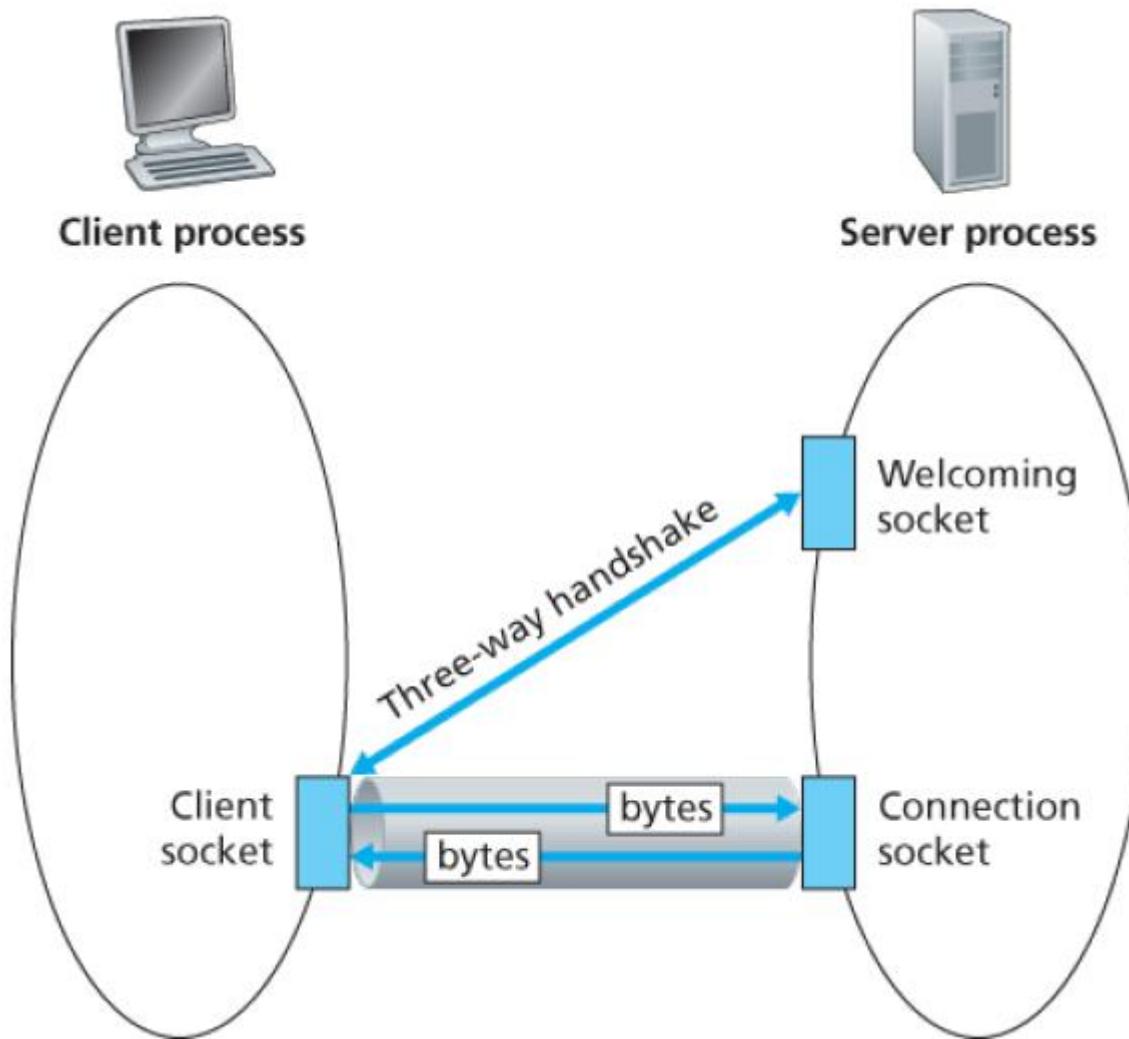
- ❖ Creating TCP socket, specifying IP address, port number of server process
- ❖ Client TCP establishes connection to server TCP

- ❖ when contacted by client, *server TCP creates new socket* for server process to communicate with that particular client
 - allows server to talk with multiple clients
 - source port numbers used to distinguish clients (more in Chap 3)
- TCP socket is identified by (destination IP address, destination port number, source IP address, source port number)

Application viewpoint:

TCP provides reliable, in-order byte-stream transfer ("pipe") between client and server

Socket programming with TCP



Client/server socket interaction:

~~TCP~~

server (running on `hostid`)

client

create socket,
port=`x`, for incoming request:

`serverSocket = socket()`

wait for incoming
connection request
`connectionSocket = serverSocket.accept()`

read request from
`connectionSocket`

write reply to
`connectionSocket`

close
`connectionSocket`

TCP
connection setup

create socket,
`clientSocket = socket()`
connect to `hostid`, port=`x`

send request using
`clientSocket`

read reply from
`clientSocket`

close
`clientSocket`

Example app: TCP server

Python TCPServer

create TCP welcoming
socket

server begins listening for
incoming TCP requests

loop forever

server waits on accept()
for incoming requests, new
socket created on return

read bytes from socket (but
not address as in UDP)

close connection to this client
(but *not* welcoming socket)

```
from socket import *
serverPort = 12000
serverSocket = socket(AF_INET,SOCK_STREAM)
serverSocket.bind(("",serverPort))
serverSocket.listen(1)
print 'The server is ready to receive'
while True:
    connectionSocket, addr = serverSocket.accept()
    sentence = connectionSocket.recv(1024).decode()
    capitalizedSentence = sentence.upper()
    connectionSocket.send(capitalizedSentence.
                           encode())
connectionSocket.close()
```

`connectionSocket` is identified by
- destination IP address and port number
- source IP address and port number

TCP socket



Example app: TCP client

Python TCPClient

```
from socket import *
serverName = 'servername'
serverPort = 12000
create TCP socket for server → clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName,serverPort))
sentence = raw_input('Input lowercase sentence:')
No need to attach server name, port → clientSocket.send(sentence.encode())
modifiedSentence = clientSocket.recv(1024)
print ('From Server:', modifiedSentence.decode())
clientSocket.close()
```

Chapter 2: summary

our study of network apps now complete!

- application architectures
 - client-server
 - P2P
- application service requirements:
 - reliability, bandwidth, delay
- Internet transport service model
 - connection-oriented, reliable: TCP
 - unreliable, datagrams: UDP
- specific protocols:
 - HTTP
 - SMTP, POP, IMAP
 - DNS
 - P2P: BitTorrent
- video streaming, CDNs
- socket programming:
TCP, UDP sockets

Chapter 2: summary

most importantly: learned about protocols!

- ❖ typical request/reply message exchange:
 - client requests info or service
 - server responds with data, status code
- ❖ message formats:
 - *headers*: fields giving info about data
 - *data*: info being communicated

important themes:

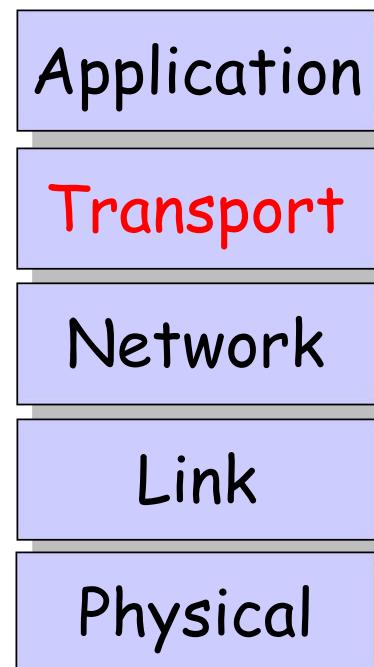
- centralized vs. distributed
- stateless vs. stateful
- reliable vs. unreliable message transfer

Chapter 3: Transport Layer

Our goals:

- ❖ understand principles behind transport layer services:
 - multiplexing, demultiplexing
 - reliable data transfer
 - flow control
 - congestion control

- ❖ learn about Internet transport layer protocols:
 - UDP: connectionless transport
 - TCP: connection-oriented reliable transport
 - TCP congestion control



Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

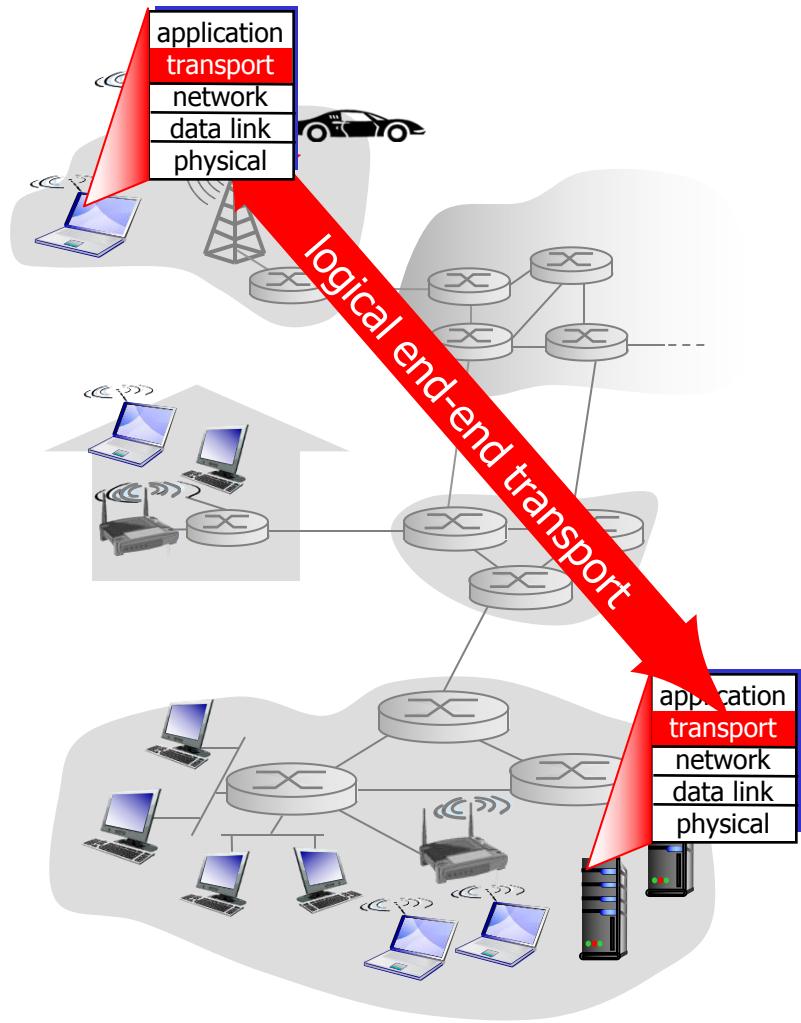
- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

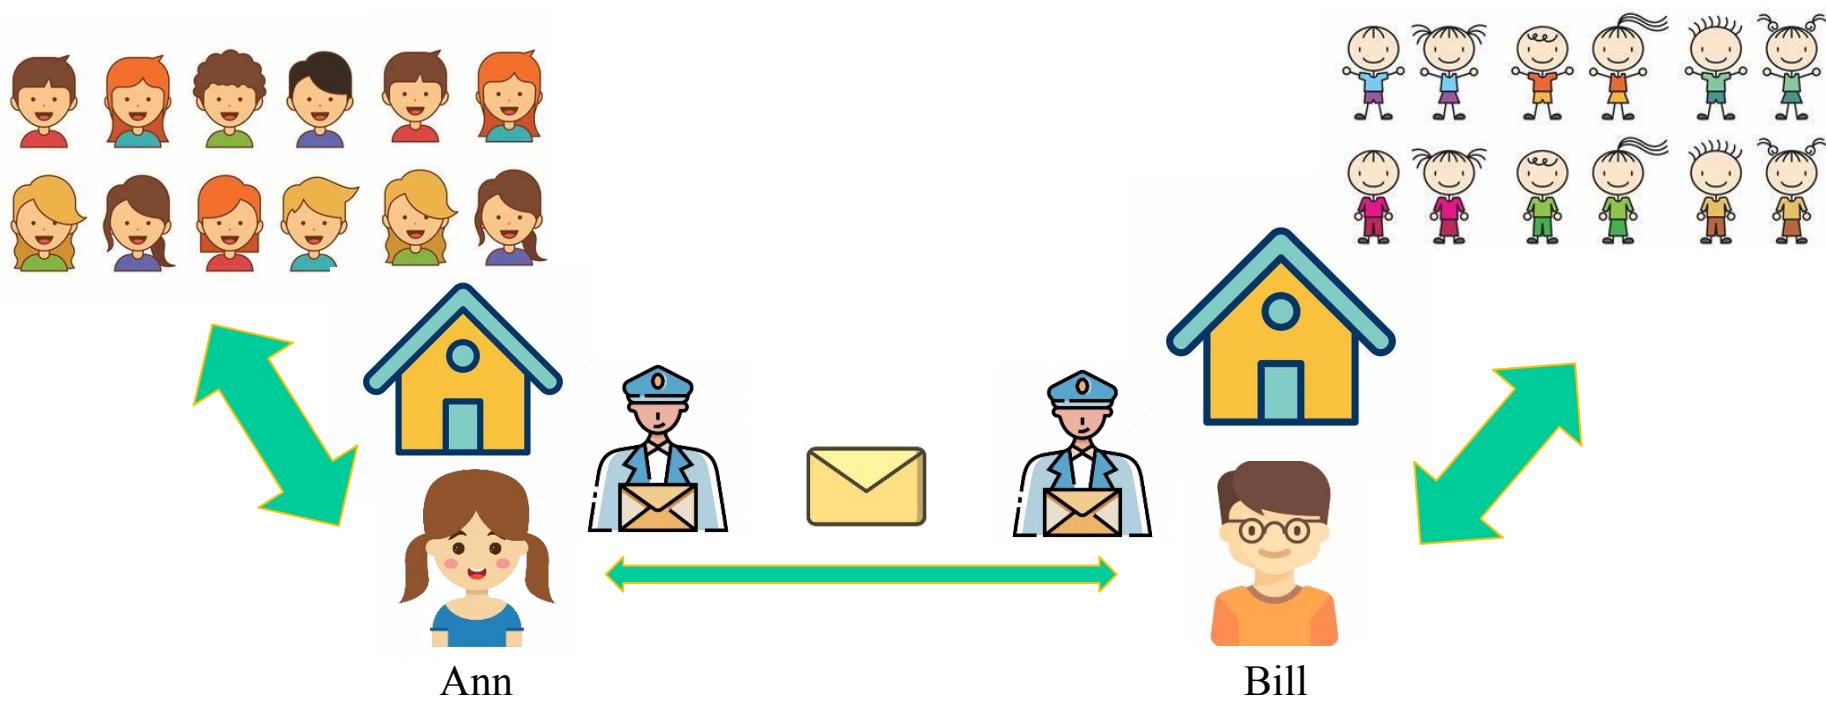
Transport services and protocols

- ❖ Provide *logical communication* between app processes running on different hosts
- ❖ Transport protocols run in end systems
 - send side: breaks app messages into *segments*, passes to network layer
 - rcv side: reassembles segments into messages, passes to app layer
- ❖ more than one transport protocol available to apps
 - Internet: TCP and UDP



Transport vs. network layer

- ❖ Network layer: logical communication between **hosts**
- ❖ Transport layer: logical communication between **processes**
 - relies on, enhances, network layer services



Transport vs. network layer



household analogy:

- ❖ hosts = houses
- ❖ processes = kids
- ❖ app messages = letters in envelopes
- ❖ transport protocol = Ann and Bill
- ❖ network-layer protocol = postal service

Susan and Harvey substitute for them and provide different delivery services?

Transport vs. network layer

- ❖ The services that a transport protocol can provide are often **constrained by** the service model of the underlying network-layer protocol.
 - delay or bandwidth guarantees
- ❖ Certain services can be offered by a transport protocol **even when** the underlying network protocol **doesn't offer** the corresponding service at the network layer.
 - Reliable data transfer; security

Internet transport-layer protocols

Network layer: Internet protocol (IP) is a best effort delivery service, unreliable

UDP: unreliable, unordered delivery:

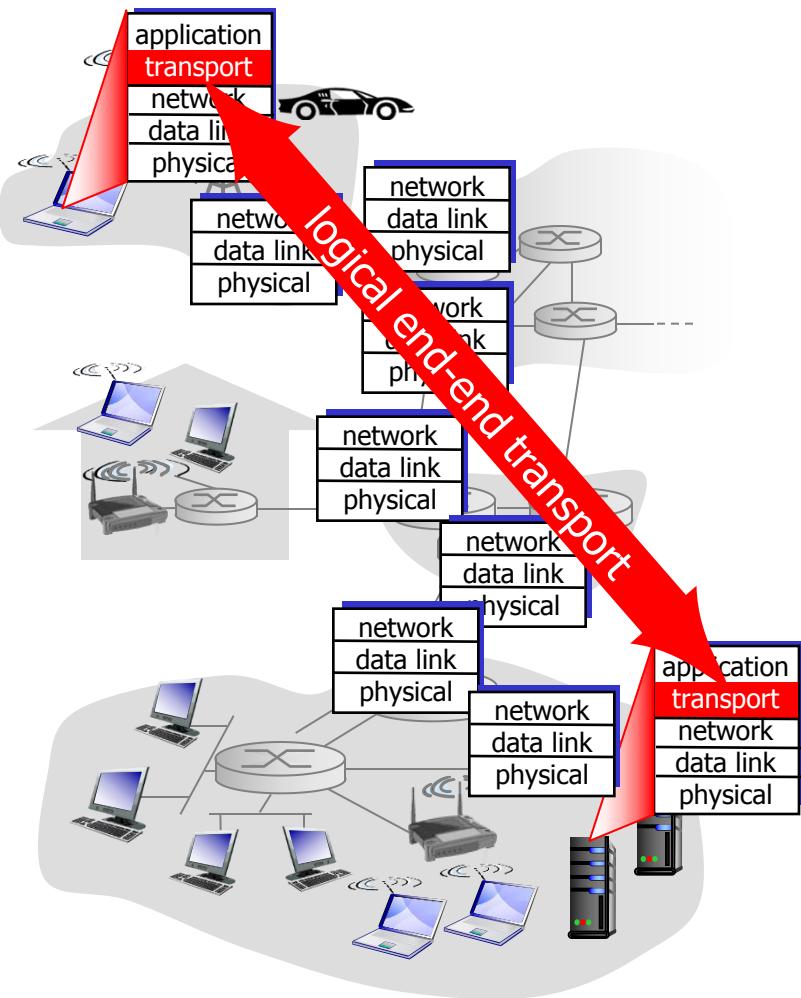
- ❖ no-frills extension
- ❖ process-to-process data delivery and error checking

TCP : reliable, in-order delivery

- ❖ congestion control
- ❖ flow control
- ❖ connection setup

Services not available:

- ❖ delay guarantees
- ❖ bandwidth guarantees



Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

Multiplexing/demultiplexing

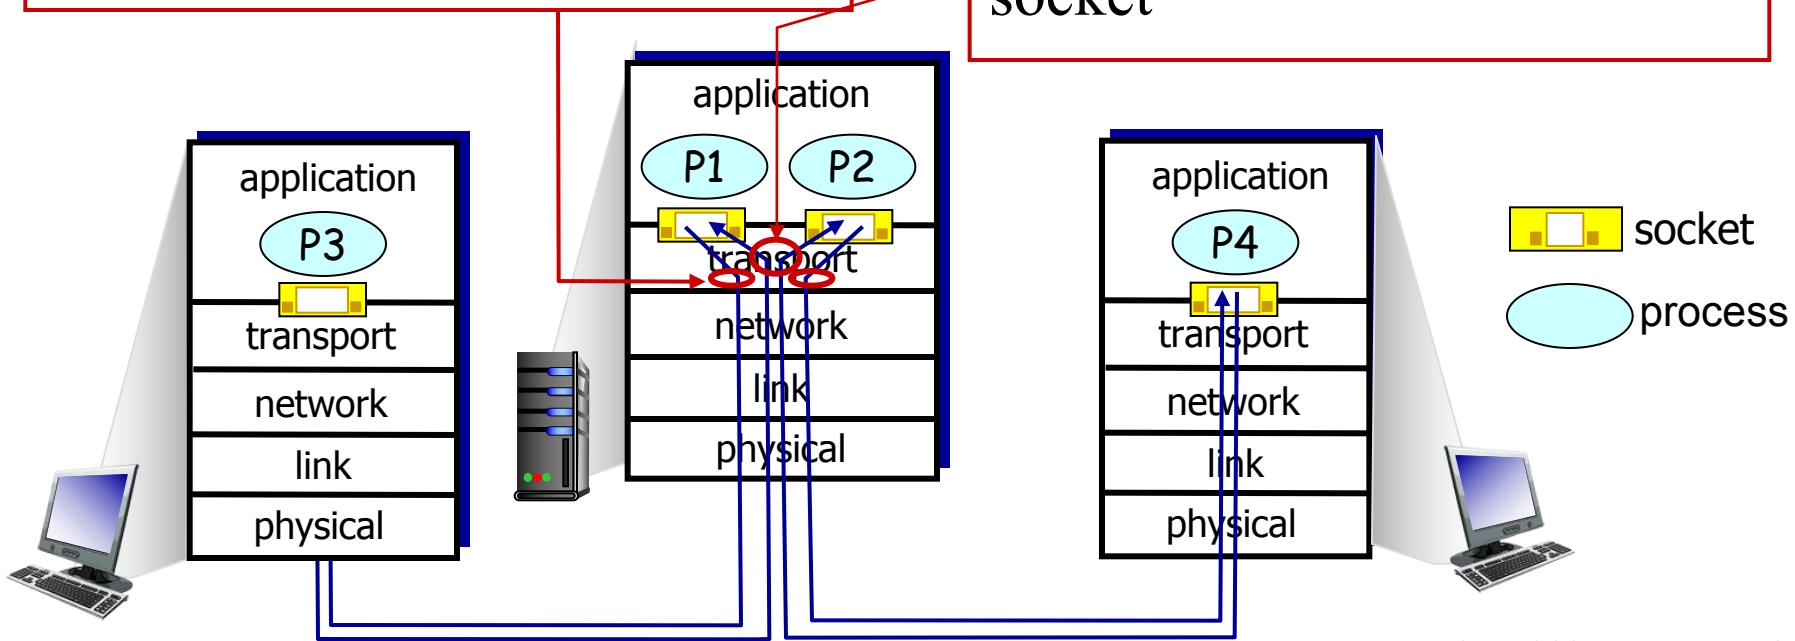
Multiplexing and demultiplexing: extending the host-to-host delivery service to a process-to-process delivery service for applications running on the hosts.

Multiplexing at sender:

handle data from multiple sockets, add transport header (later used for demultiplexing)

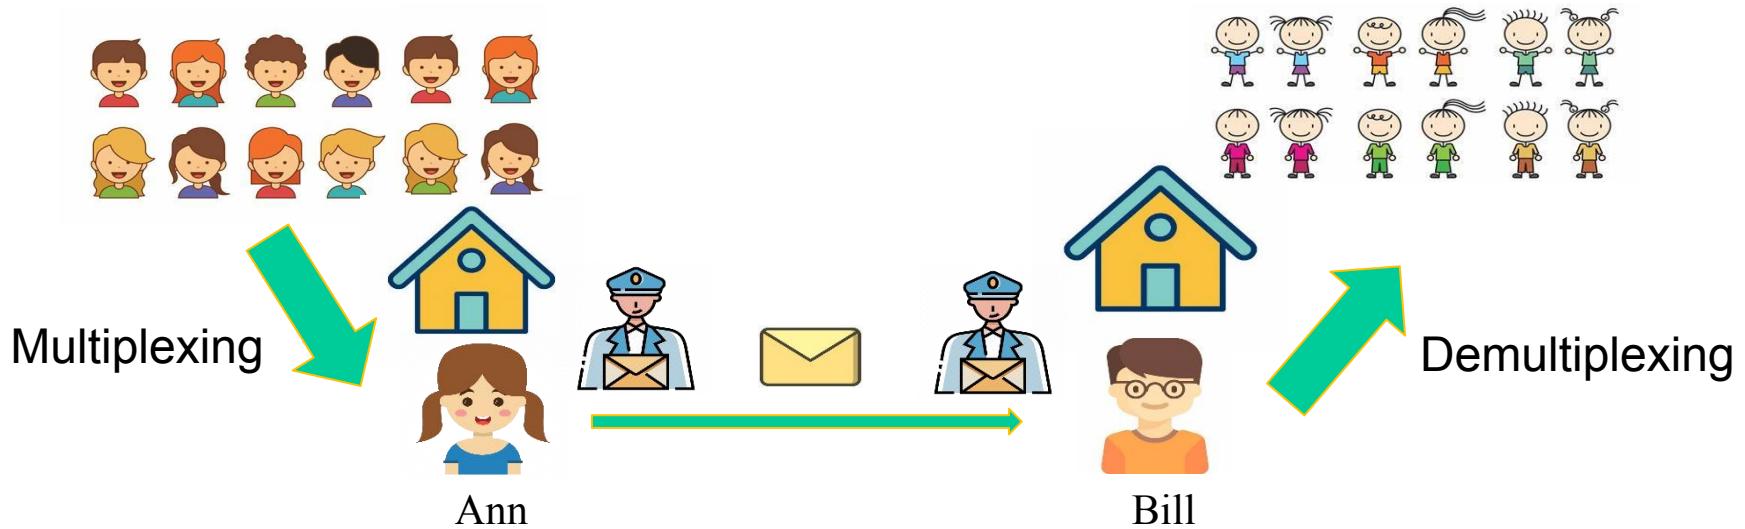
Demultiplexing at receiver:

use header info to deliver received segments to correct socket



Ann and Bill example?

Multiplexing/demultiplexing



How demultiplexing works?

- Each child must have an identifier (e.g., name, ID)

How demultiplexing works

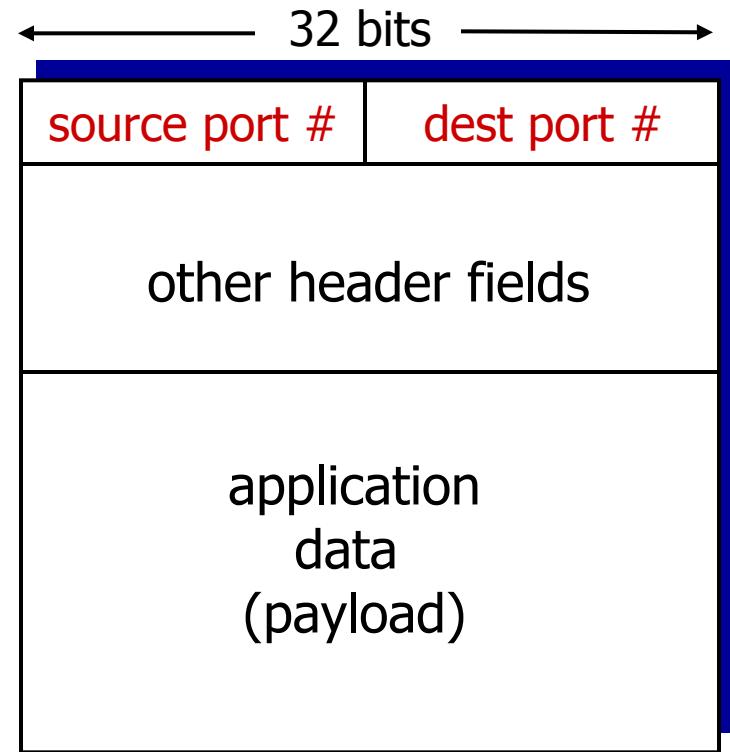
Multiplexing/demultiplexing:

- ❖ sockets have unique identifiers
- ❖ each segment have special fields that indicate the socket to which the segment is to be delivered.

Sending host: Host uses *IP addresses & port numbers* to direct segment to appropriate socket

Receiving host: Host receives IP datagrams from network layer

- ❖ each datagram has source IP address, destination IP address
- ❖ each datagram carries one transport-layer segment



TCP/UDP segment format

Port number: a 16-bit number, ranging from 0 to 65535:

- well-known port numbers: 0-1023

Overview

- ❖ UDP: Connectionless demux
- ❖ TCP: Connection-oriented demux

UDP: Connectionless demux

- ❖ *recall:* created socket has local port #:

```
clientSocket = socket(AF_INET, SOCK_DGRAM)
```

Automatically assigns a
port number

```
clientSocket.bind(('', 19157))
```

Optional; at the server side, usually assign
port number

- ❖ *recall:* when creating datagram to send into UDP socket, must specify destination IP address and destination port #

```
clientSocket.sendto(message.encode(), (serverName, serverPort))
```

UDP socket is fully identified by a **two-tuple** consisting of a destination IP address and a destination port number.

when host receives UDP segment:

- ❖ checks destination port # in segment
- ❖ directs UDP segment to socket with that port #



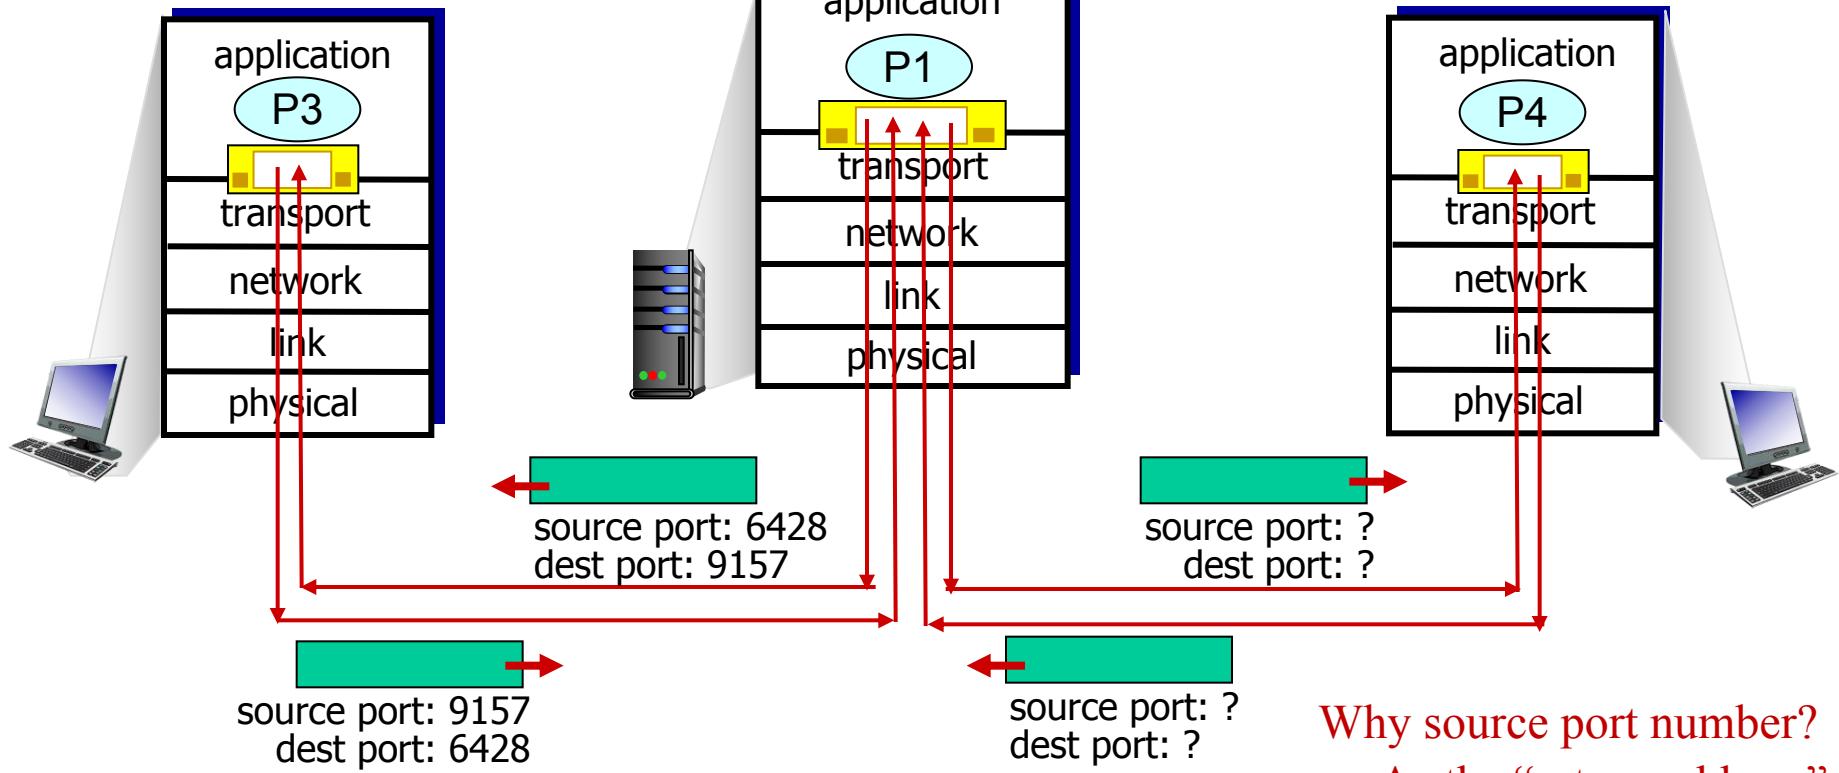
IP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at destination

Connectionless demux: example

```
mysocket2 =  
    socket(AF_INET,  
           SOCK_DGRAM)  
  
mysocket2.bind  
    (('', 9157))
```

```
serversocket =  
    socket(AF_INET,  
           SOCK_DGRAM)  
  
serversocket.bind  
    (('', 6428))
```

```
mysocket1 =  
    socket(AF_INET,  
           SOCK_DGRAM)  
  
mysocket1.bind  
    (('', 5775))
```



Overview

- ❖ UDP: Connectionless demux
- ❖ TCP: Connection-oriented demux

TCP: Connection-oriented demux

- ❖ Server **creates a welcome socket** with port no.12000
`serversocket = socket(AF_INET, SOCK_STREAM)`
`serversocket.bind(‘,12000))`
- ❖ Client **connects to the server**, the request is a TCP segment with a flag bit = 1
`clientsocket = socket(AF_INET, SOCK_STREAM)`
`clientsocket.connect((ServerName,12000))`
- ❖ Server **creates a new socket** to accept the connection
`connectionsocket, addr = serversocket.accept()`

The server may maintain TCP connections with multiple clients, each has a different **connectionsocket**. How to demux?

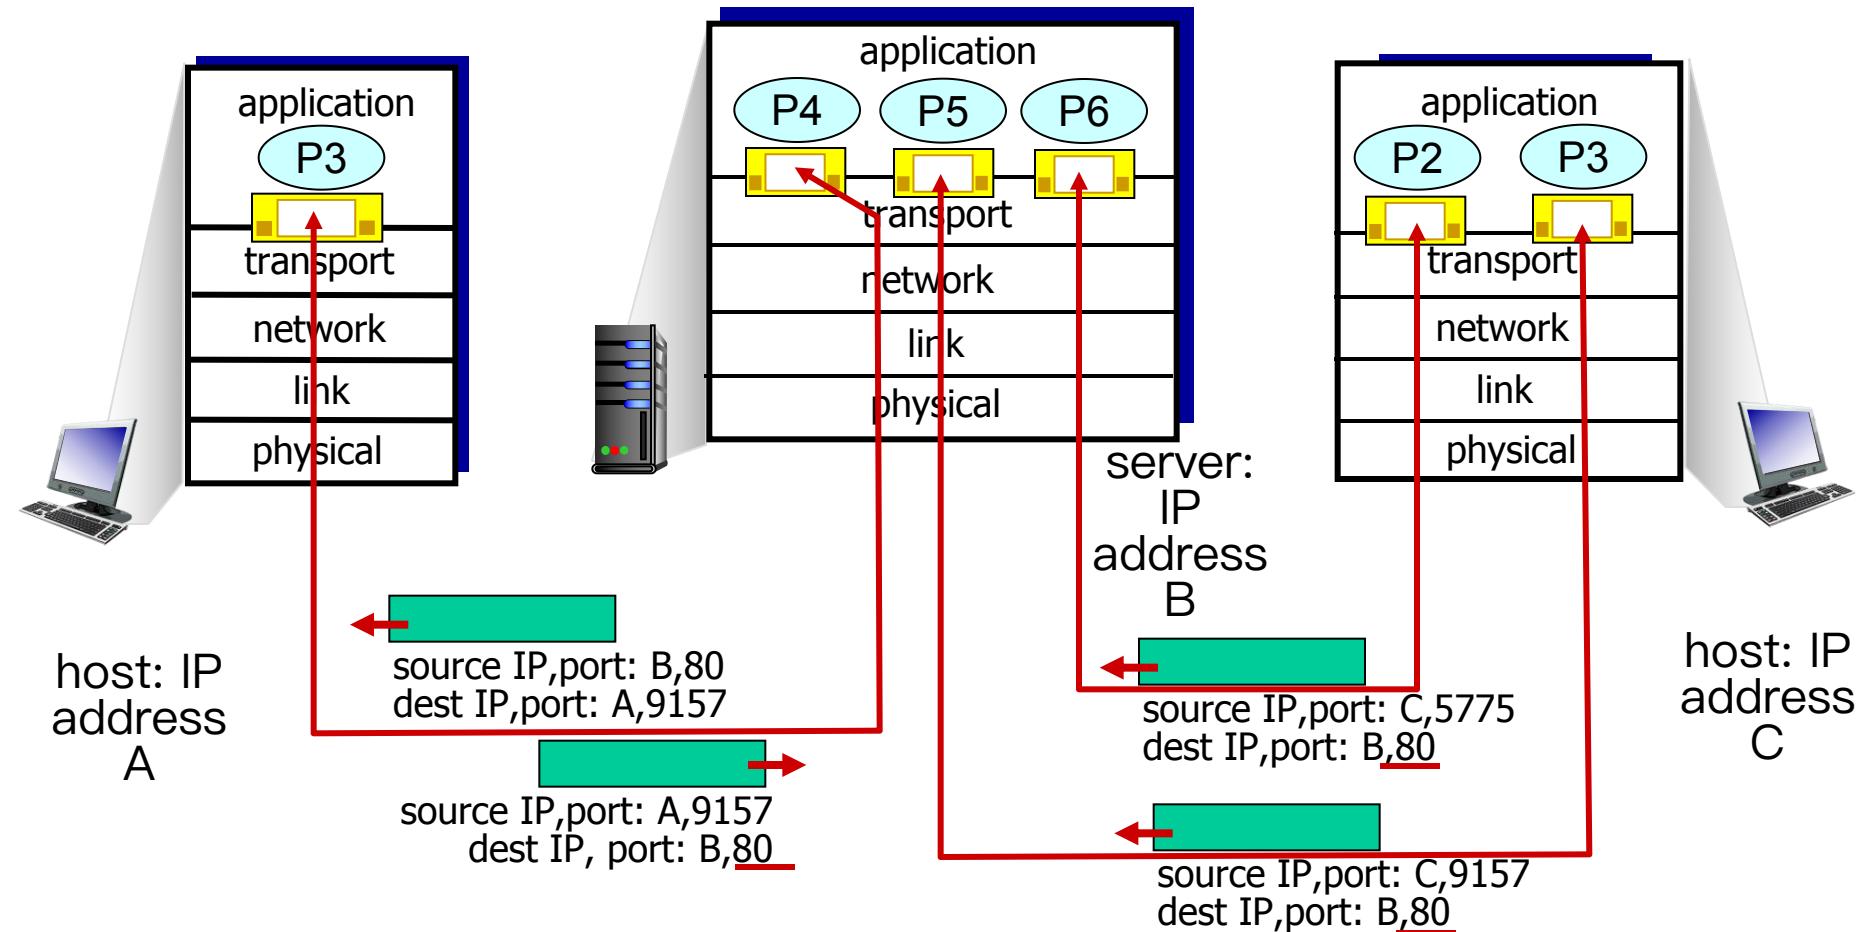
Connection-oriented demux

- ❖ TCP socket identified by 4-tuple:
 - source IP address, source port number, dest IP address, dest port number
- ❖ Demux: receiver uses all four values to direct segment to appropriate socket
- ❖ Server host may support many simultaneous TCP sockets:
 - each socket identified by its own 4-tuple

Web servers have different sockets for each connecting client

- Both the initial connection-establishment segments and the segments carrying HTTP requests will have destination port 80.
- non-persistent HTTP will have different socket for each request

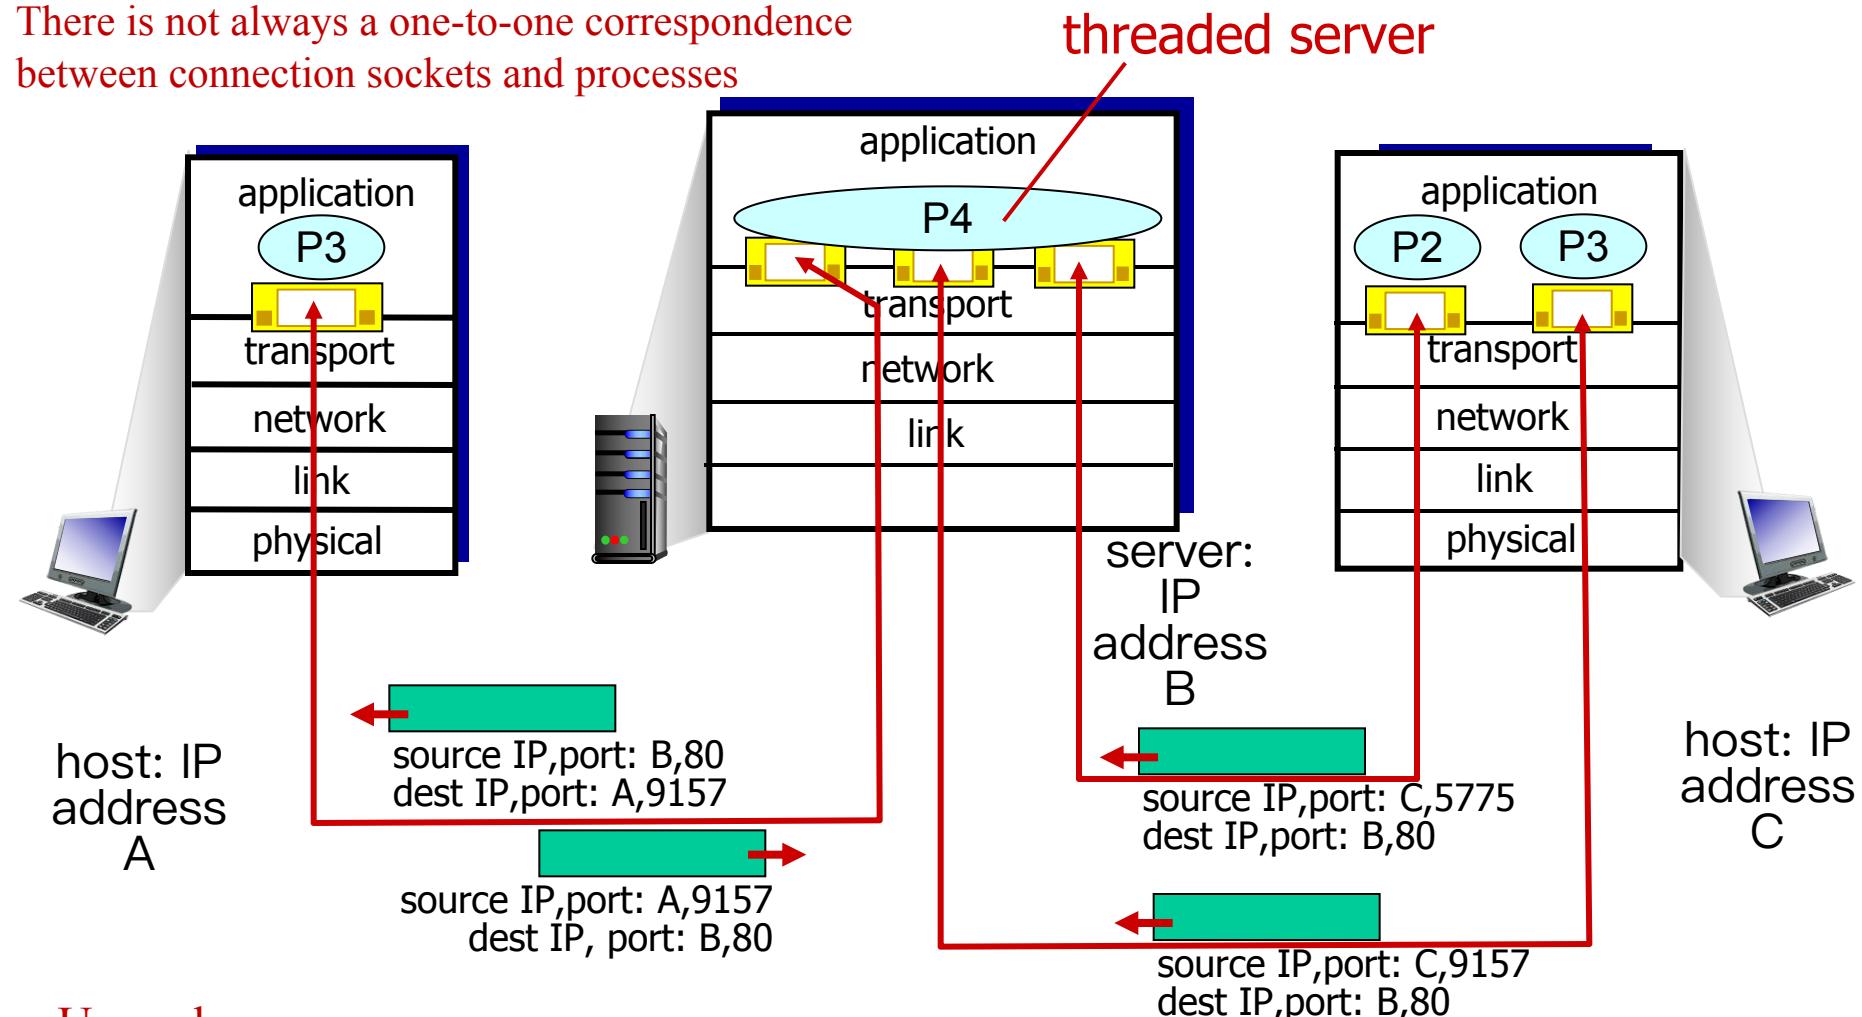
Connection-oriented demux: example



three segments, all destined to IP address: B, dest port: 80
are demultiplexed to *different* sockets

Connection-oriented demux: example

There is not always a one-to-one correspondence
between connection sockets and processes



Use only one process

- ❖ Create a new thread with a new connection socket for each new client connection.
- ❖ A thread can be viewed as a lightweight subprocess.

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

UDP: User Datagram Protocol

- ❖ “No frills,” “bare bones” Internet transport protocol
 - Multiplexing/demultiplexing; light error checking
- ❖ “Best effort” service, UDP segments may be:
 - Lost, delivered out-of-order to app
- ❖ *connectionless:*
 - No handshaking between UDP sender, receiver
 - Each UDP segment handled independently of others
 - No congestion control

Advantage?

- No congestion control: Immediately pass the segment to network layer
- No connection-establish delay
- No connection state: server can support more clients
- Smaller packet overhead

Disadvantage?

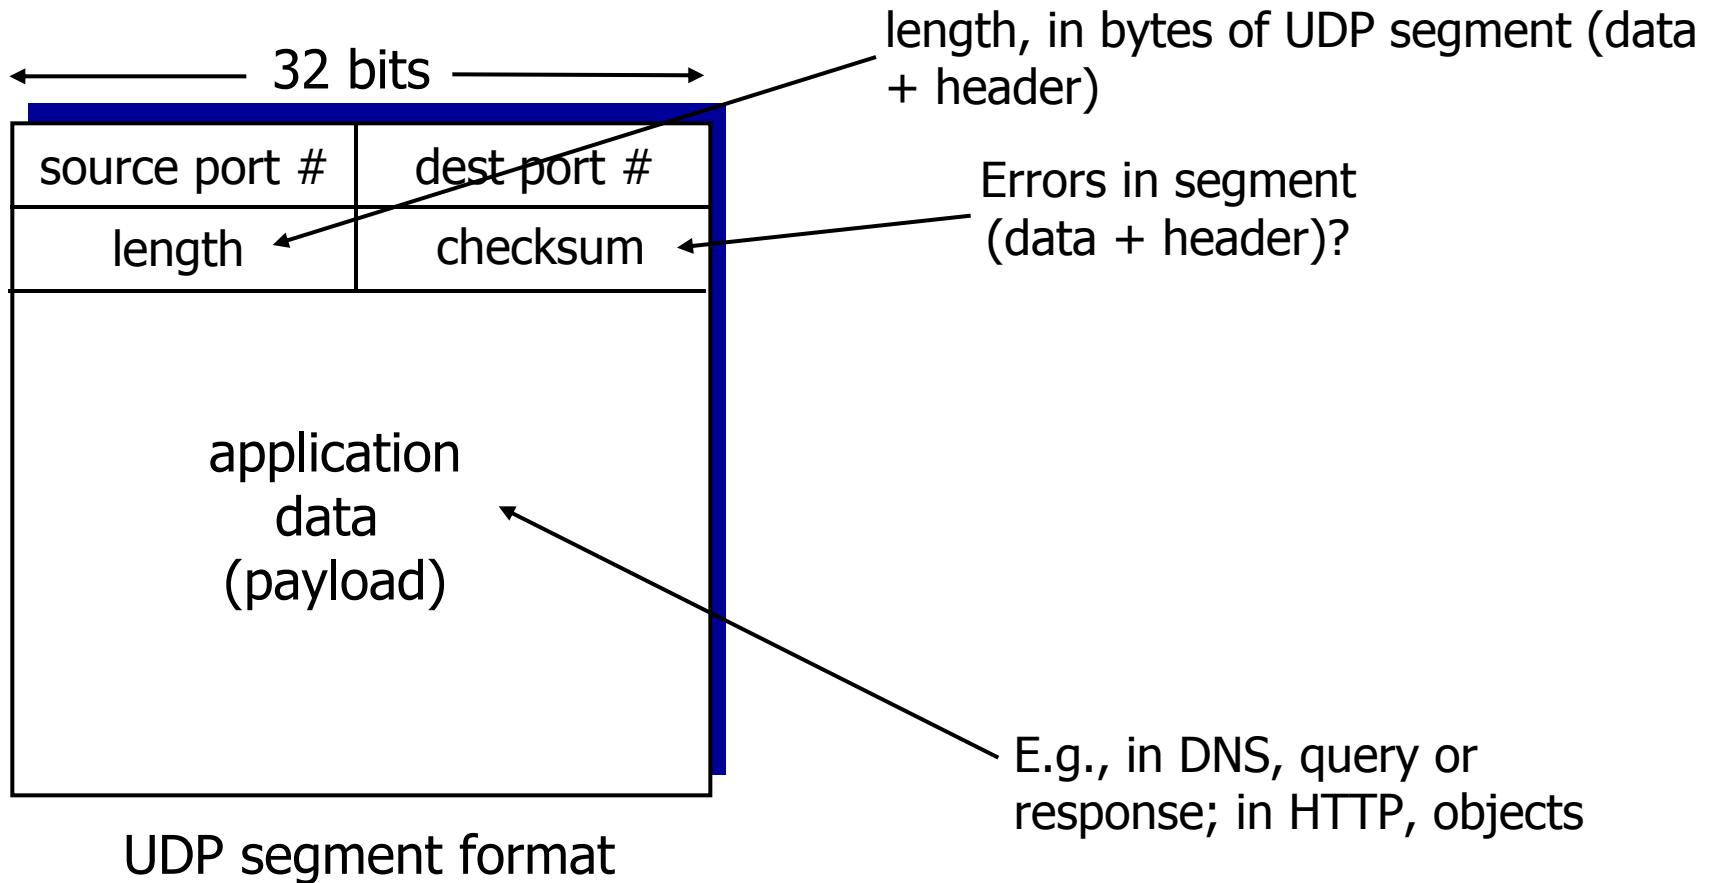
- No congestion control: congestion, overflow, fairness
- Not reliable

UDP: User Datagram Protocol

- ❖ UDP is used in:
 - streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS

- ❖ reliable transfer over UDP:
 - add reliability at application layer
 - application-specific error recovery!

UDP: segment header



UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment (from source to destination)

Sender:

- ❖ treat segment contents, including header fields, as sequence of 16-bit integers
- ❖ checksum: **1s complement of the sum** of segment contents
- ❖ sender puts checksum value into UDP checksum field

Receiver:

- ❖ check the sum of the segment
 - All bits are equal to 1 - no error detected. *But maybe errors nonetheless? More later*
 - Otherwise: error detected

Internet checksum: example

At the sender side, determine the check sum of the following three 16-bit words

0110011001100000

0101010101010101

1000111100001100

The sum of first two of these 16-bit words is

0110011001100000

0101010101010101

1011101110110101

Adding the third word to the above sum gives

1011101110110101

1000111100001100

1010010101100001

0100101011000010

Note: when adding numbers, a carryout from the most significant bit needs to be added to the result

The 1s complement is obtained by converting all the 0s to 1s and converting all the 1s to 0s.

Checksum: 1011010100111101

Receiver Side?

Internet checksum: example

At the receiver side, all four 16-bit words are added, including the checksum:

- ❖ If no errors are introduced into the packet, then clearly the sum at the receiver will be **1111111111111111**.
- ❖ If one of the bits is a 0, then we know that errors have been introduced into the packet.

Check the sum of the segment

- All bits are equal to 1 - **no error detected**. *But maybe errors nonetheless? More later*
- Otherwise: error detected

Internet checksum: example

Why UDP provides a checksum?

- no guarantee that all the links provide error checking
- bit errors could be introduced when segments are in memory
- “functions placed at the lower levels may be redundant or of little value when compared to the cost of providing them at the higher level.”

Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

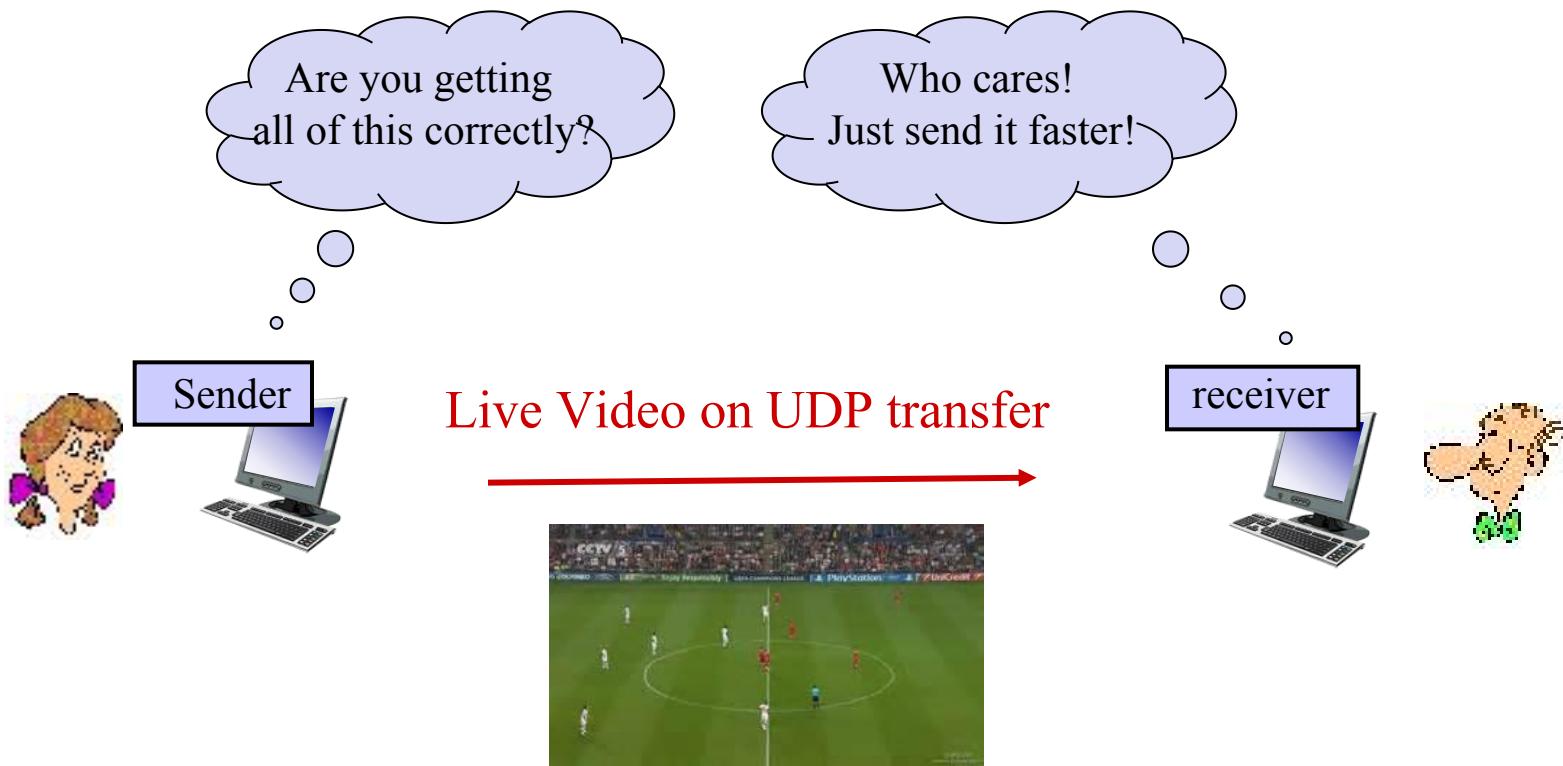
- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

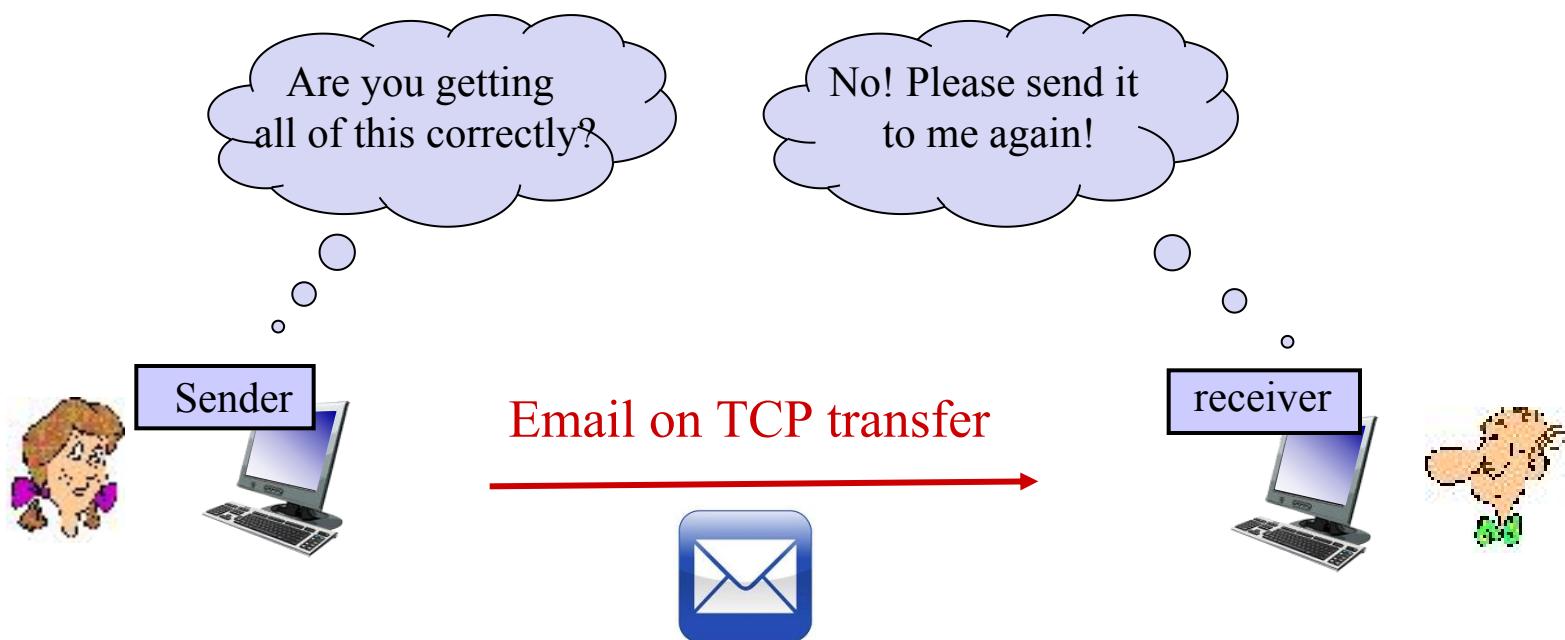
UDP Transfer

- ❖ UDP cannot guarantee reliable data transfer
- ❖ But, it's faster!



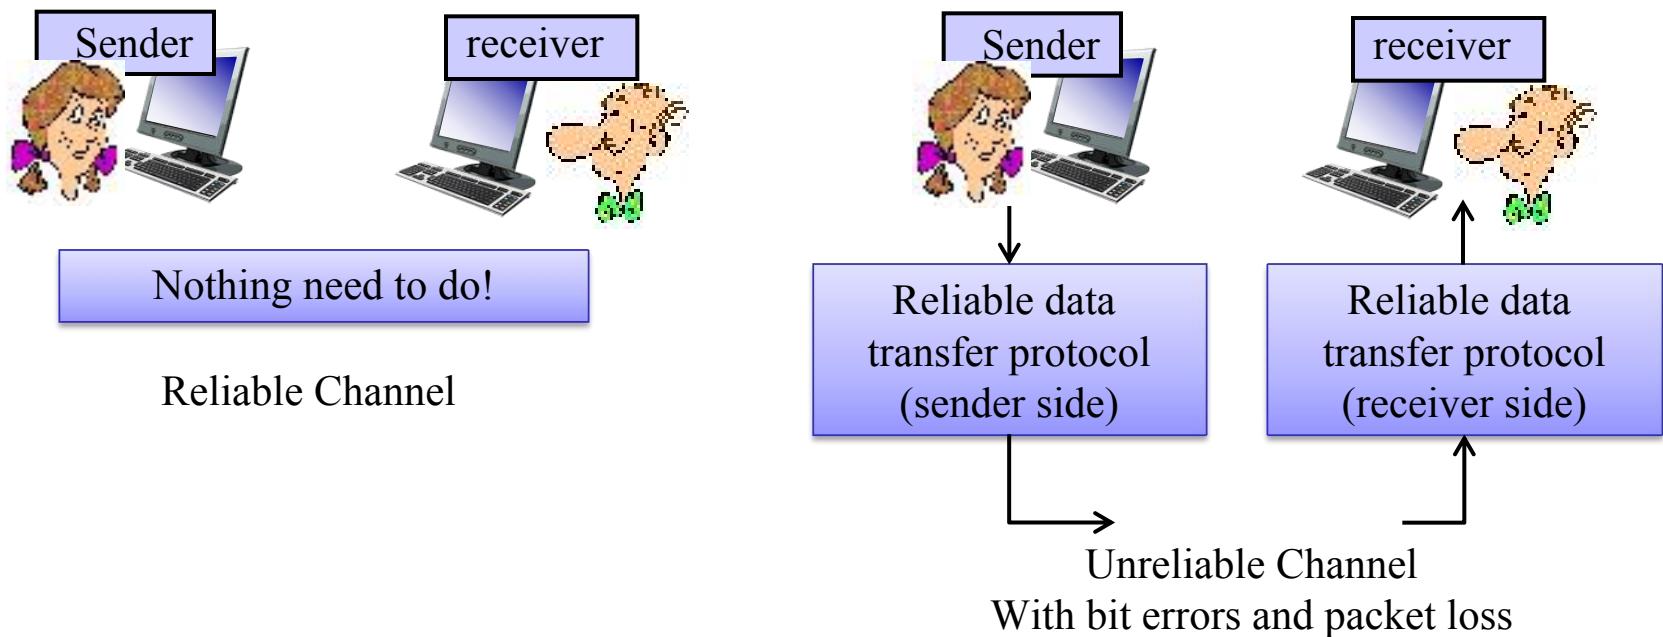
TCP Transfer

- ❖ TCP can guarantee reliable data transfer
- ❖ But, it's slower and more complex!

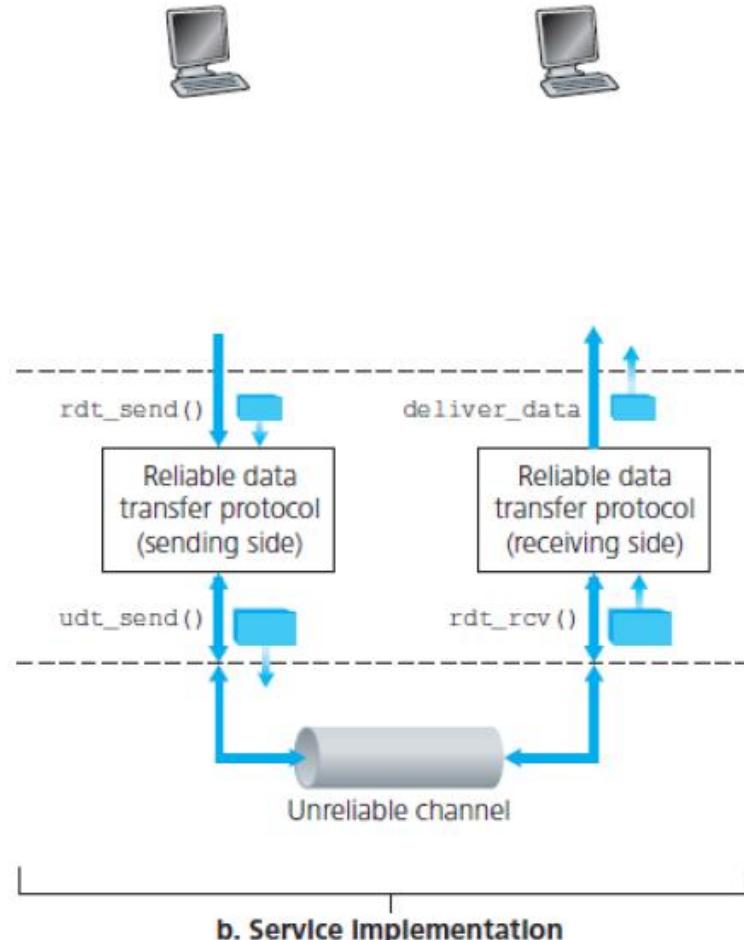
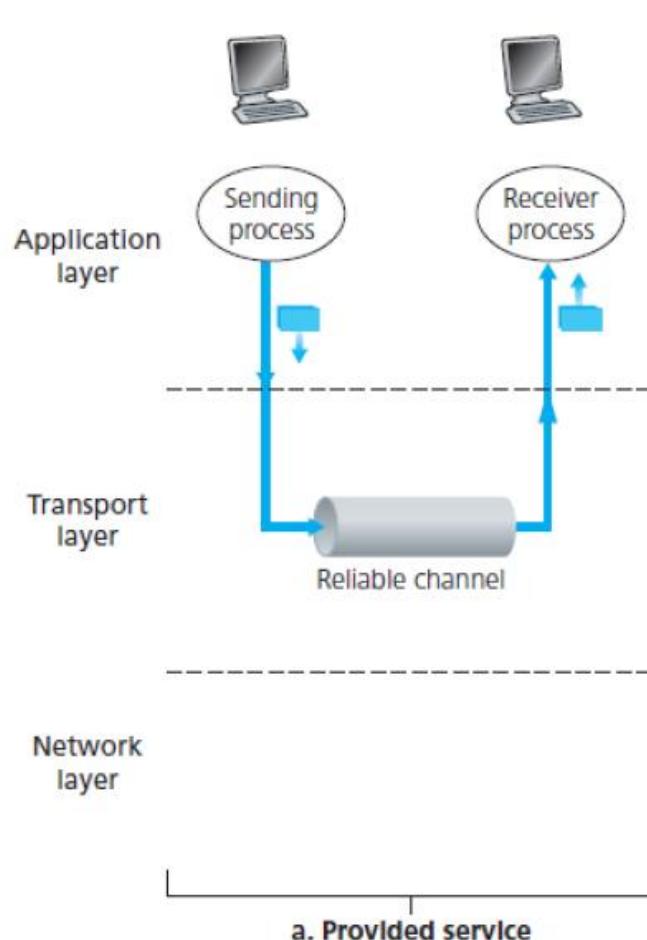


Reliable Data Transfer (rdt)

- ❖ In top-10 list of important networking topics!
- ❖ Reliable data transfer over **unreliable channel**:
 - Bit flip, lost, out-of-order
 - In this section, assume unreliable channel not reorder packets



Reliable Data Transfer (rdt)

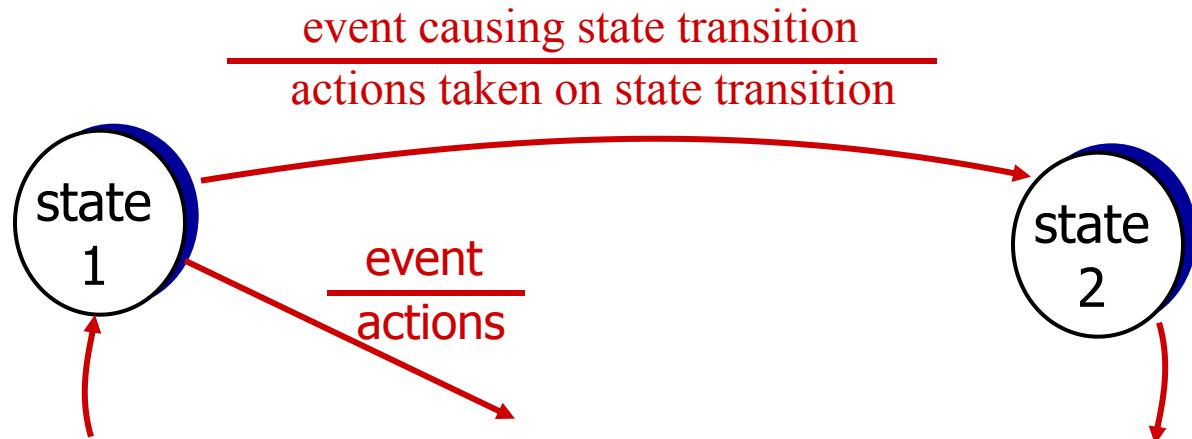


Reliable data transfer: getting started

We'll:

- ❖ Incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ Consider only unidirectional data transfer
 - but control info will flow on both directions!
- ❖ Use finite state machines (FSM) to specify sender, receiver

state: when in this “state”
next state uniquely
determined by next event



Overview

Roadmap:

- ❖ Perfectly reliable channel: rdt1.0
- ❖ Channel with bit error:
 - bit error in packet: rdt 2.0
 - bit error in ACK: 2.1
 - NAK-free: 2.2
- ❖ Lossy channel: rdt 3.0

Summary of Techniques

- Checksum
- Sequence number
- ACK packets
- Retransmission
- Timeout

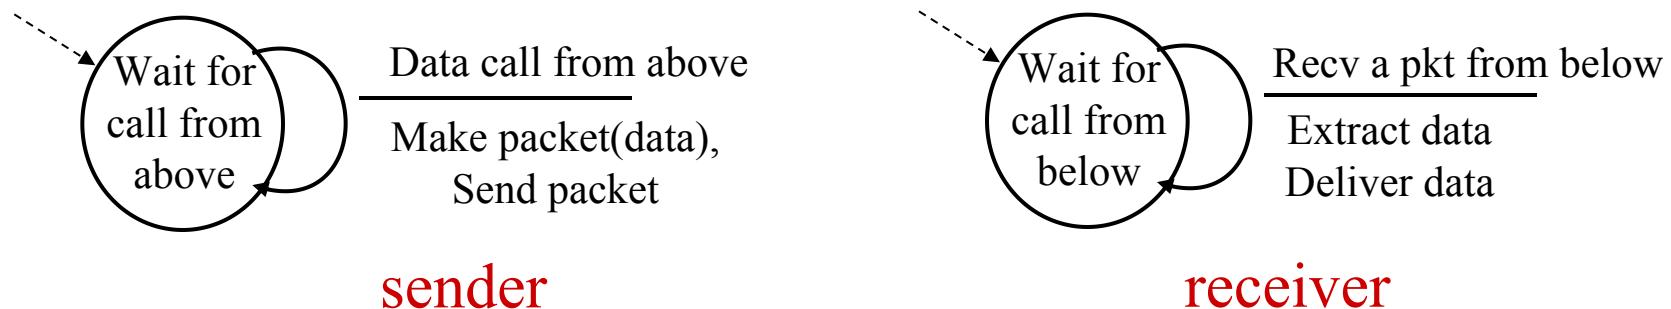
rdt1.0: reliable transfer over a reliable channel

- ❖ Underlying channel perfectly reliable

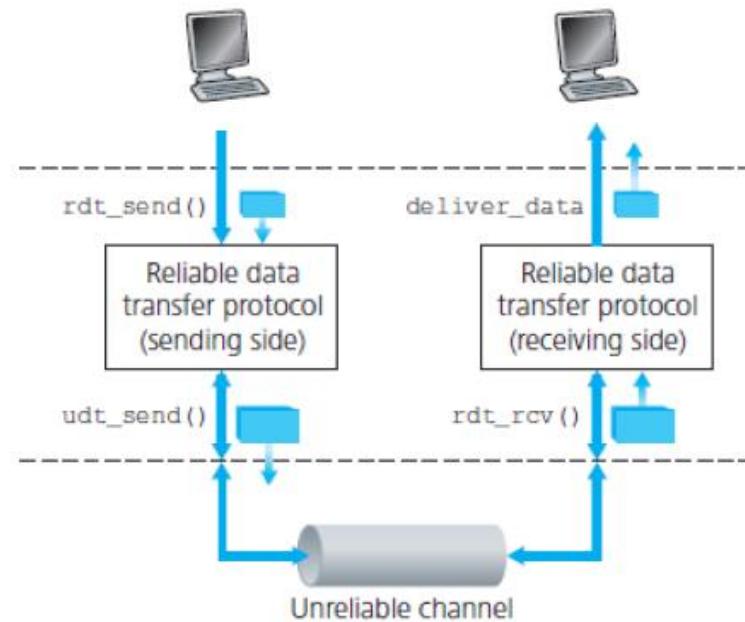
- no bit errors
- no loss of packets

- ❖ Rdt 1.0:

- sender sends data into underlying channel
- receiver reads data from underlying channel
- Reliable channel, no need for feedback (no control message)



rdt1.0: reliable transfer over a reliable channel



Wait for
call from
above

`rdt_send(data)`
`packet=make_pkt(data)`
`udt_send(packet)`

sender

Wait for
call from
below

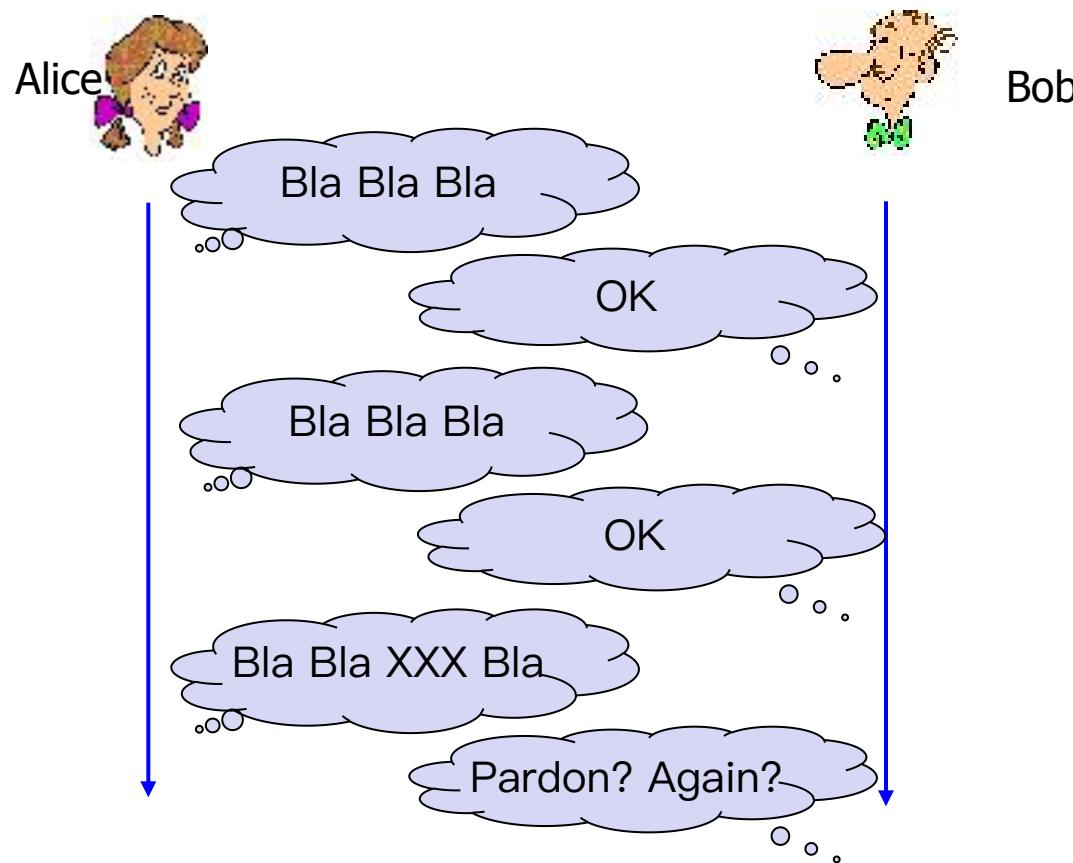
`rdt_rcv(packet)`
`extract(packet,data)`
`deliver_data(data)`

receiver

rdt2.0: channel with bit errors

- ❖ Underlying channel may flip bits ($0 \rightarrow 1$) in packet

How do humans recover from “errors” during conversation?



rdt2.0: channel with bit errors

- ❖ Underlying channel may **flip bits** ($0 \rightarrow 1$) in packet

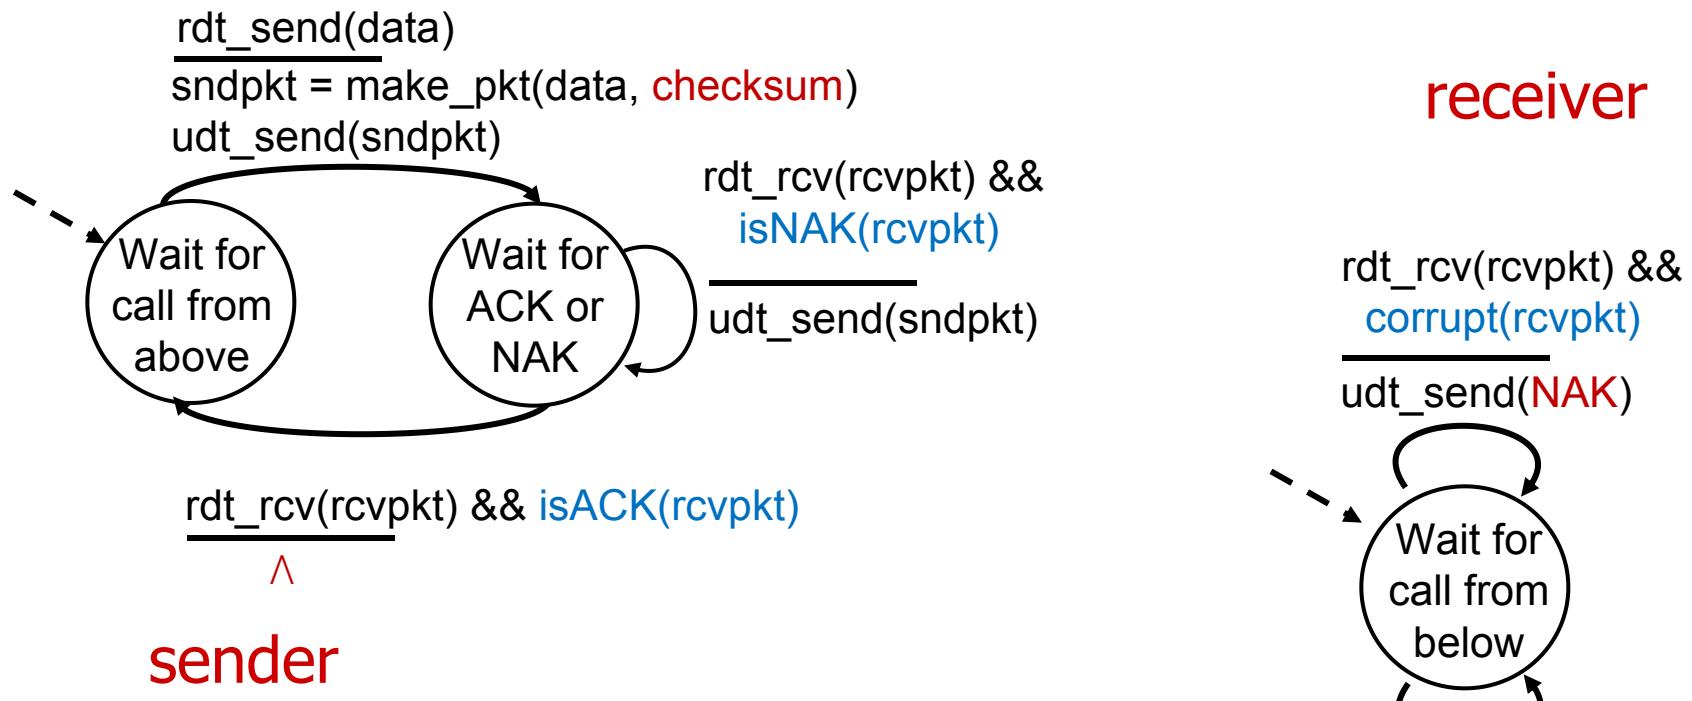
How do humans recover from “errors” during conversation?

- ❖ The question: how to recover from errors?
 - *acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
 - *negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
 - sender retransmits pkt on receipt of NAK
- ❖ new mechanisms in **rdt2.0** (beyond **rdt1.0**):
 - error detection
 - receiver feedback: control msgs (ACK,NAK) rcvr->sender
 - retransmission

rdt2.0: channel with bit errors

- ❖ Key mechanisms:
 - error detection
 - **feedback**: control msgs (ACK, NAK) from receiver to sender
 - retransmission
- ❖ Error detection: checksum
- ❖ Feedback messages:
 - *acknowledgements (ACKs)*: receiver explicitly tells sender that pkt received OK
 - *negative acknowledgements (NAKs)*: receiver explicitly tells sender that pkt had errors
 - sender retransmits pkt on receipt of NAK

rdt2.0: FSM specification

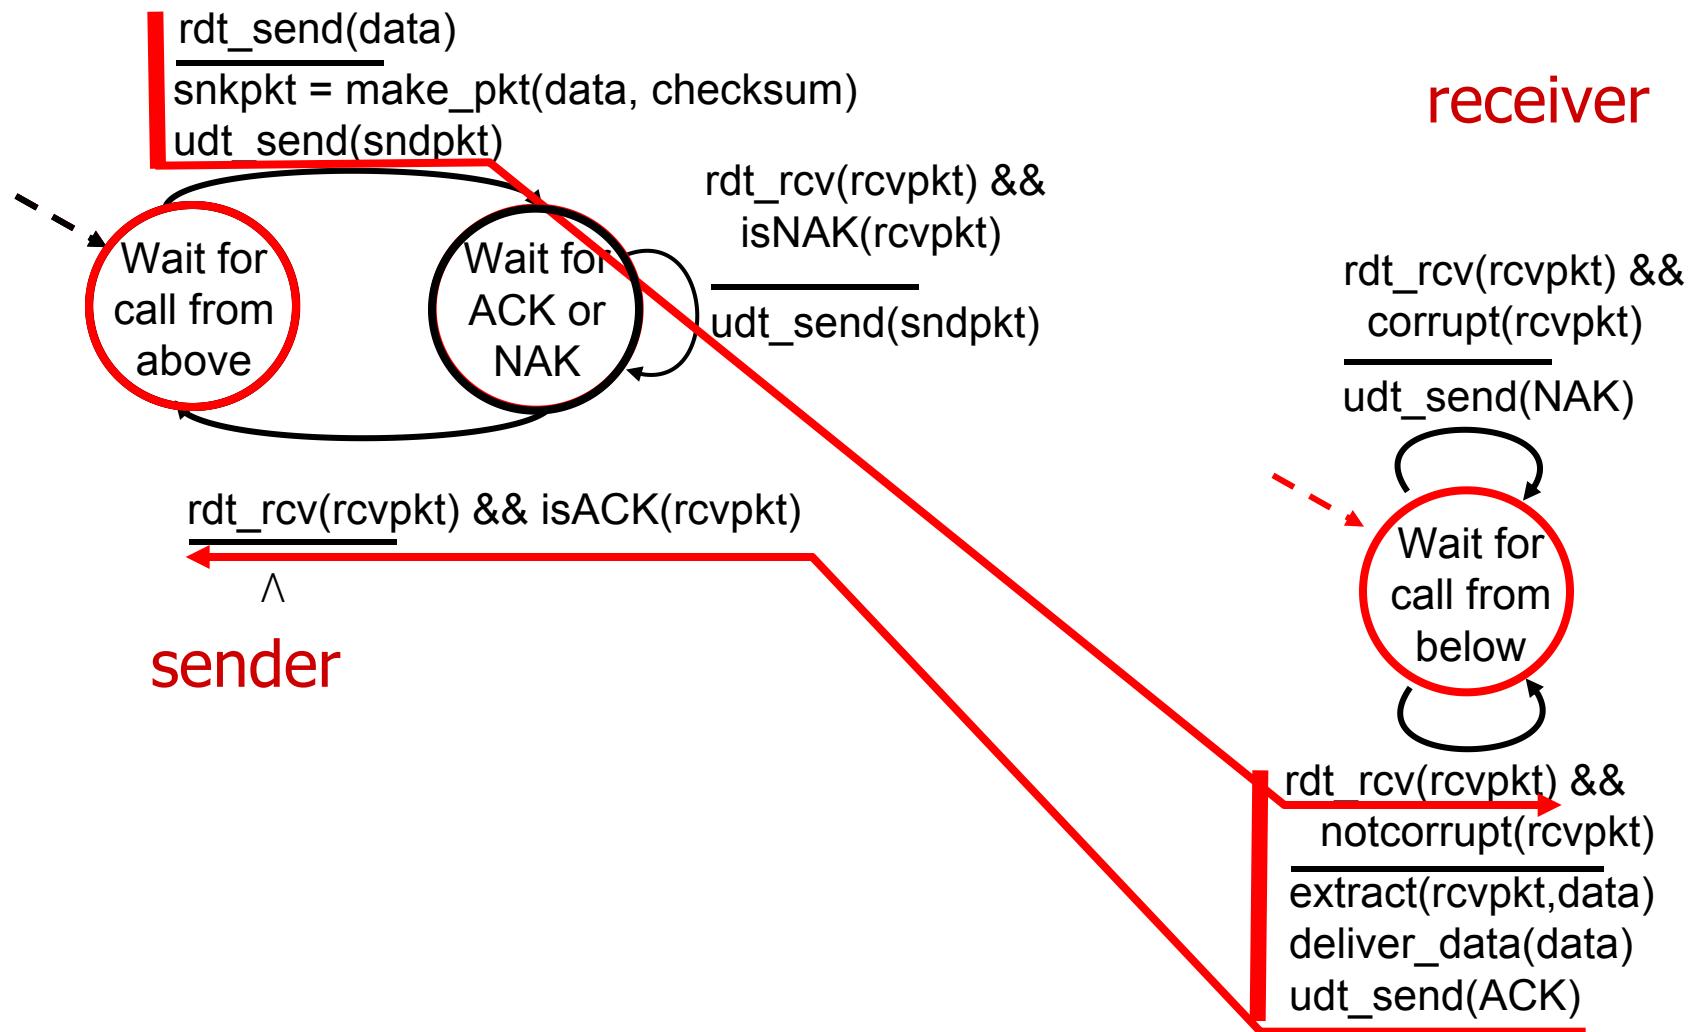


Stop and wait

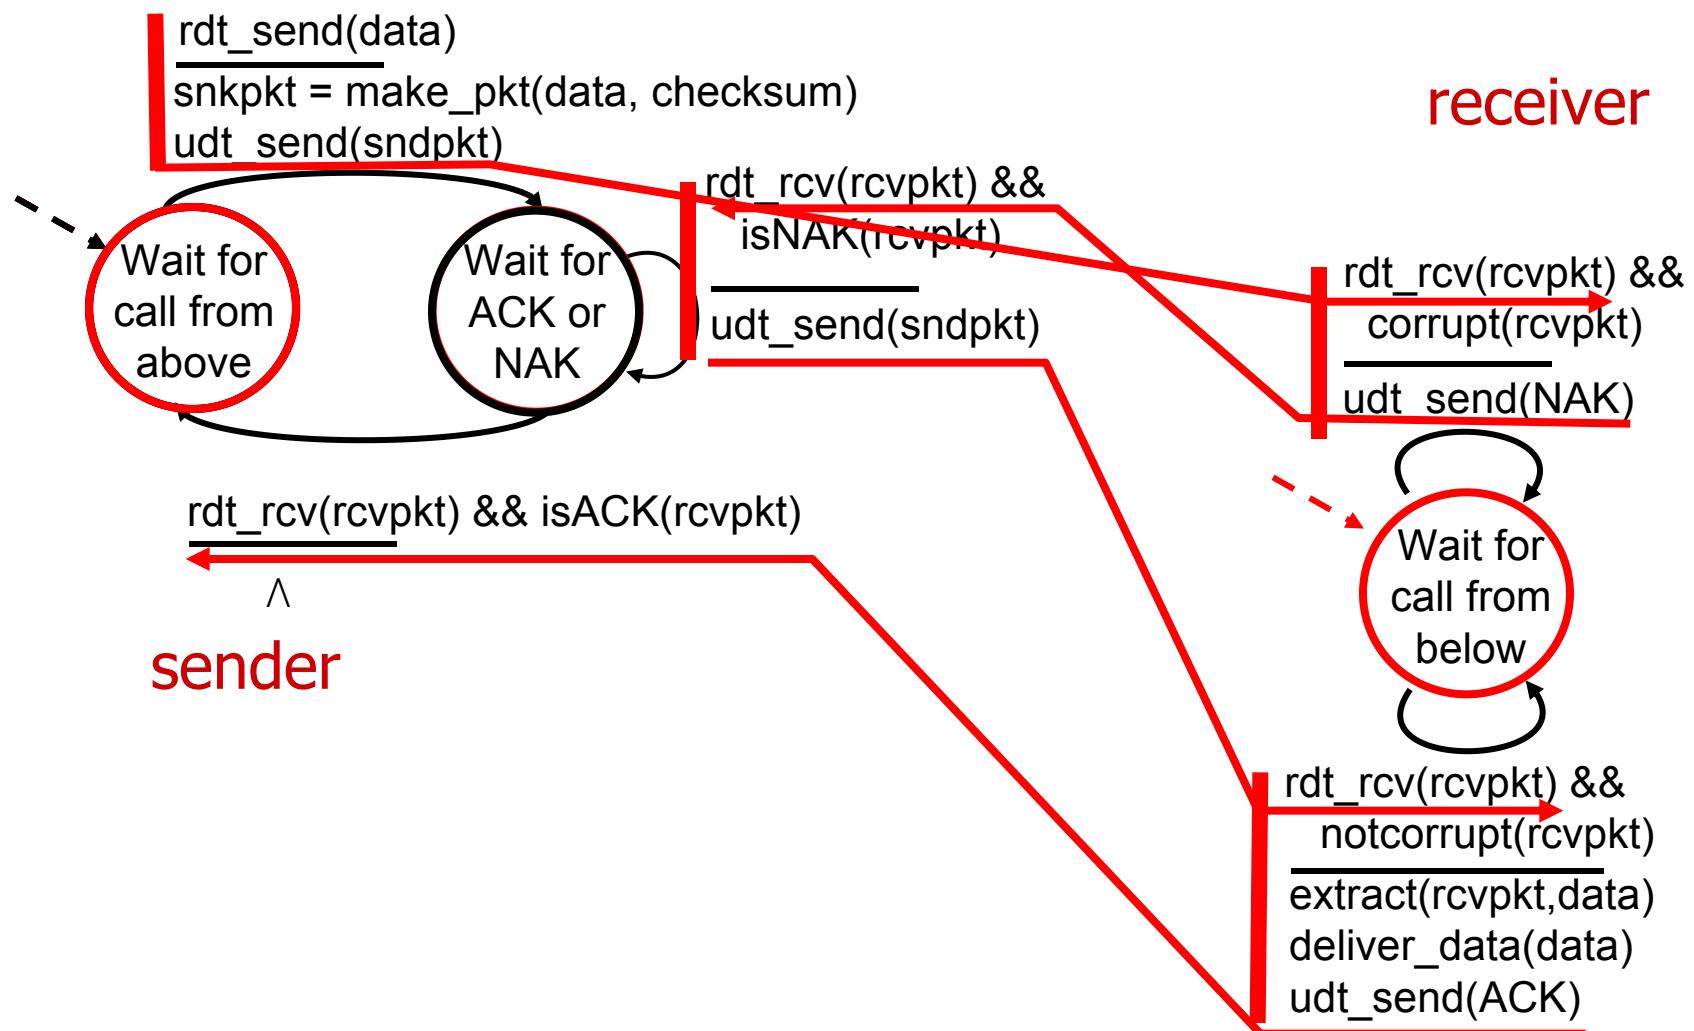
Sender sends one packet,
then waits for receiver
response

rdt_rcv(rcvpkt) &&
notcorrupt(rcvpkt)
extract(rcvpkt,data)
deliver_data(data)
udt_send(ACK)

rdt2.0: operation with no errors



rdt2.0: error scenario



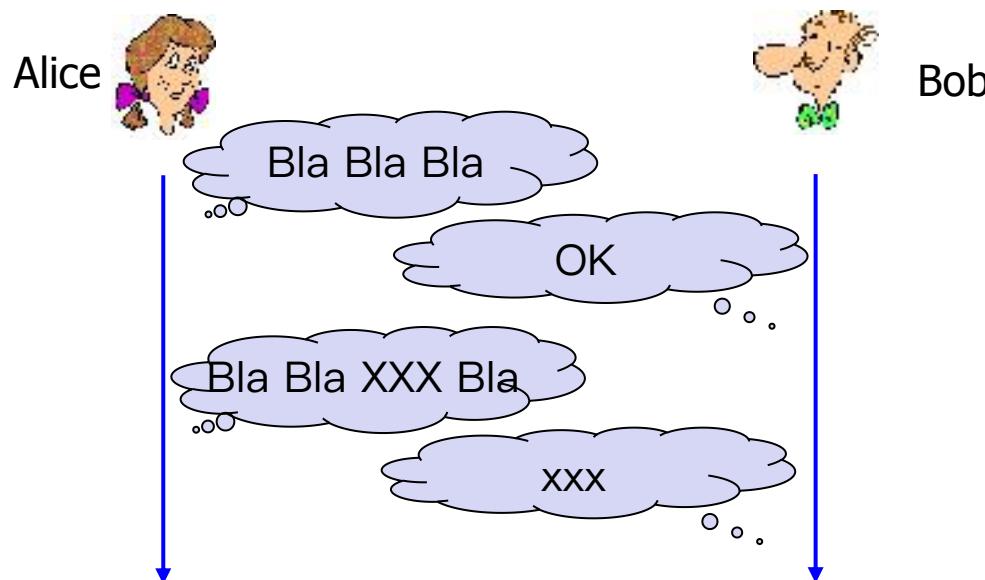
rdt2.0 has a fatal flaw!

The possibility that ACK or NAK packet could be corrupted:

- ❖ Checksum bits

Handling corrupted ACKs or NAKs:

- ❖ Option 1: “blabla...”, “OK”, “What did you say?”, “OK”
 - “What did you say?”, “What did you say?”, ...
- ❖ Option 2: add enough checksum to recover
- ❖ Option 3: when garbled ACK or NAK, retransmit



rdt2.0 has a fatal flaw!

The possibility that ACK or NAK packet could be corrupted:

- ❖ Checksum bits

Handling corrupted ACKs or NAKs:

- ❖ Option 1: “blabla...”, “OK”, “What did you say?”, “OK”
 - “What did you say?”, “What did you say?”, ...
- ❖ Option 2: add enough checksum to recover
- ❖ **Option 3: when garbled ACK or NAK, retransmit**

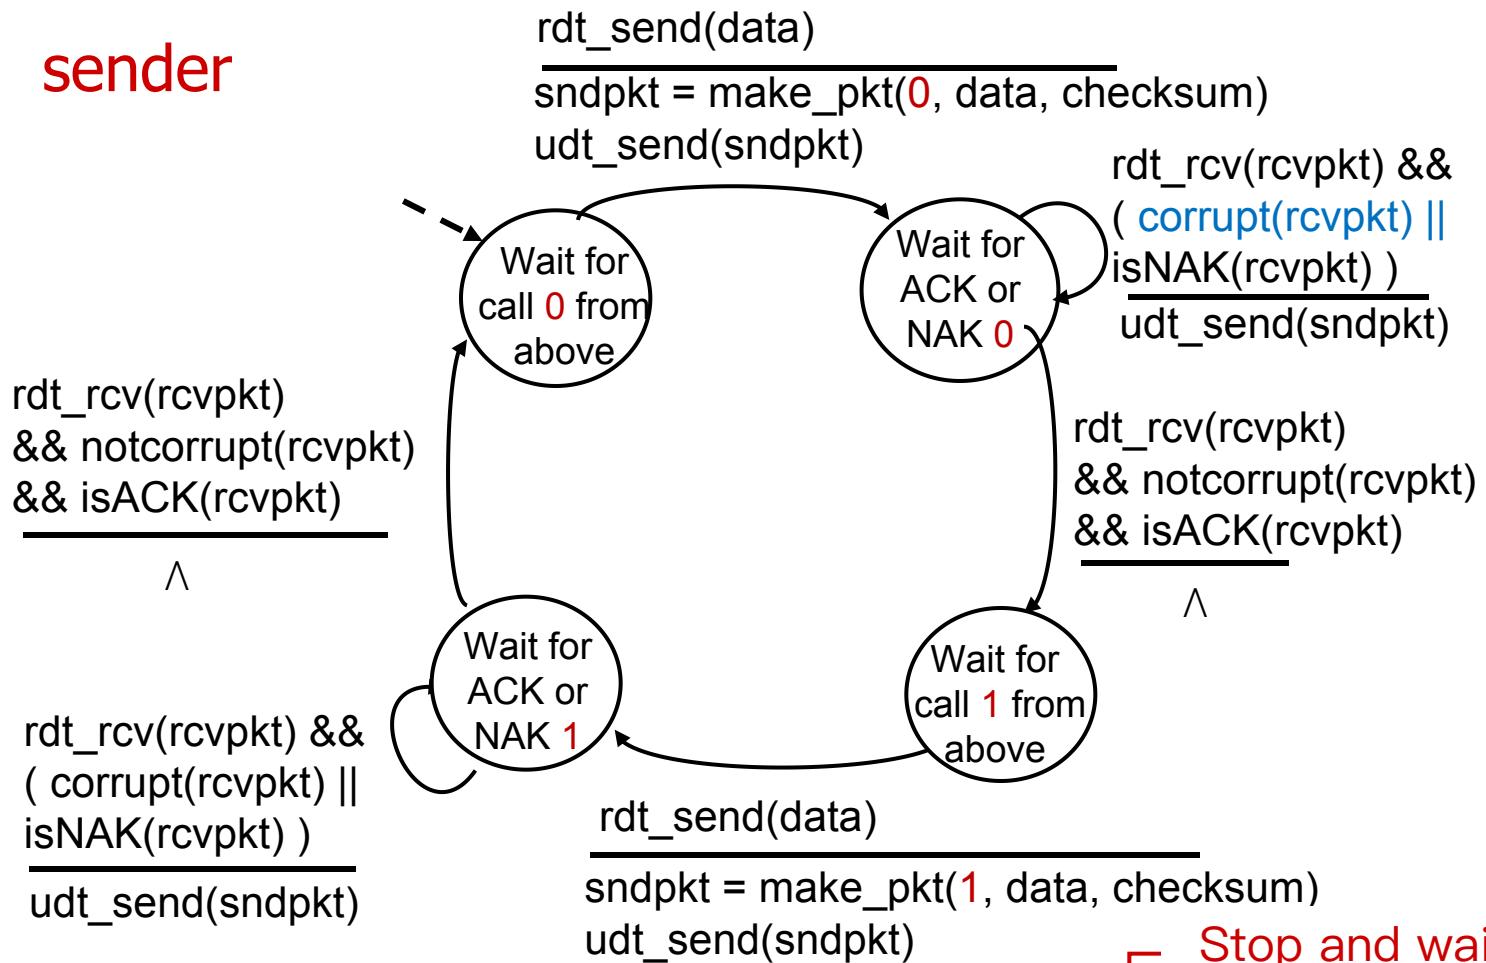
Problem: can't just retransmit: new data or retransmission?
possible duplicate

Handling duplicates:

- ❖ sender retransmits current pkt if ACK/NAK corrupted
- ❖ sender adds *sequence number* to each pkt
- ❖ receiver discards (doesn't deliver up) duplicate pkt

rdt2.1: sender, handles garbled ACK/NAKs

sender



Stop and wait

Sender sends one packet, then
waits for receiver response

Two sequence number would be sufficient!

rdt2.1: receiver, handles garbled ACK/NAKs

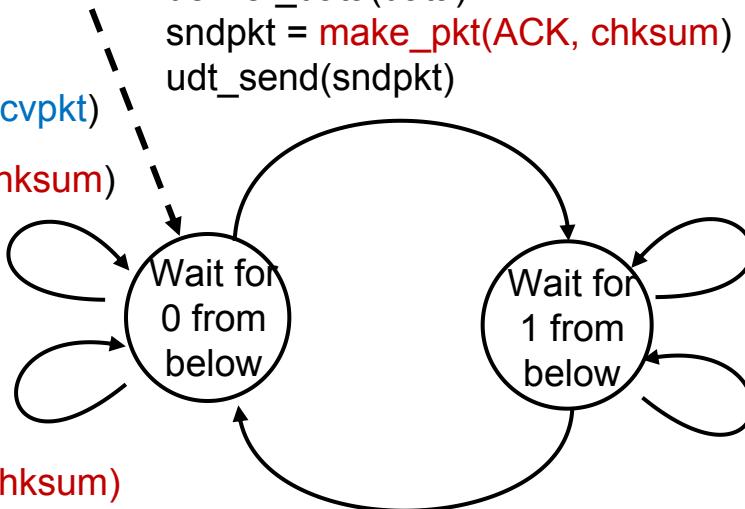
receiver

rdt_rcv(rcvpkt) && (corrupt(rcvpkt)

 sndpkt = make_pkt(NAK, chksum)
 udt_send(sndpkt)

rdt_rcv(rcvpkt) &&
 not corrupt(rcvpkt) &&
 has_seq1(rcvpkt)

 sndpkt = make_pkt(ACK, chksum)
 udt_send(sndpkt)



rdt_rcv(rcvpkt) && notcorrupt(rcvpkt)
 && has_seq0(rcvpkt)

extract(rcvpkt,data)
 deliver_data(data)
 sndpkt = make_pkt(ACK, chksum)
 udt_send(sndpkt)

No sequence number

rdt_rcv(rcvpkt) && (corrupt(rcvpkt)

 sndpkt = make_pkt(NAK, chksum)
 udt_send(sndpkt)

rdt_rcv(rcvpkt) &&
 not corrupt(rcvpkt) &&
 has_seq0(rcvpkt)

rdt_rcv(rcvpkt) && notcorrupt(rcvpkt)
 && has_seq1(rcvpkt)

extract(rcvpkt,data)
 deliver_data(data)
 sndpkt = make_pkt(ACK, chksum)
 udt_send(sndpkt)

Saying seq1 has been
received, please send the
next one

rdt2.1: discussion

sender:

- ❖ seq # added to pkt
- ❖ must check if received ACK/NAK corrupted
- ❖ twice as many states
 - state must “remember” whether “expected” pkt should have seq # of 0 or 1

receiver:

- ❖ must check if received packet is duplicate
 - state indicates whether 0 or 1 is expected pkt seq #
- ❖ note: receiver can *not* know if its last ACK/NAK received OK at sender

Two seq. #'s (0,1) will suffice. Why?

rdt2.2: a NAK-free protocol

- ❖ same functionality as rdt2.1, **using ACKs only**
- ❖ instead of NAK, receiver sends ACK for last pkt received OK
 - receiver must *explicitly* include seq # of pkt being ACKed
- ❖ **duplicate ACK** at sender results in same action as NAK: *retransmit current pkt*

rdt2.2: sender, receiver fragments

sender

rdt_send(data)

sndpkt = make_pkt(0, data, checksum)

udt_send(sndpkt)

Wait for
call 0 from
above

Wait for
ACK
0

sender FSM
fragment

rdt_rcv(rcvpkt) &&
(corrupt(rcvpkt) ||
isACK(rcvpkt,1))
udt_send(sndpkt)

rdt_rcv(rcvpkt)
&& notcorrupt(rcvpkt)
&& **isACK(rcvpkt,0)**

Λ

rdt_rcv(rcvpkt) &&
(corrupt(rcvpkt) ||
has_seq1(rcvpkt))
udt_send(sndpkt)

Wait for
0 from
below

receiver FSM
fragment

rdt_rcv(rcvpkt) && notcorrupt(rcvpkt)
&& has_seq1(rcvpkt)

extract(rcvpkt,data)

deliver_data(data)

sndpkt = make_pkt(ACK, 1, chksum)

udt_send(sndpkt)

receiver

rdt3.0: channels with errors and loss

New assumption: underlying channel can also lose packets (data, ACKs)

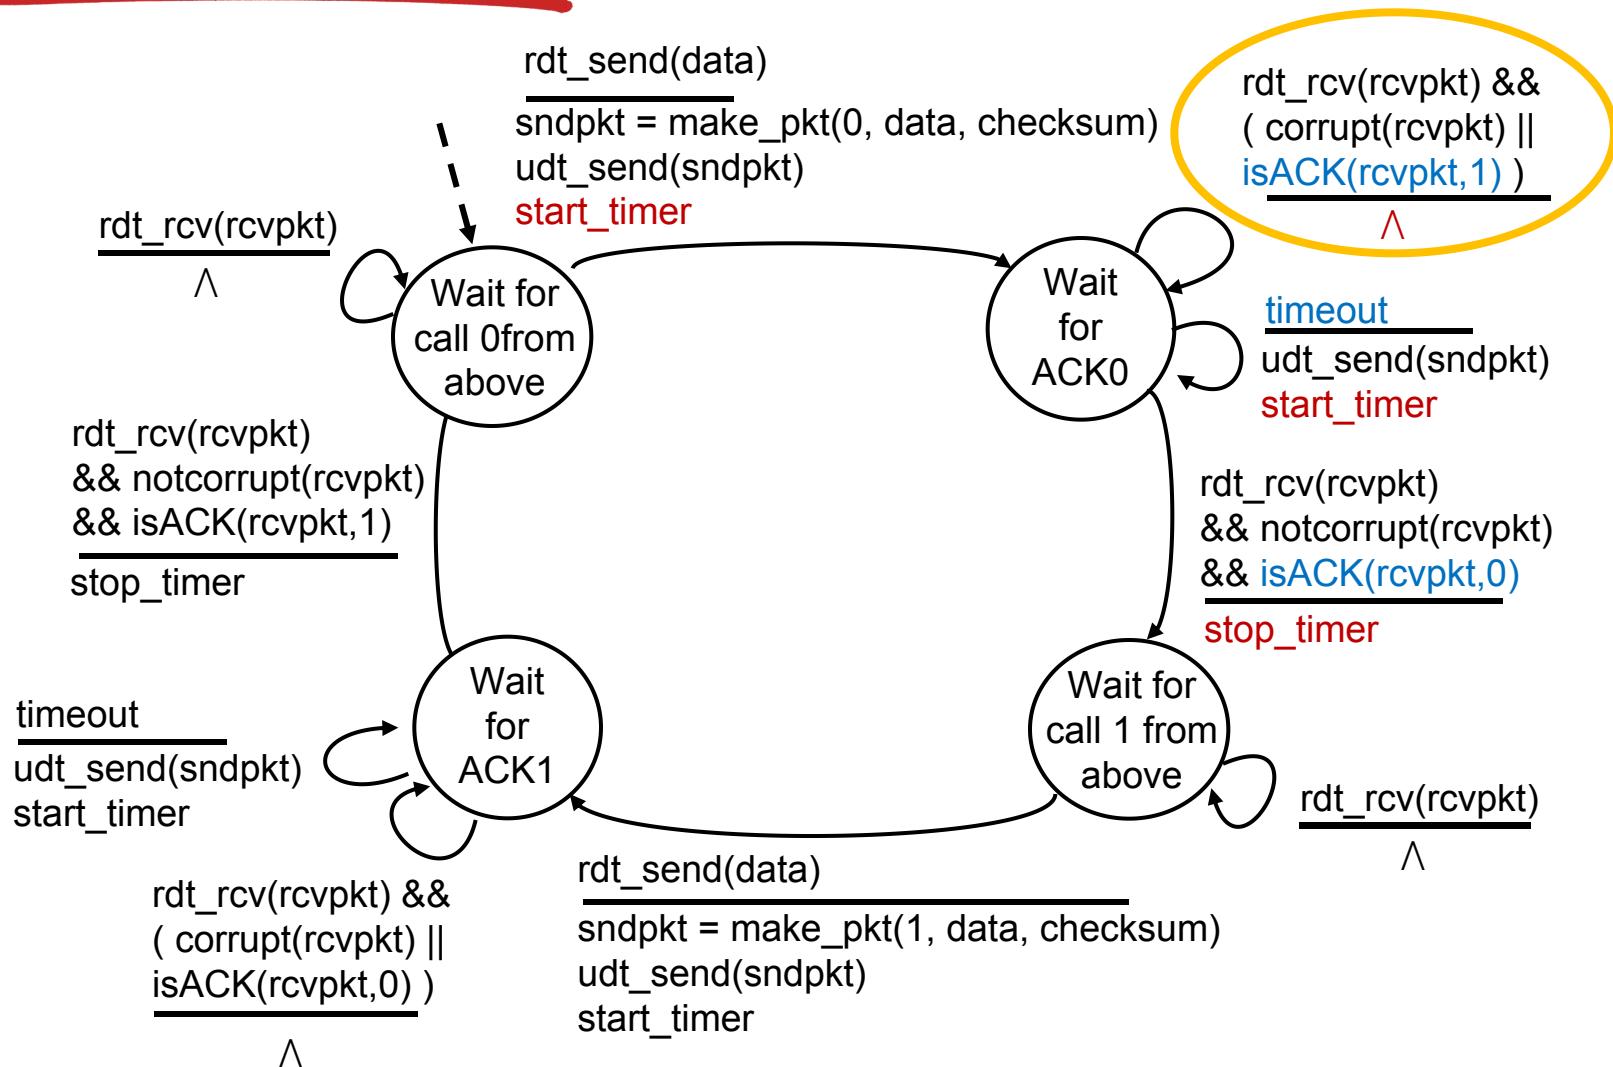
- checksum, seq. #, ACKs, retransmissions will be of help ... but not enough

Approach: sender waits “reasonable” amount of time for ACK

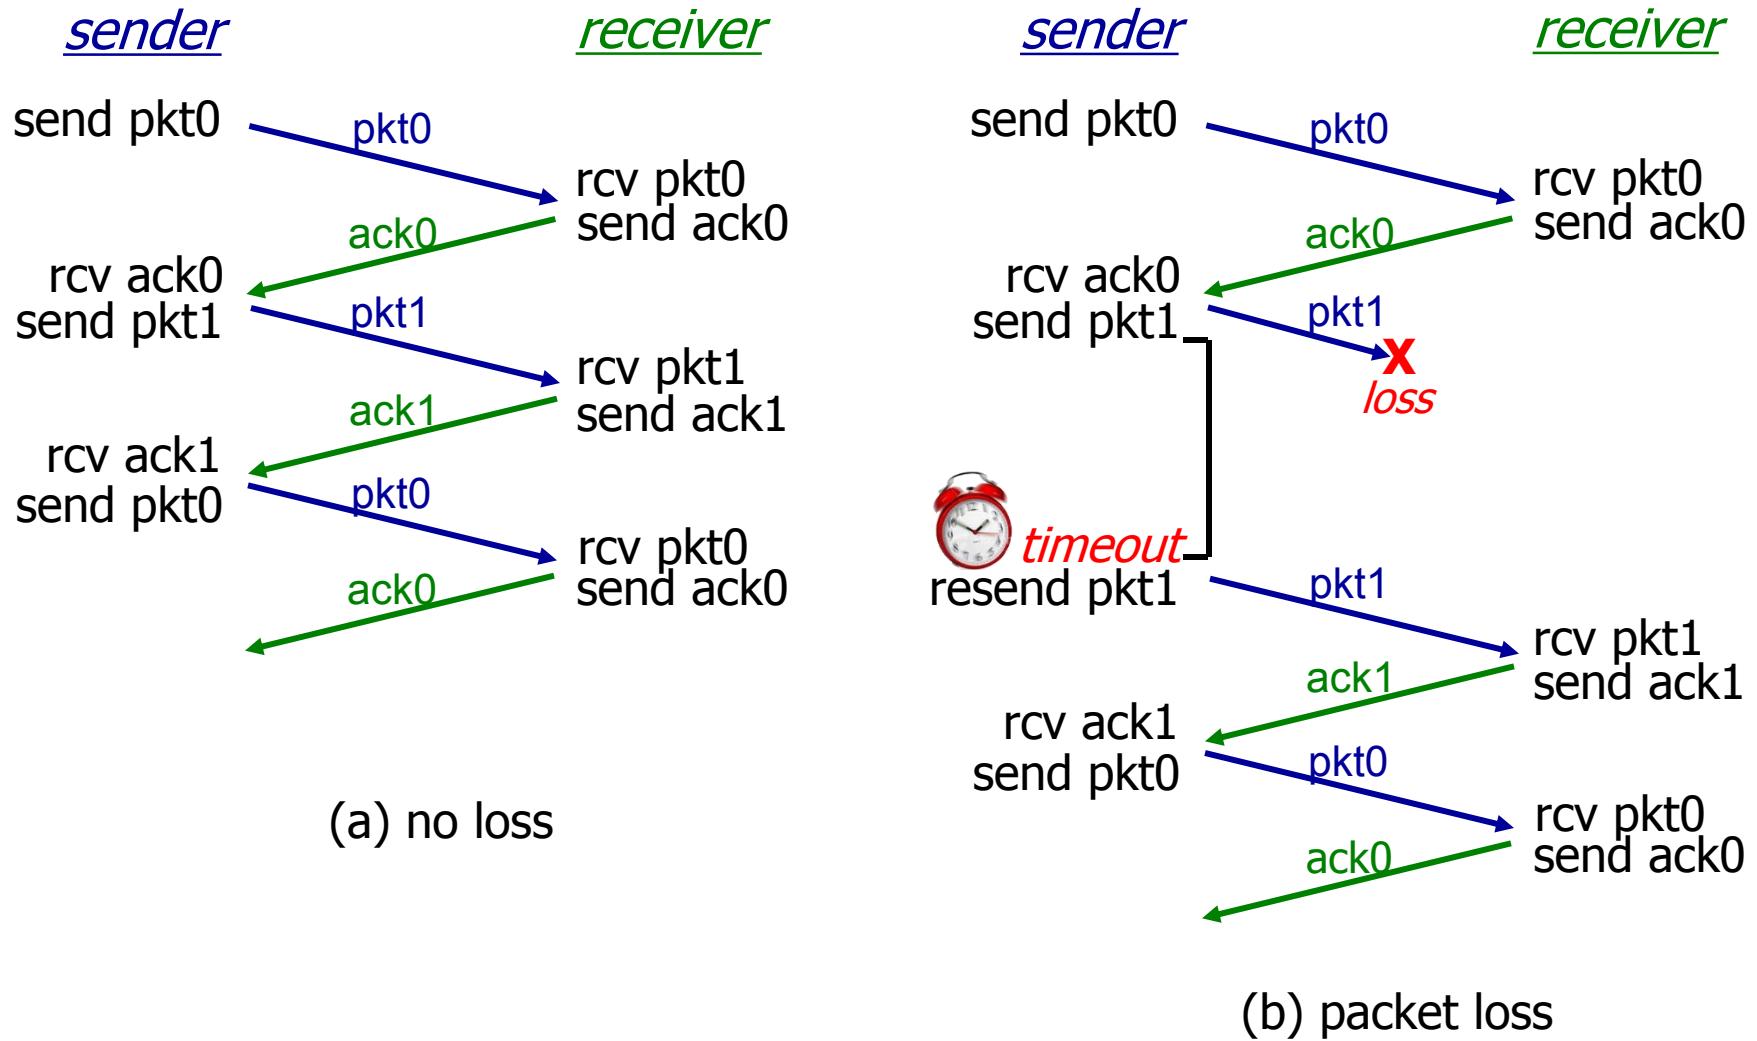
- ❖ retransmits if no ACK received in this time
- ❖ if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but seq. #'s already handles this
 - receiver must specify seq # of pkt being ACKed
- ❖ requires countdown timer
 - start timer, timer interrupt, stop timer

How long should the sender wait?

rdt3.0 sender

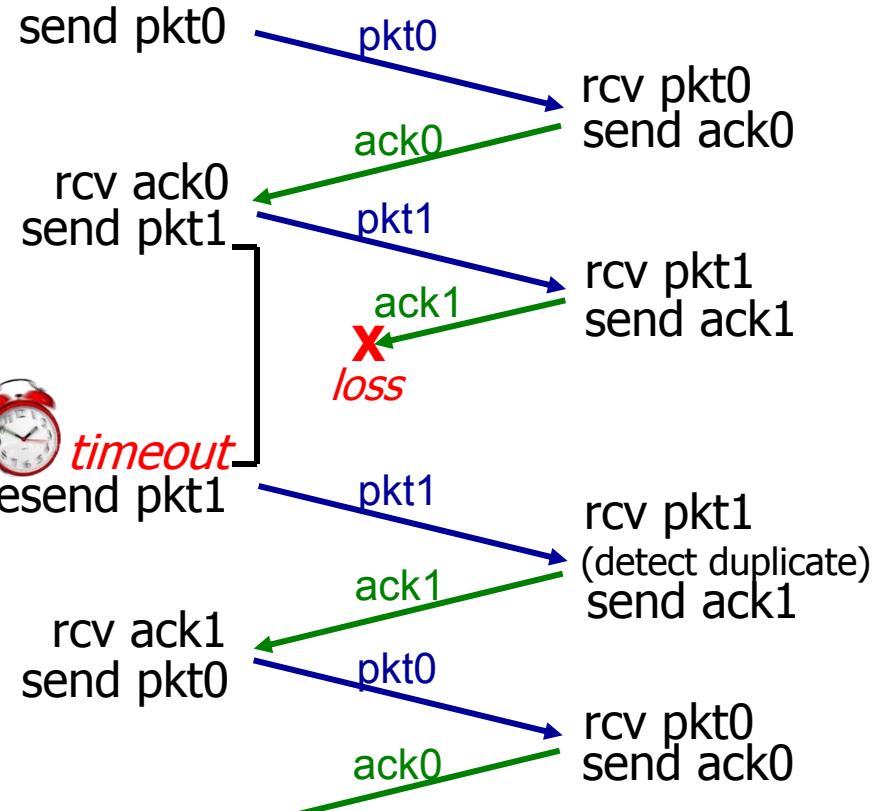


rdt3.0 in action



rdt3.0 in action

sender



(c) ACK loss

sender

send pkt0

rcv ack0
send pkt1

resend pkt1

rcv ack1
send pkt0

rcv ack1
send pkt0

pkt0

ack0

pkt1

ack1

pkt1

pkt0

ack1

ack0

receiver

rcv pkt0
send ack0

rcv pkt1
send ack1

rcv pkt1
(detect duplicate)

send ack1

rcv pkt0
send ack0



timeout

(d) premature timeout/ delayed ACK

Summary

Roadmap:

- ❖ Perfectly reliable channel: rdt1.0
- ❖ Channel with bit error:
 - bit error in packet: rdt 2.0
 - bit error in ACK: 2.1
 - NAK-free: 2.2
- ❖ Lossy channel: rdt 3.0

Summary of Techniques

- Checksum
- Sequence number
- ACK packets
- Retransmission
- Timeout

Next Lecture

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer (continue)

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control