

# CS 305: Computer Networks

## Fall 2024

### **Lecture 9: Transport Layer**

**Ming Tang**

Department of Computer Science and Engineering  
Southern University of Science and Technology (SUSTech)

# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 **connection-oriented transport: TCP**

- segment structure, RTT measurement
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

# TCP Reliable Data Transfer

- ❖ Segment structure
- ❖ Round-trip time estimation
- ❖ **Reliable data transfer**
- ❖ Flow control
- ❖ Control management

# TCP reliable data transfer

- ❖ TCP creates rdt service on top of IP's unreliable service
  - pipelined segments: window size, SendBase
  - **cumulative acks**
  - single retransmission timer
- ❖ retransmissions triggered by:
  - timeout events
  - duplicate acks

Let's initially consider simplified TCP sender:

- ignore duplicate acks
- ignore flow control, congestion control

# TCP sender events:

## *data rcvd from app:*

- ❖ create segment with seq #
- ❖ seq # is byte-stream number of first data byte in segment
- ❖ start timer if not already running
  - think of timer as for oldest unacked segment
  - expiration interval: **TimeoutInterval**

## *timeout:*

- ❖ retransmit segment that caused timeout
- ❖ restart timer

## *ack rcvd:*

- ❖ if ack acknowledges previously unacked segments
  - update what is known to be ACKed
  - start timer if there are still unacked segments

# TCP sender events:

```
NextSeqNum=InitialSeqNumber  
SendBase=InitialSeqNumber
```

```
loop (forever) {  
    switch(event)
```

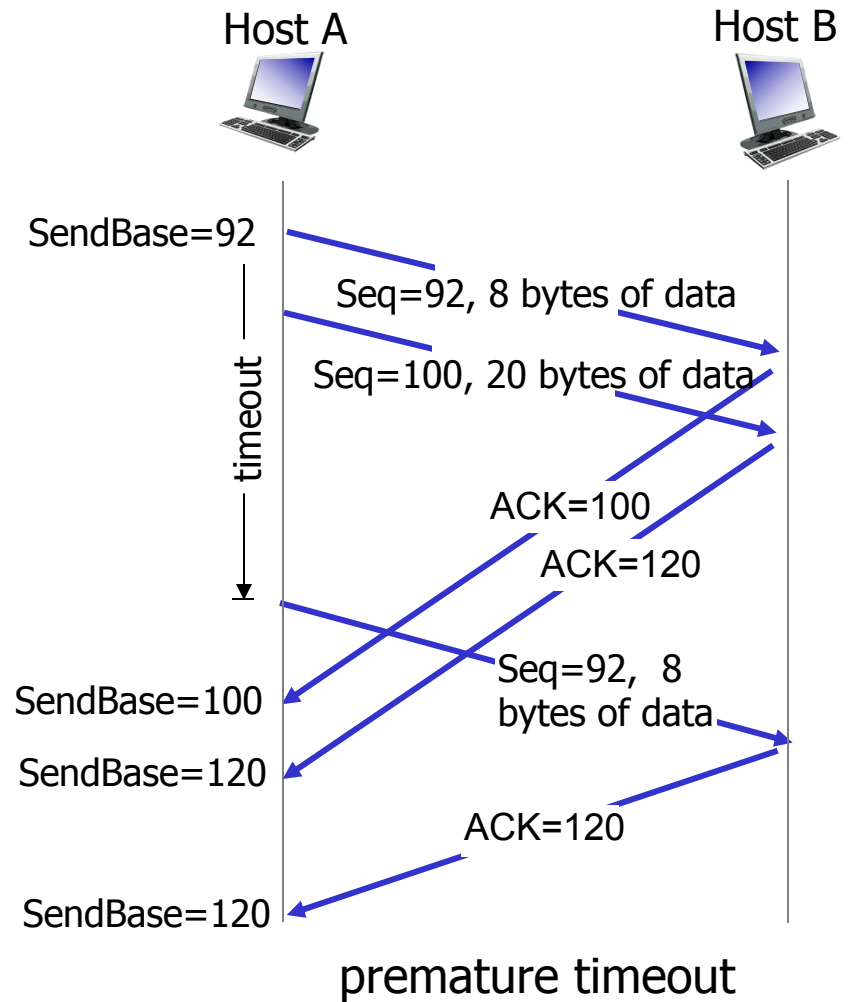

```
event: data received from application above  
    create TCP segment with sequence number NextSeqNum  
    if (timer currently not running)  
        start timer  
    pass segment to IP  
    NextSeqNum=NextSeqNum+length(data)  
    break;
```

```
event: timer timeout  
    retransmit not-yet-acknowledged segment with  
        smallest sequence number  
    start timer  
    break;
```


```
event: ACK received, with ACK field value of y  
    if (y > SendBase) {  
        SendBase=y  
        if (there are currently any not-yet-acknowledged segments)  
            start timer  
    }  
    break;
```

```
} /* end of loop forever */
```

# TCP: retransmission scenarios



# TCP: retransmission scenarios





# TCP receiver [RFC 1122, RFC 2581]

<i>event at receiver</i>	<i>TCP receiver action</i>
arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed	delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK
arrival of in-order segment with expected seq #. One other segment has ACK pending	immediately send single cumulative ACK, ACKing both in-order segments
arrival of out-of-order segment higher-than-expect seq. # . Gap detected	immediately send <i>duplicate ACK</i> , indicating seq. # of next expected byte
arrival of segment that partially or completely fills gap	immediate send ACK, provided that segment starts at lower end of gap

# Double Timeout Interval

Each time TCP retransmits, it sets the **next timeout interval to twice** the previous value

- ❖ Suppose *TimeoutInterval* associated with the oldest not yet acknowledged segment is 0.75 sec.
- ❖ when the timer first expires, TCP will then retransmit this segment and set the new expiration time to 1.5 sec.
- ❖ If the timer expires again, TCP will again retransmit this segment, now setting the expiration time to 3.0 sec.

When **the timer is started** after either of the two other events (that is, data received from application above, and ACK received):

- ❖ the *TimeoutInterval* is derived from the most recent values of *EstimatedRTT* and *DevRTT*.

Part of congestion control: being polite

# TCP fast retransmit

- ❖ time-out period often relatively long:
  - long delay before resending lost packet
- ❖ detect lost segments via duplicate ACKs.
  - sender often sends many segments back-to-back
  - if segment is lost, there will likely be many duplicate ACKs.

## *TCP fast retransmit*

if sender receives **3 duplicate ACKs** for same data  
( “triple duplicate ACKs” ), resend unacked segment with smallest seq #

- likely that unacked segment lost, so don't wait for timeout

# TCP fast retransmit

```
NextSeqNum=InitialSeqNumber
SendBase=InitialSeqNumber

loop (forever) {
    switch(event)


        event: data received from application
            create TCP segment with sequence
            if (timer currently not running)
                start timer
            pass segment to IP
            NextSeqNum=NextSeqNum+length(data)
            break;

        event: timer timeout
            retransmit not-yet-acknowledged segment ,
                smallest sequence number
            start timer
            break;

        event: ACK received, with ACK field value of y
            if (y > SendBase) {
                SendBase=y
                if (there are currently any not yet
                    acknowledged segments)
                    start timer
            }
            else { /* a duplicate ACK for already ACKed
                segment */
                increment number of duplicate ACKs
                received for y
                if (number of duplicate ACKs received
                    for y==3)
                    /* TCP fast retransmit */
                    resend segment with sequence number y
            }
            break;

    } /* end of loop forever */
```

# TCP fast retransmit




fast retransmit after sender receipt of triple duplicate ACK

# TCP Reliable Data Transfer

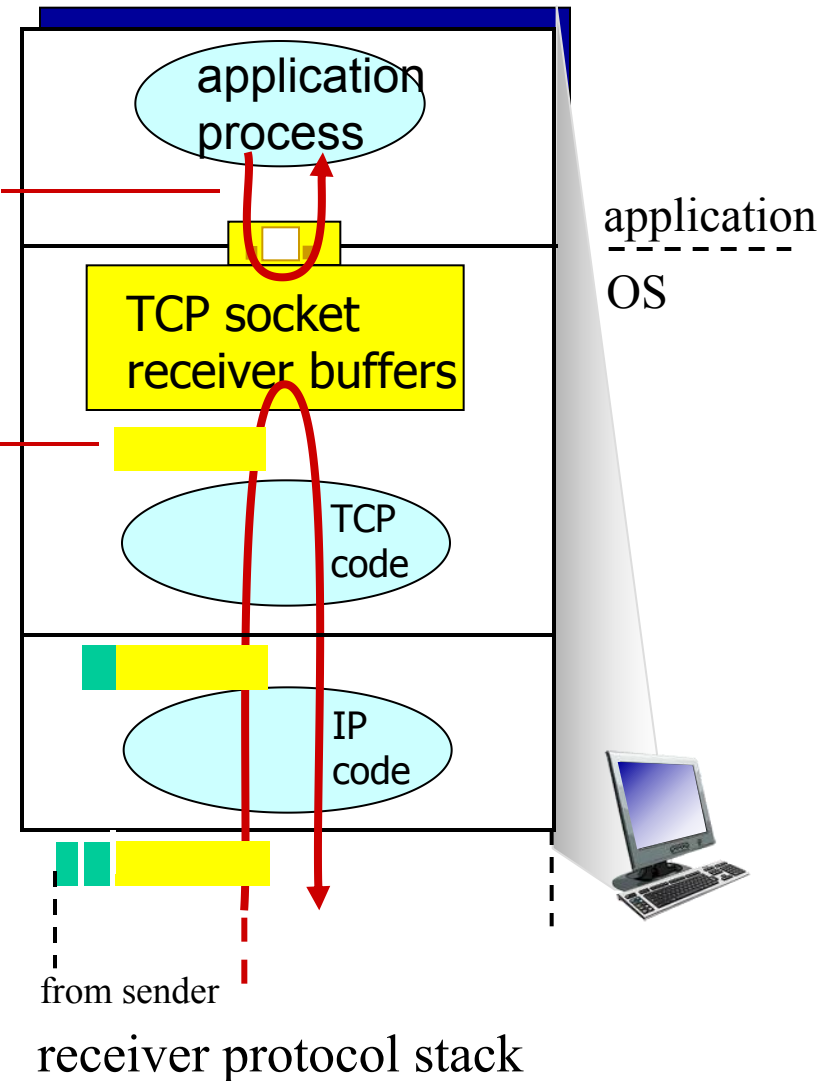
- ❖ Segment structure
- ❖ Round-trip time estimation
- ❖ Reliable data transfer
- ❖ **Flow control**
- ❖ Control management

# TCP: Overview

---



- TCP connection
- TCP grab data from the sender buffer
- TCP receives a segment at the other end, place it in receiver buffer
- application reads the stream from the receive buffer

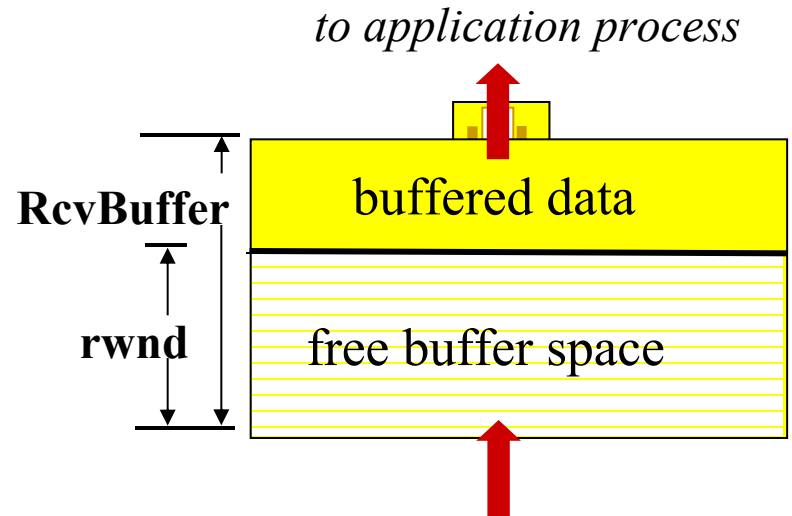




# TCP flow control

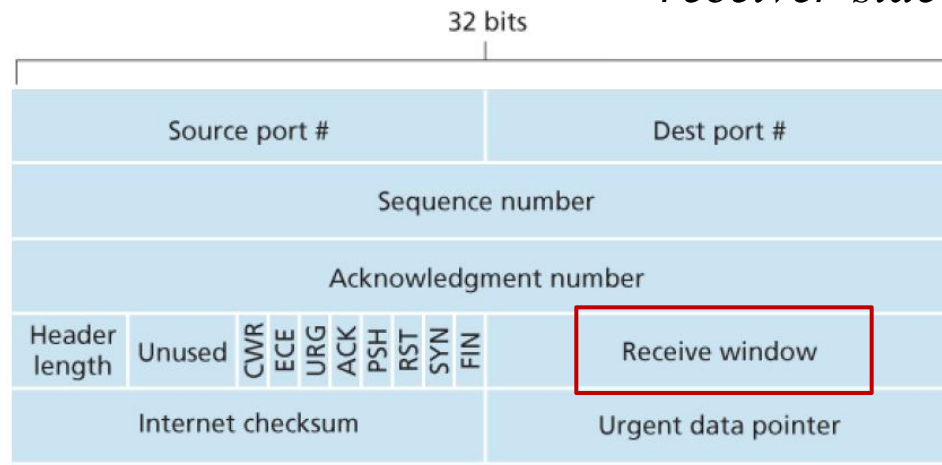
**Receiver** “advertises” free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments

- ❖ **RcvBuffer** size set via socket options (typical default is 4096 bytes)
- ❖ many operating systems autoadjust **RcvBuffer**



*TCP segment payloads*  
*receiver-side buffering*

$$\text{rwnd} = \text{RcvBuffer} - [\text{LastByteRcvd} - \text{LastByteRead}]$$

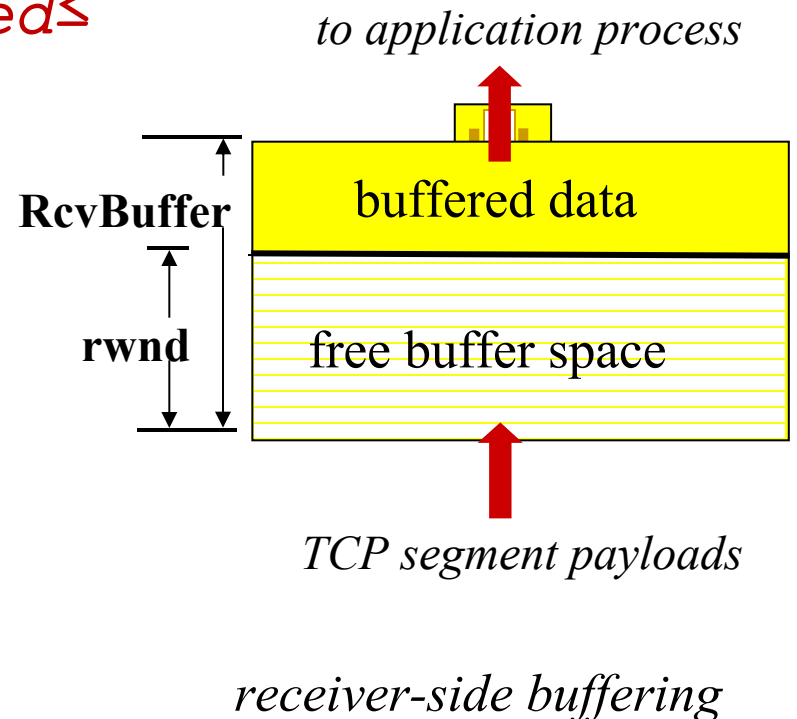


# TCP flow control

**Sender** limits amount of unacked ( “in-flight” ) data to receiver's **rwnd** value

*$LastByteSent - LastByteAcked \leq rwnd$*

Guarantees receive buffer will not overflow




# TCP Reliable Data Transfer

- ❖ Segment structure
- ❖ Round-trip time estimation
- ❖ Reliable data transfer
- ❖ Flow control
- ❖ Control management

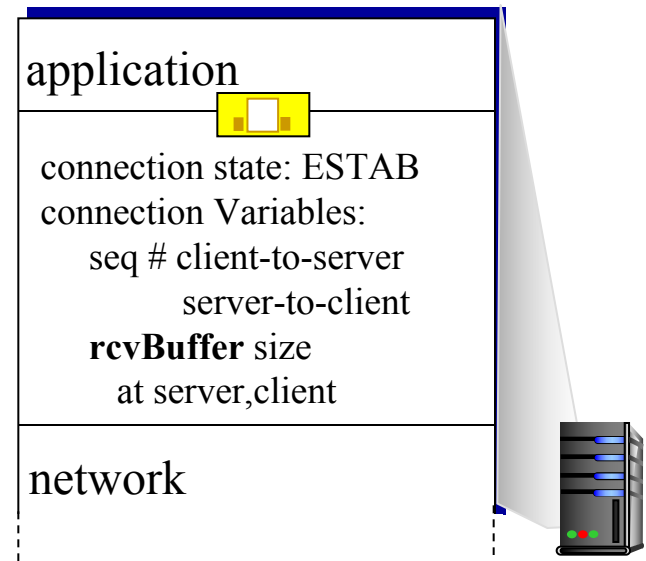
# Connection Management

before exchanging data, sender/receiver “handshake” :

- ❖ agree to establish connection (each knowing the other willing to establish connection)
- ❖ agree on connection parameters




```
clientSocket = socket(AF_INET, SOCK_STREAM);  
clientSocket.connect((hostname,port number));
```



```
connectionSocket = welcomeSocket.accept();
```

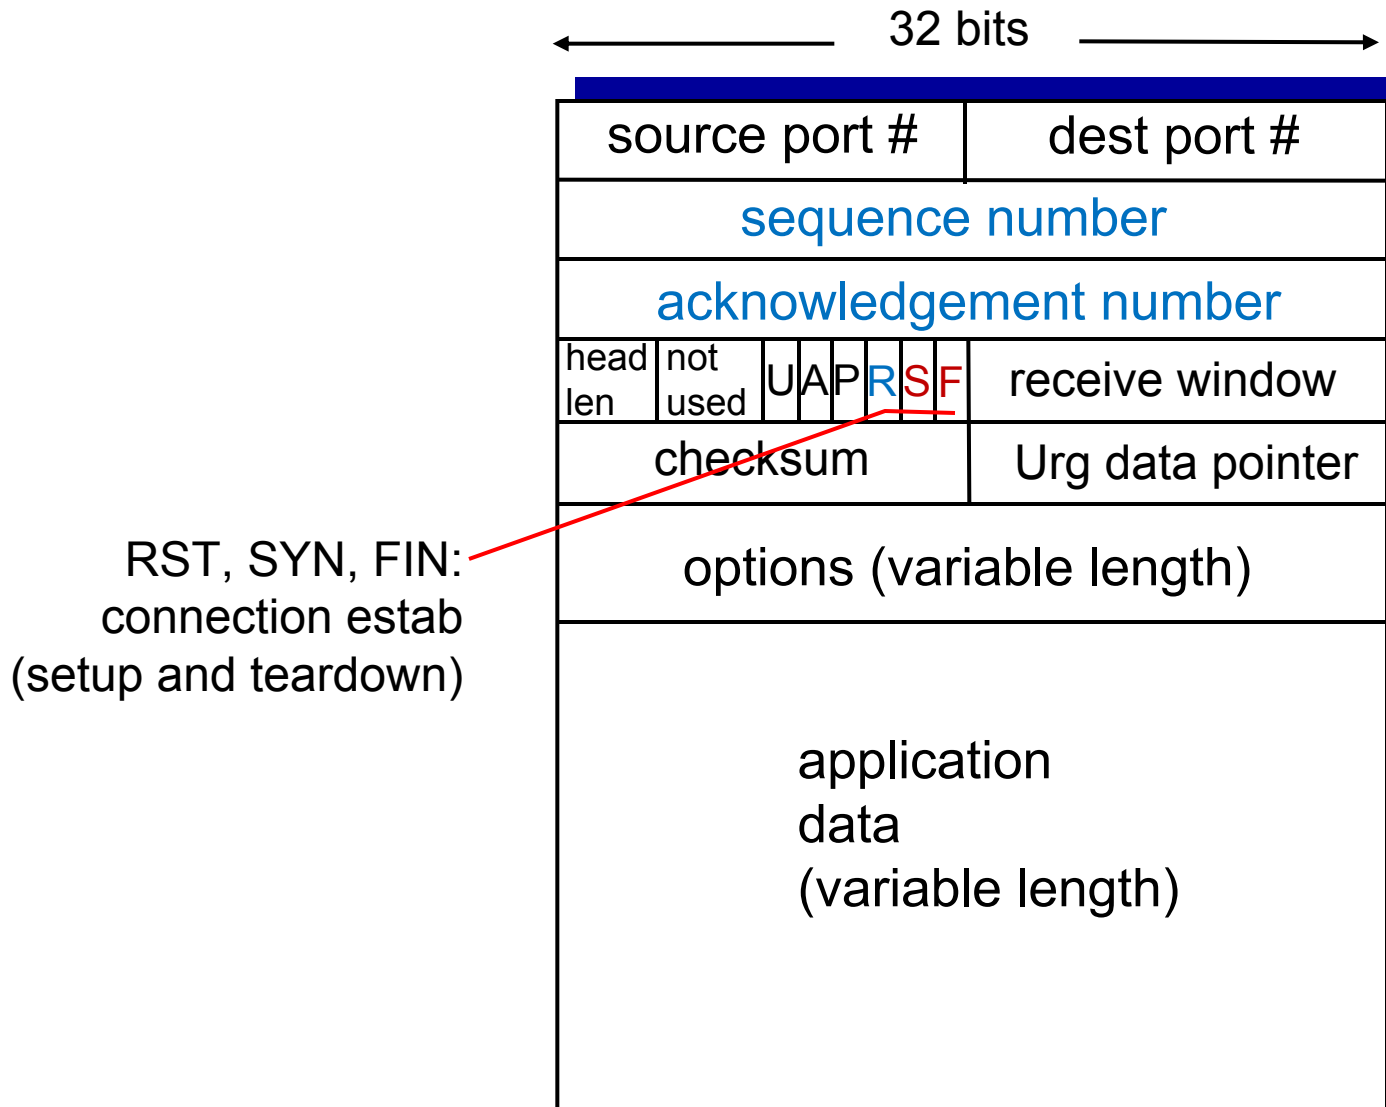
# Agreeing to establish a connection

3-way handshake:



- ❖ T1: B knows A's transmitter and B's receiver is OK
- ❖ T2: A knows A's transceiver and B's transceiver is OK, B has no more information than T1
- ❖ T3: Both A and B know their transceiver are OK, they can start the communication!

# TCP segment structure



# TCP 3-way handshake

## client state

LISTEN

SYNSENT

ESTAB

choose init seq num, x  
send TCP SYN msg

received SYNACK(x)  
indicates server is live;  
send ACK for SYNACK;  
this segment may contain  
client-to-server data



SYNbit=1, Seq=x

SYNbit=1, Seq=y  
ACKbit=1; ACKnum=x+1

SYNbit=0  
ACKbit=1, ACKnum=y+1



choose init seq num, y  
send TCP SYNACK  
msg, acking SYN

received ACK(y)  
indicates client is live

## server state

LISTEN

SYN RCVD

ESTAB

Once these three steps have been completed, the client and server hosts can send segments containing data to each other.

- In each of these future segments, SYNbit=0

# TCP: closing a connection

## client state

ESTAB

clientSocket.close()

FIN\_WAIT\_1

can no longer  
send but can  
receive data

FIN\_WAIT\_2

wait for server  
close

TIMED\_WAIT

timed wait

CLOSED



FINbit=1, seq=x

ACKbit=1; ACKnum=x+1

FINbit=1, seq=y

ACKbit=1; ACKnum=y+1

can still  
send data

can no longer  
send data

## server state

ESTAB

CLOSE\_WAIT

LAST\_ACK

CLOSED

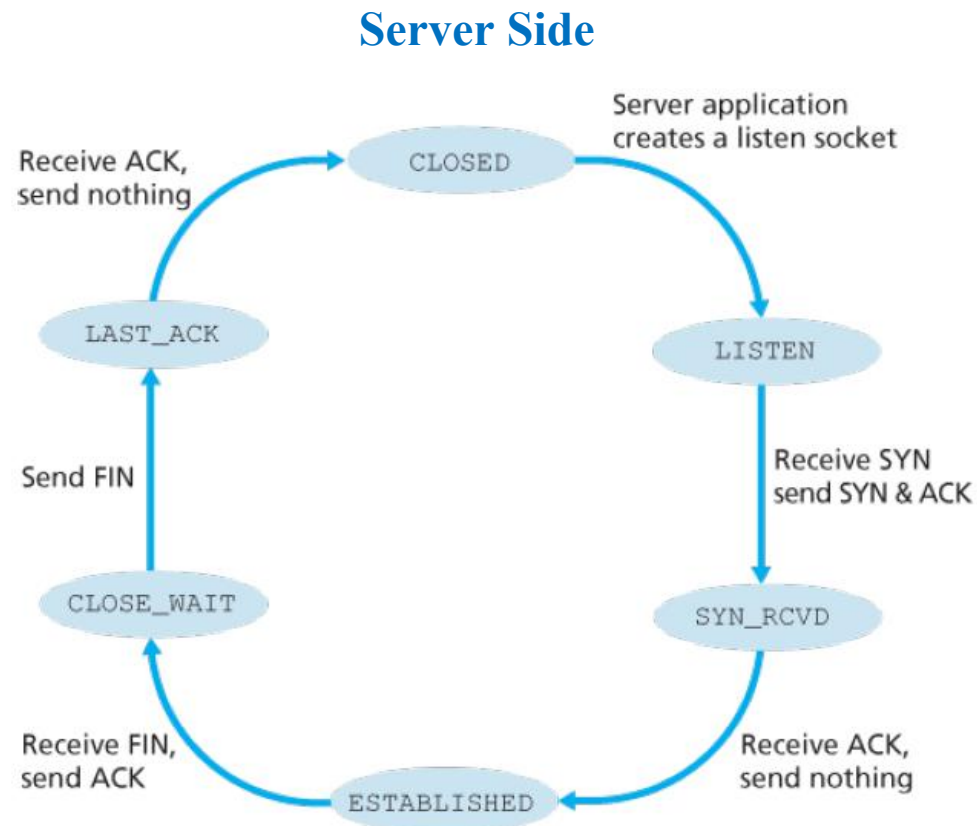
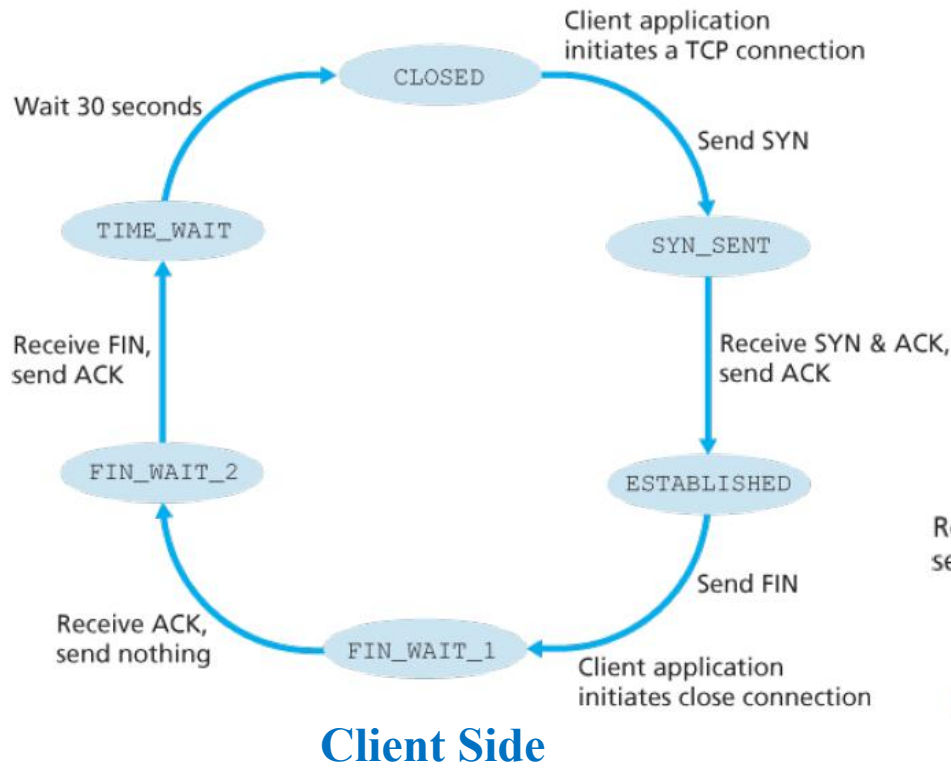
The TIME\_WAIT state lets the TCP client resend the final acknowledgment in case the ACK is lost.



# TCP: closing a connection

- ❖ Four-way handshaking
  - Either of the two processes participating in a TCP connection can end the connection.
- ❖ client, server each close their side of connection
  - send TCP segment with FIN bit = 1
- ❖ respond to received FIN with ACK
- ❖ Why FIN and ACK can not be sent in one msg as SYNACK in connection establishment?
  - The other side may still have packets need to be sent. It can not send FIN until the transmission is finished.

# TCP States



# Reset Segment

---

When a host receives a TCP segment whose port numbers or source IP address do not match with any of the ongoing sockets.

- ❖ Then the host will send a special reset segment to the source. RST flag bit is set to 1.
- ❖ “I don’t have a socket for that segment. Please do not resend the segment.”

# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control




# Principles of congestion control

## Congestion:

- ❖ informally: “too many sources sending too much data too fast for *network* to handle”
- ❖ different from flow control!
- ❖ manifestations:
  - lost packets (buffer overflow at routers)
  - long delays (queueing in router buffers)
- ❖ a top-10 problem!

# Causes/costs of congestion: scenario 1

- ❖ two senders, two receivers
- ❖ one router, **infinite** buffers
- ❖ output link capacity:  $R$
- ❖ **no retransmission**




*large queuing delays are experienced as the packet-arrival rate nears the link capacity.*

- ❖ maximum per-connection throughput:  $R/2$
- ❖ large delays as arrival rate,  $\lambda_{in}$ , approaches capacity

# Causes/costs of congestion: scenario 2



- ❖ one router, *finite* buffers
- ❖ sender retransmission of timed-out packet
  - application-layer input = application-layer output:  $\lambda_{in} = \lambda_{out}$
  - transport-layer input includes *retransmissions* :  $\lambda'_{in} \geq \lambda_{in}$



# Causes/costs of congestion: scenario 2

Idealization: perfect knowledge

- ❖ Sender sends only when router buffers available
- ❖ No loss occurs:  $\lambda'_{in} = \lambda_{in}$






# Causes/costs of congestion: scenario 2

*Idealization: known loss* packets can be lost, dropped at router due to full buffers

❖ sender only resends if packet *known* to be lost

❖  $\lambda'_{in} \geq \lambda_{in}$




# Causes/costs of congestion: scenario 2

*Idealization: known loss* packets can be lost, dropped at router due to full buffers

❖ sender only resends if packet *known* to be lost



❖  $\lambda'_{in} \geq \lambda_{in}$



# Causes/costs of congestion: scenario 2

## *Realistic: duplicates*

- ❖ packets can be lost, dropped at router due to full buffers
- ❖ sender times out prematurely, sending *two* copies, both of which are delivered



# Cause and Cost of Congestion

## **Cause**

- Shared link; limited link capacity
- Sending at a high rate

## **Cost of Congestion**

- Delay
- Packet lost and retransmission
- Unneeded retransmission: waste
- “upstream” transmission capacity was wasted (to be explained)

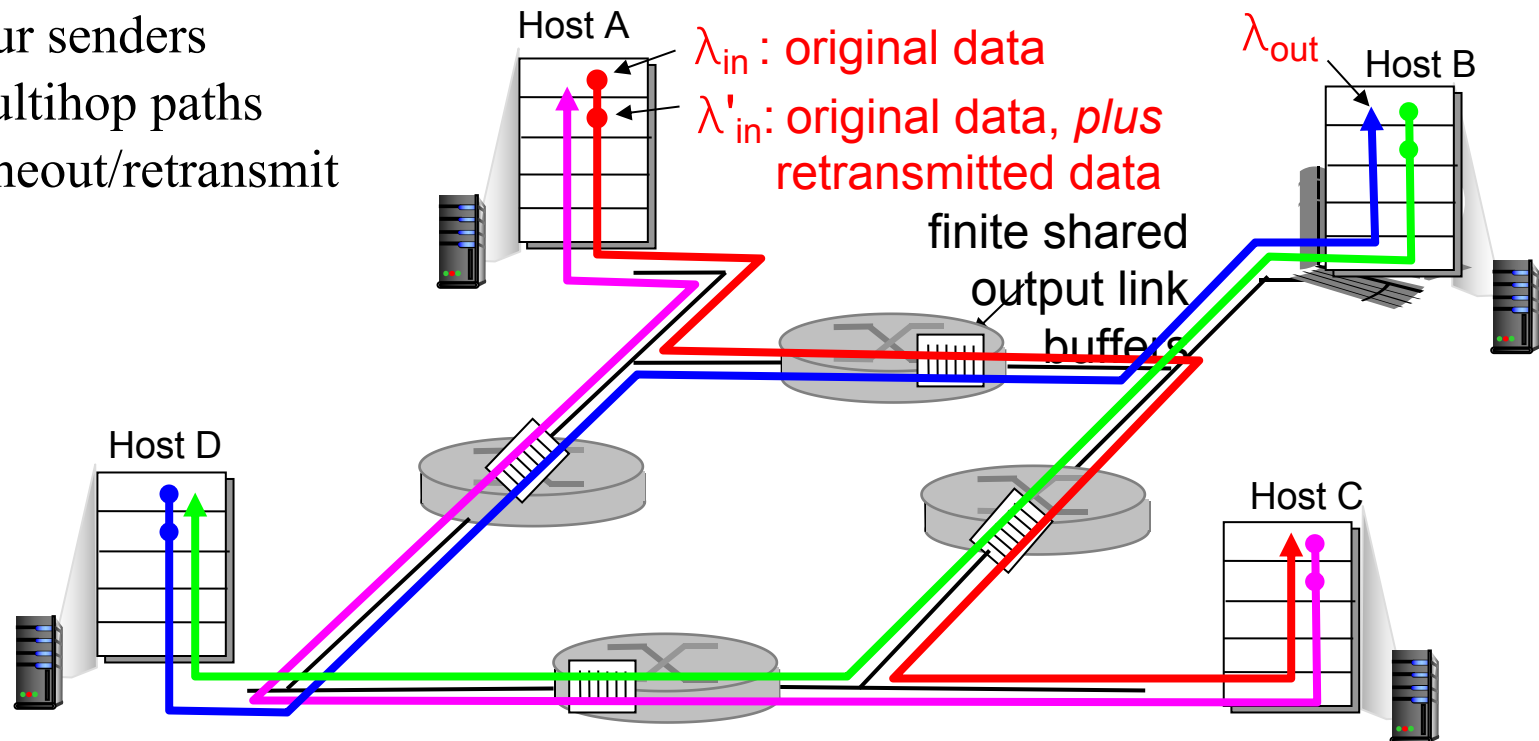
# Causes/costs of congestion: scenario 3

For small values of  $\lambda_{in}^-$ :

- buffer overflows are rare
- the throughput  $\lambda_{out}$  approximately equals the offered load

$$\lambda'_{in} = \lambda_{in}.$$

- ❖ four senders
- ❖ multihop paths
- ❖ timeout/retransmit




# Causes/costs of congestion: scenario 3



Q: what happens as  $\lambda_{in}$  increases ?

A: as green  $\lambda_{in}$  increases, all arriving red pkts at upper queue are dropped, red throughput goes 0

- ❖ four senders
- ❖ multihop paths
- ❖ timeout/retransmit



# Causes/costs of congestion: scenario 3



another “cost” of congestion:

- ❖ when packet dropped, any “upstream” transmission capacity used for that packet was wasted!

# Cause and Cost of Congestion

---

## **Cause**

- Shared link; limited link capacity
- Sending at a high rate

## **Cost of Congestion**

- Delay
- Packet lost and retransmission
- Unneeded retransmission: waste
- “upstream” transmission capacity was wasted



# Approaches to Congestion Control

## **End-to-end congestion control:**

- TCP segment loss or round-trip segment delay
- TCP decreases its **window size** accordingly

## **Network-assisted congestion control:**

- **routers provide feedback** to the sender and/or receiver
- a single bit indicating congestion at a link; the maximum host sending rate the router can support

# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

# TCP congestion control: additive increase multiplicative decrease

- TCP use **end-to-end** congestion control
- have each sender **limits the rate** at which it sends traffic into its connection as a function of **perceived network congestion**

## Questions for achieving congestion control:

**Q1:** How does a TCP **sender limit the rate** at which it sends traffic into its connection?

**Q2:** How does a TCP sender **perceive that there is congestion** on the path between itself and the destination?

**Q3:** What algorithm should the sender use to **change its send rate** as a function of perceived end-to-end congestion?

# TCP congestion control: additive increase multiplicative decrease

**Questions for achieving congestion control:**

**Q1:** How does a TCP sender limit the rate at which it sends traffic into its connection?

Congestion window: **cwnd**



$$\text{LastByteSent} - \text{LastByteAcked} \leq \min\{\text{cwnd}, \text{rwnd}\}$$

**Q2:** How does a TCP sender perceive that there is congestion on the path between itself and the destination?

Timeout; three duplicate ACKs

**Q3:** What algorithm should the sender use to change its send rate as a function of perceived end-to-end congestion?

# Congestion window: Throughput



- ❖ sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- ❖ **cwnd** is dynamic, function of perceived network congestion

*TCP sending rate:*

- ❖ *roughly*: send cwnd bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

# TCP congestion control: additive increase multiplicative decrease

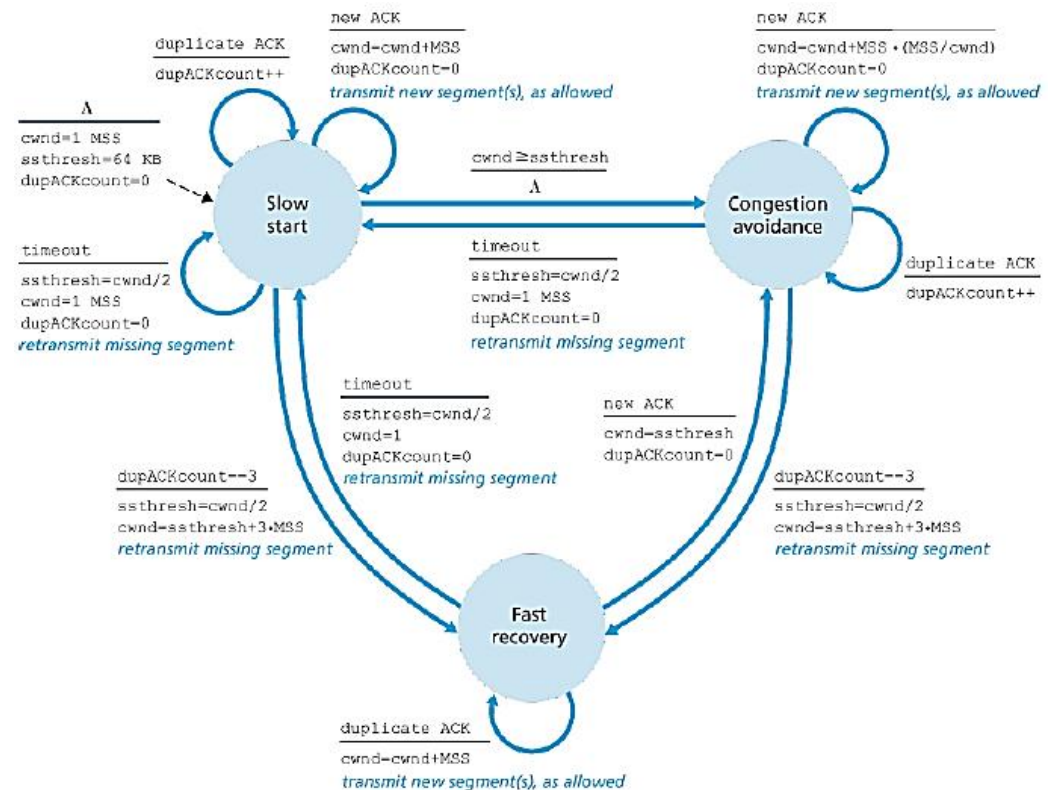

**Q3:** What algorithm should the sender use to **change its send rate** as a function of perceived end-to-end congestion?

- A lost segment → congestion → decrease rate
- An acknowledged segment → the network is fine → increase rate
- Bandwidth **probing**: network condition may change

# TCP Congestion Control: details


The congestion control algorithm has three major components:

- ❖ **Slow start:** exponentially increase
- ❖ **Congestion avoidance:** linearly increase
- ❖ **Fast recovery:**



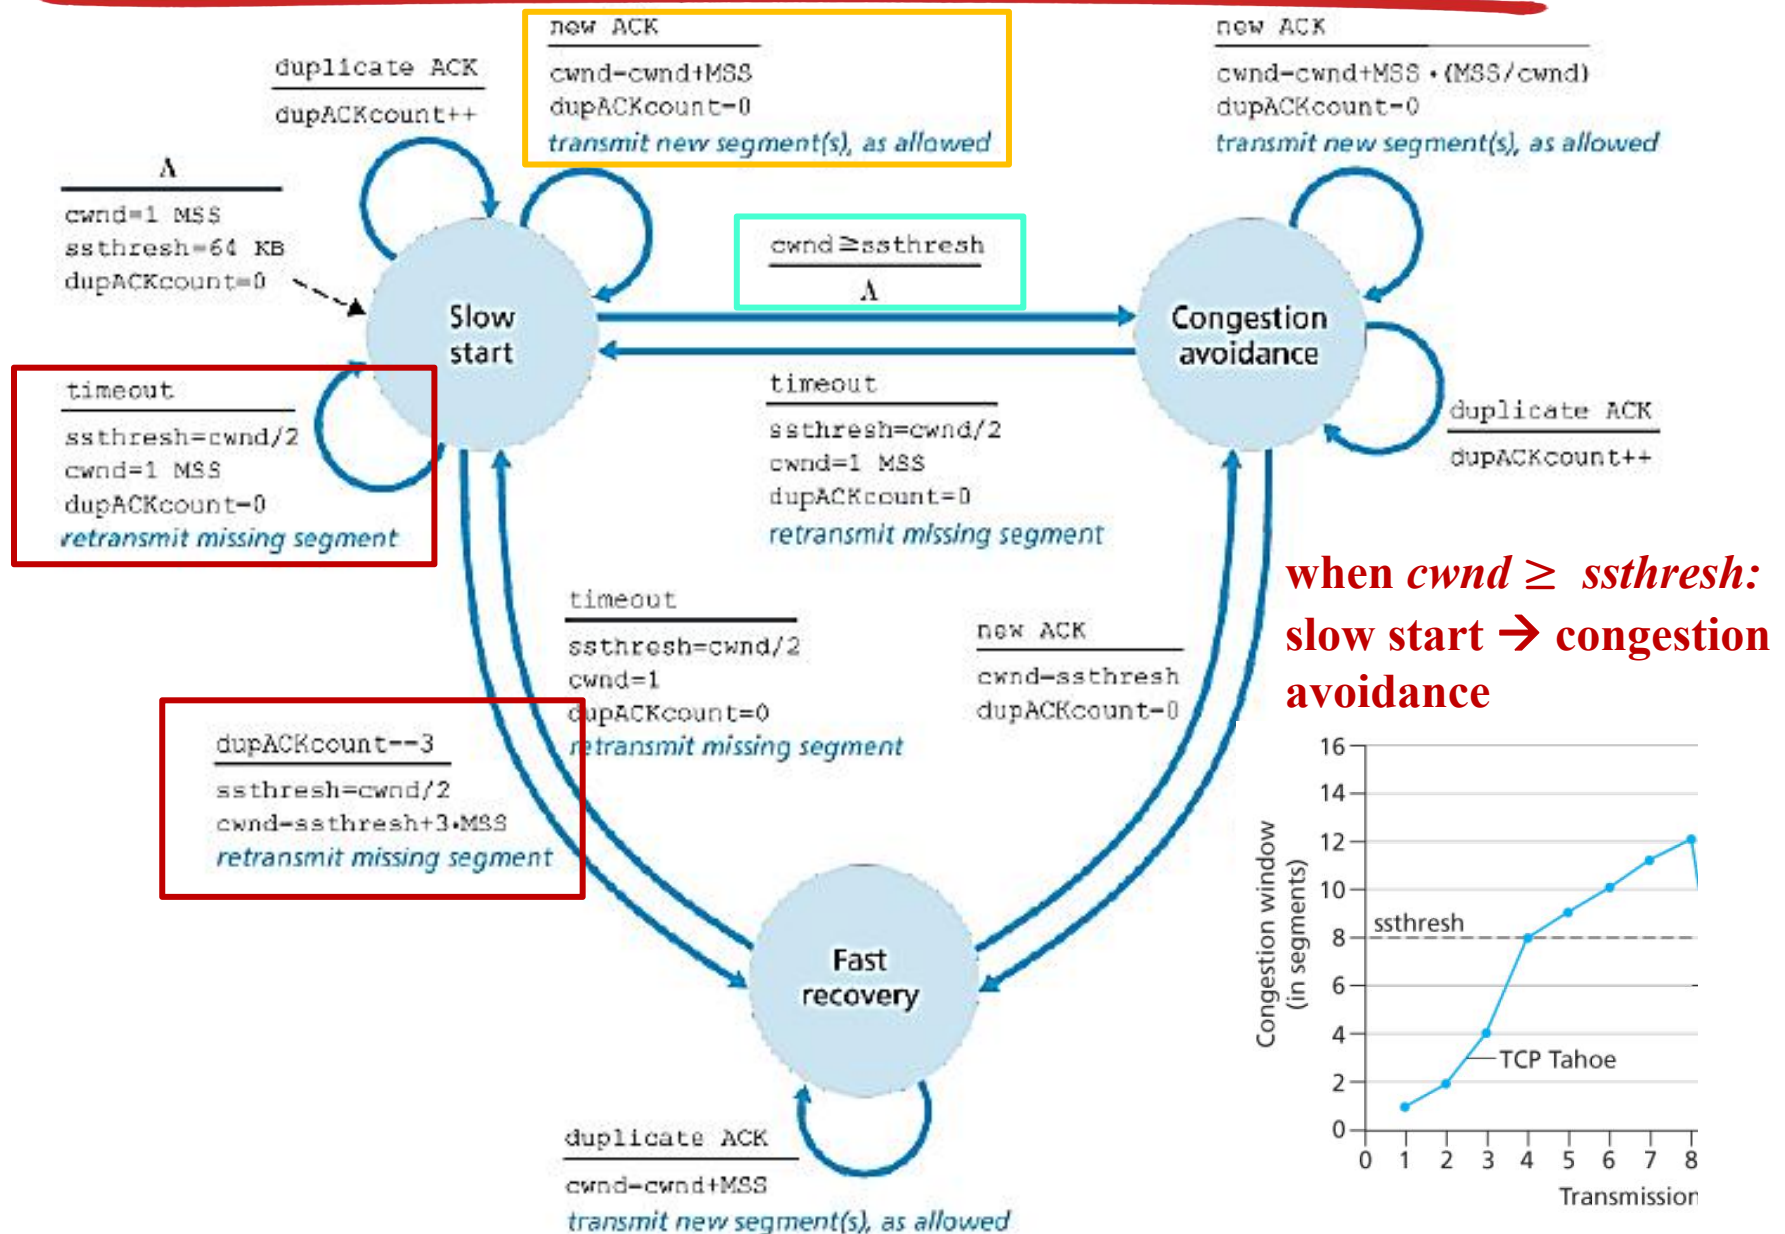
# TCP Slow Start

- ❖ when **connection begins** or **timeout** occurs, increase rate **exponentially** until packet lost:
  - initially **cwnd** = 1 MSS
  - double **cwnd** every RTT
  - done by incrementing **cwnd** for every ACK received
- ❖ Summary: initial rate is slow but ramps up exponentially fast





# TCP Congestion Control: FSM



# TCP: detecting, reacting to loss


- ❖ loss indicated by timeout:
  - **cwnd** set to 1 MSS; **ssthresh** = **cwnd**/2
  - window then grows exponentially (as in slow start) to threshold, then grows linearly
- ❖ loss indicated by 3 duplicate ACKs:
  - TCP Tahoe always sets **cwnd** to 1 (timeout or 3 duplicate acks) → Slow start
    - **cwnd** set to 1 MSS; **ssthresh** = **cwnd**/2
  - TCP RENO
    - dup ACKs indicate network capable of delivering some segments → Fast Recovery
    - $ssthresh = cwnd / 2$ ;  $cwnd = ssthresh + 3MSS$

# TCP: switching from slow start to Congestion Avoidance

**Q:** when should the **exponential increase** switch to linear?

**A:** when **cwnd** gets to 1/2 of its value before timeout

**Thus, when timeout occurs,  $ssthresh = cwnd/2$**




# TCP Congestion Avoidance

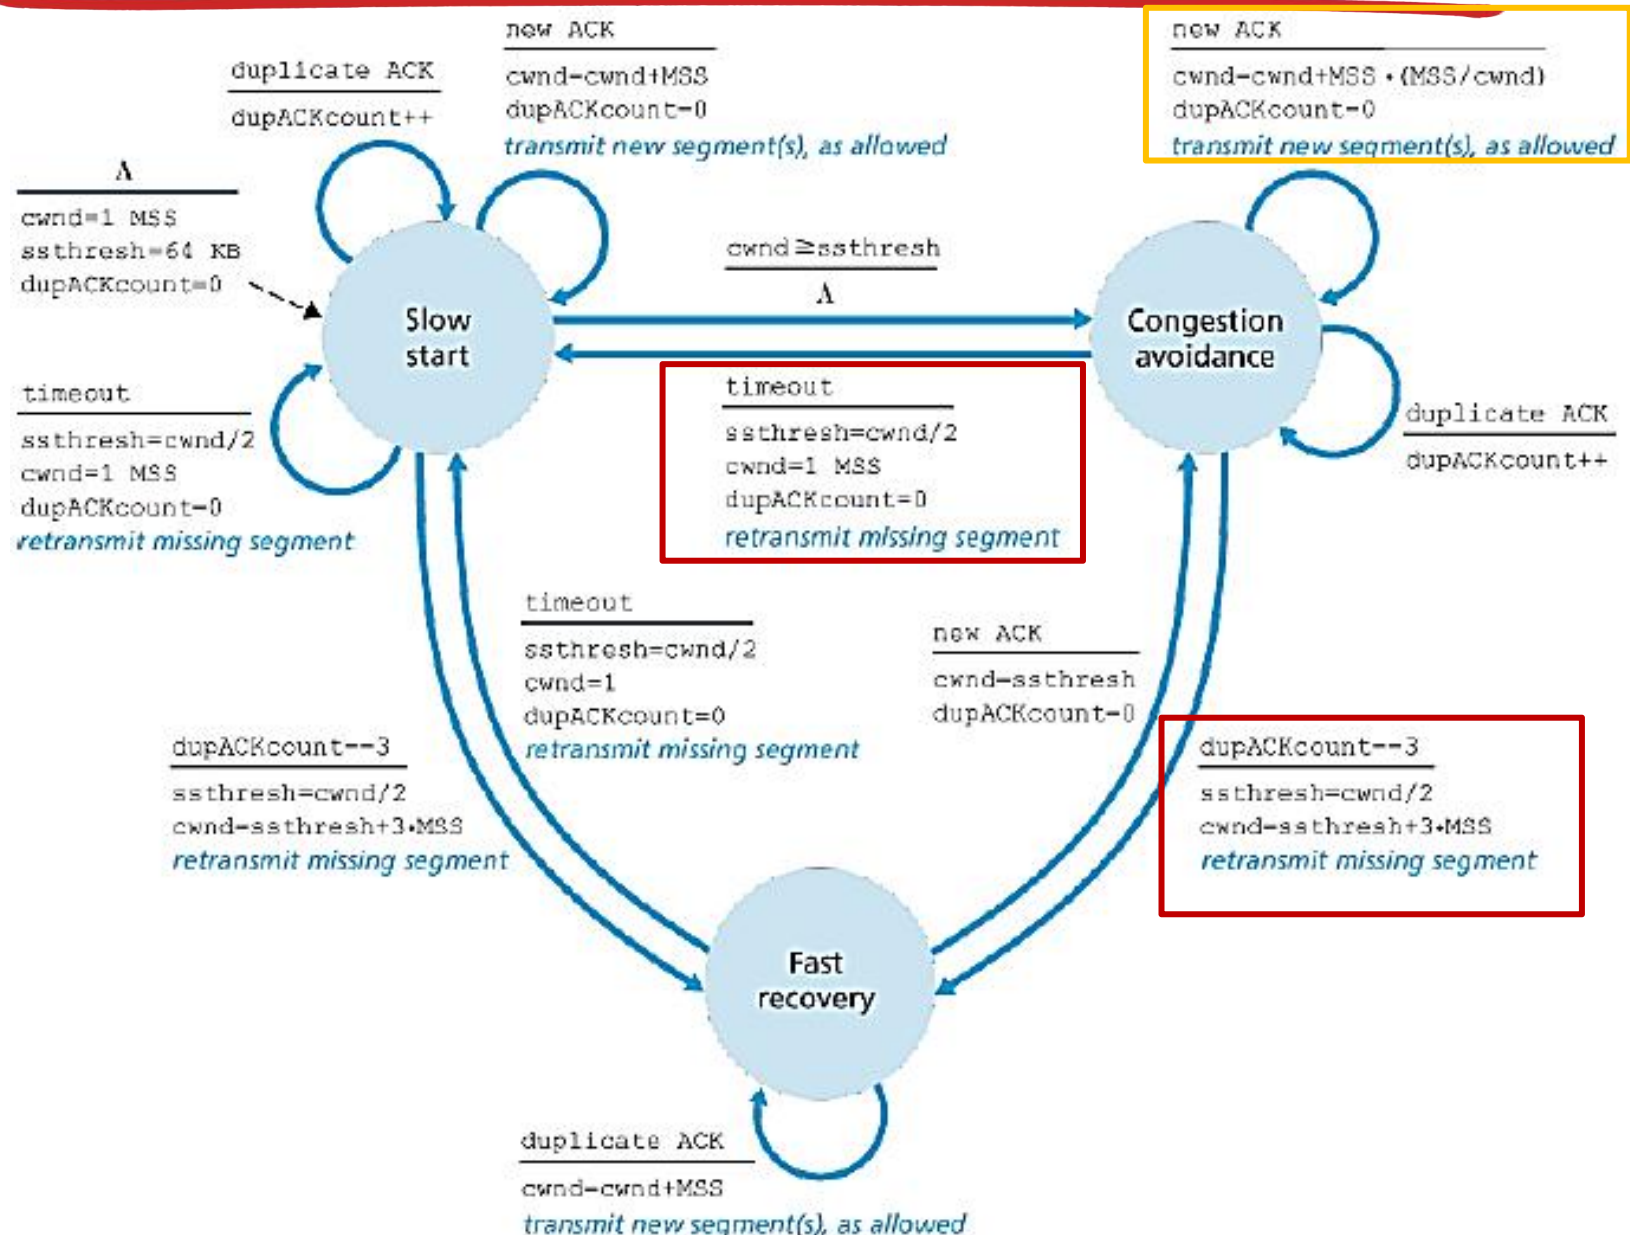
Trigger:  $cwnd \geq ssthresh$

Increases  $cwnd$  **linearly**: by one MSS every RTT

- ❖ Increase  $cwnd$  by  $(MSS/cwnd)MSS$  bytes whenever a new acknowledgment arrives.
- ❖ E.g., if MSS is 1,460 bytes and  $cwnd$  is 14,600 bytes, then 10 segments are being sent within an RTT.
  - Each arriving ACK (assuming one ACK per segment) increases the congestion window size by 1/10 MSS,




# TCP Congestion Control: FSM



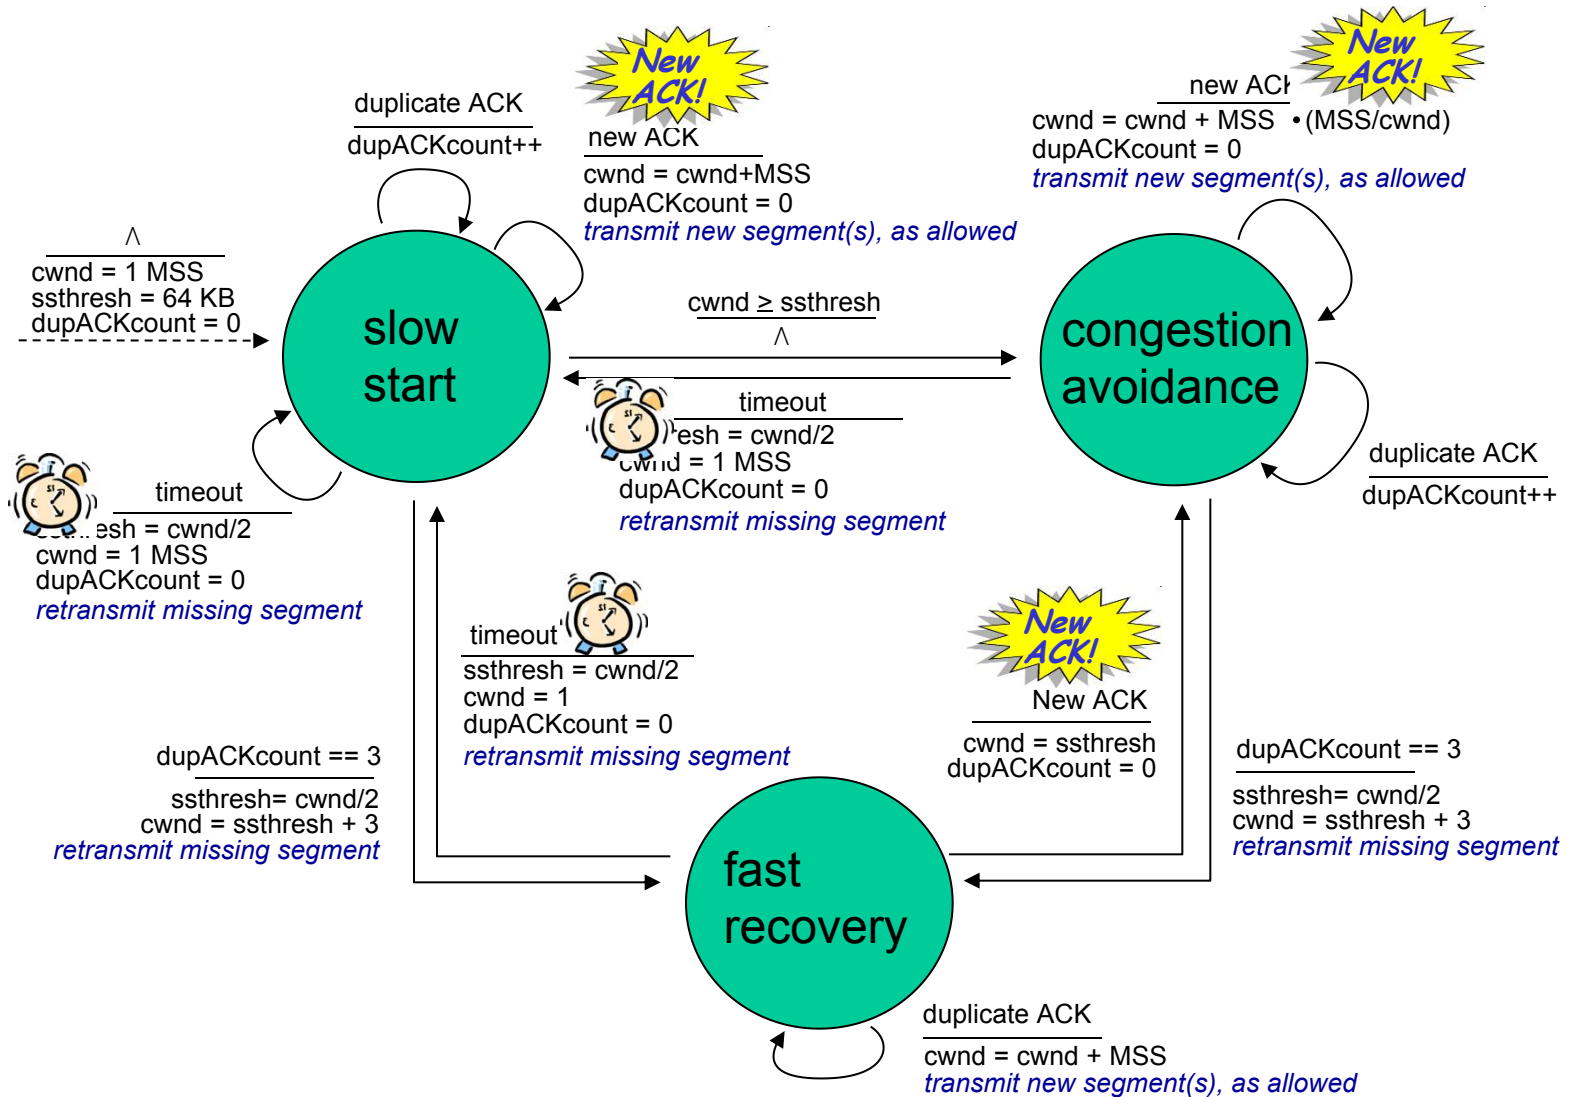
# TCP Fast Recovery

- ❖ Trigger (RENO) : triple duplicate ACKs
- ❖  $ssthresh = cwnd / 2; cwnd = ssthresh + 3MSS$
- ❖ The value of  $cwnd$  is increased by 1 MSS for every duplicate ACK received for the missing segment that caused TCP to enter the fast-recovery state
- ❖ when an ACK arrives for the missing segment, enter congestion avoidance

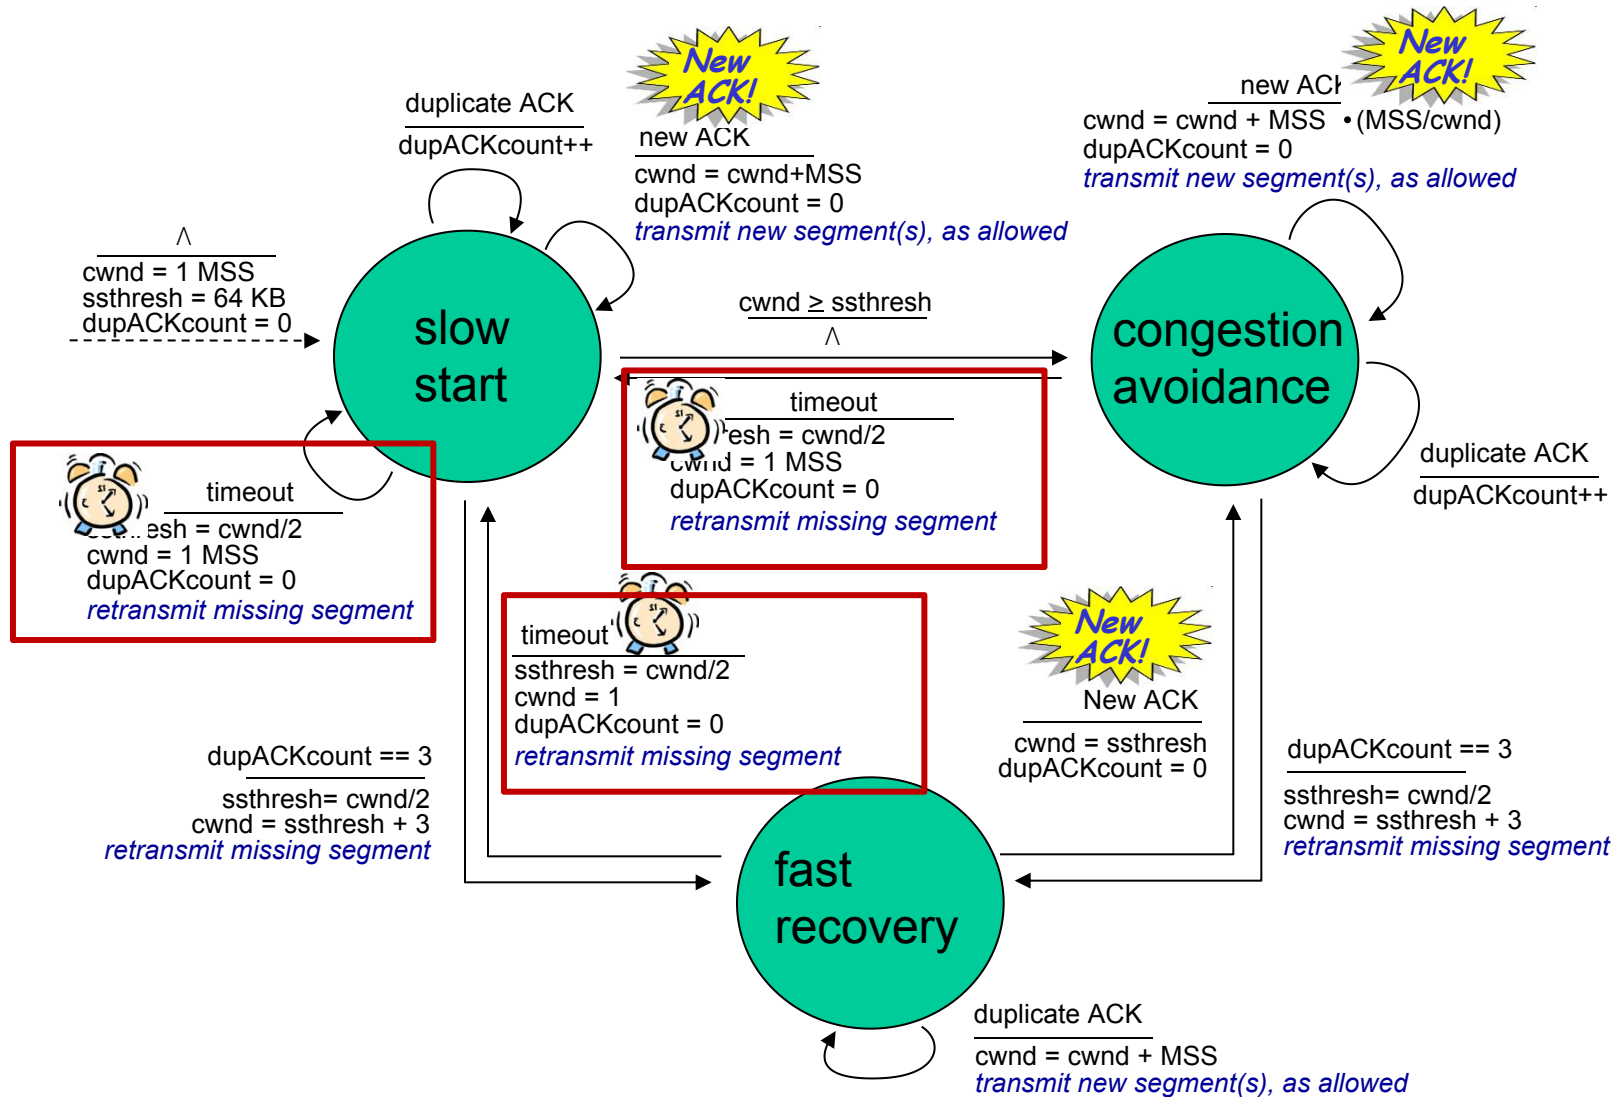




# Summary: TCP Congestion Control

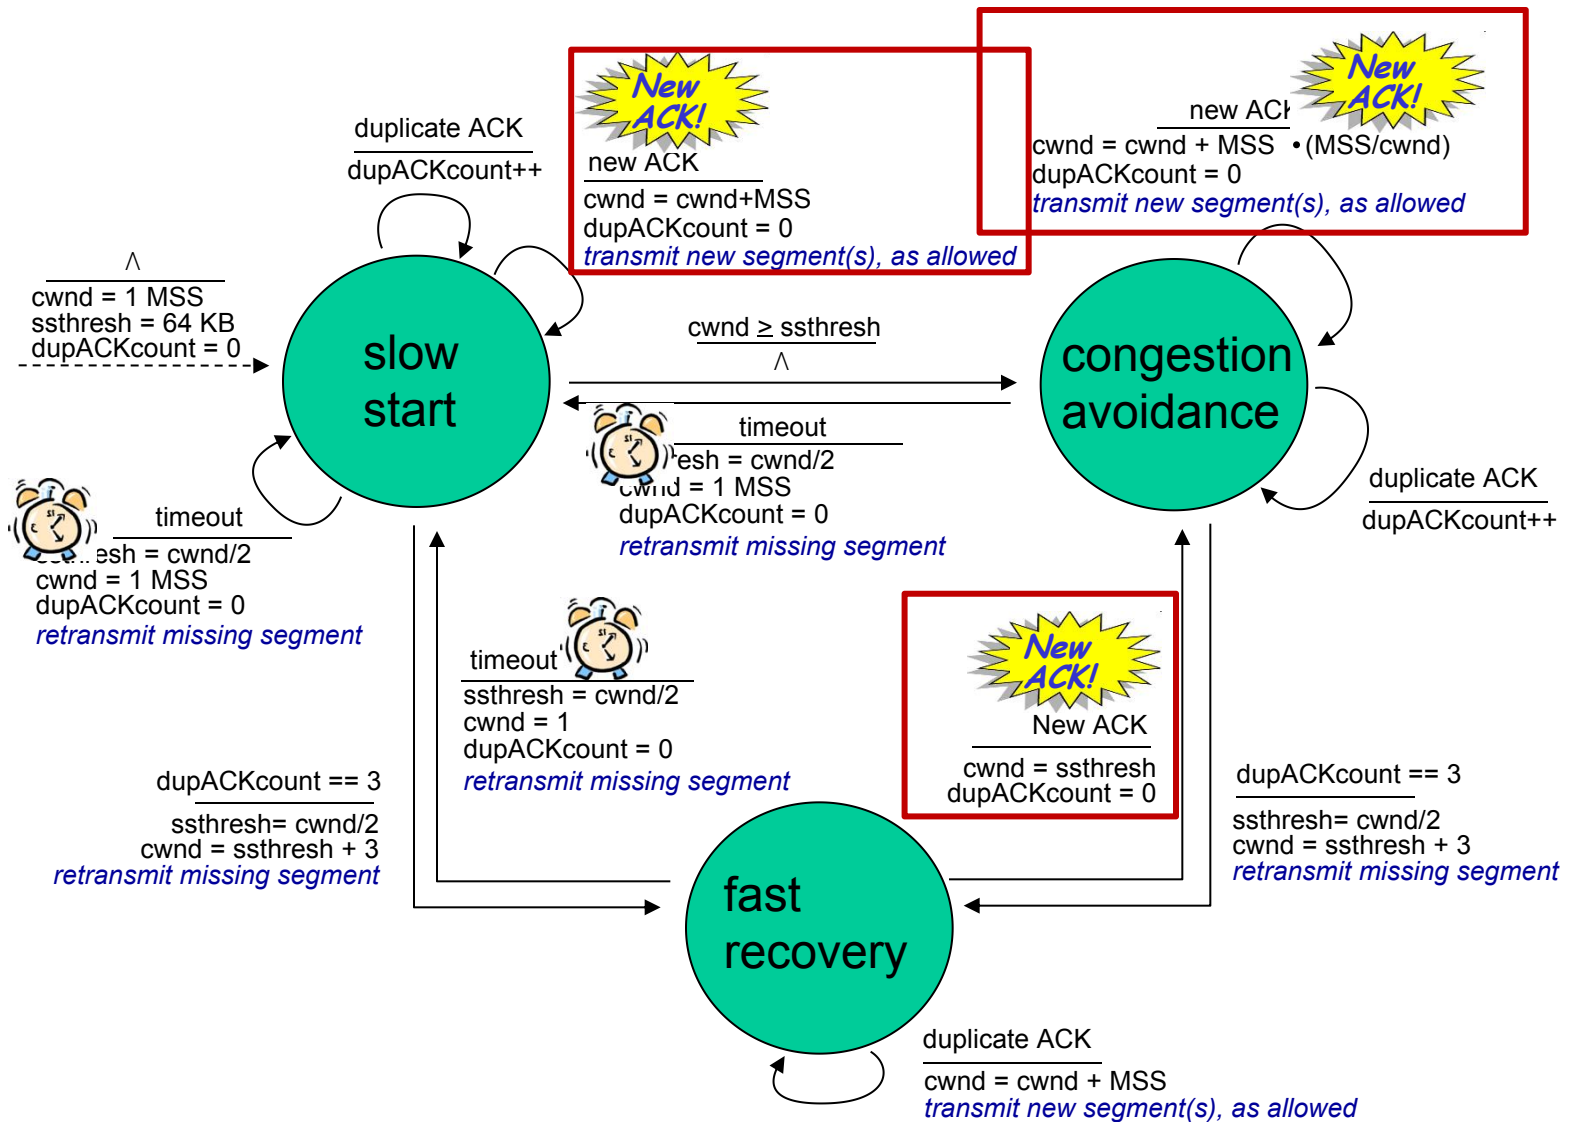


# Summary: TCP Congestion Control

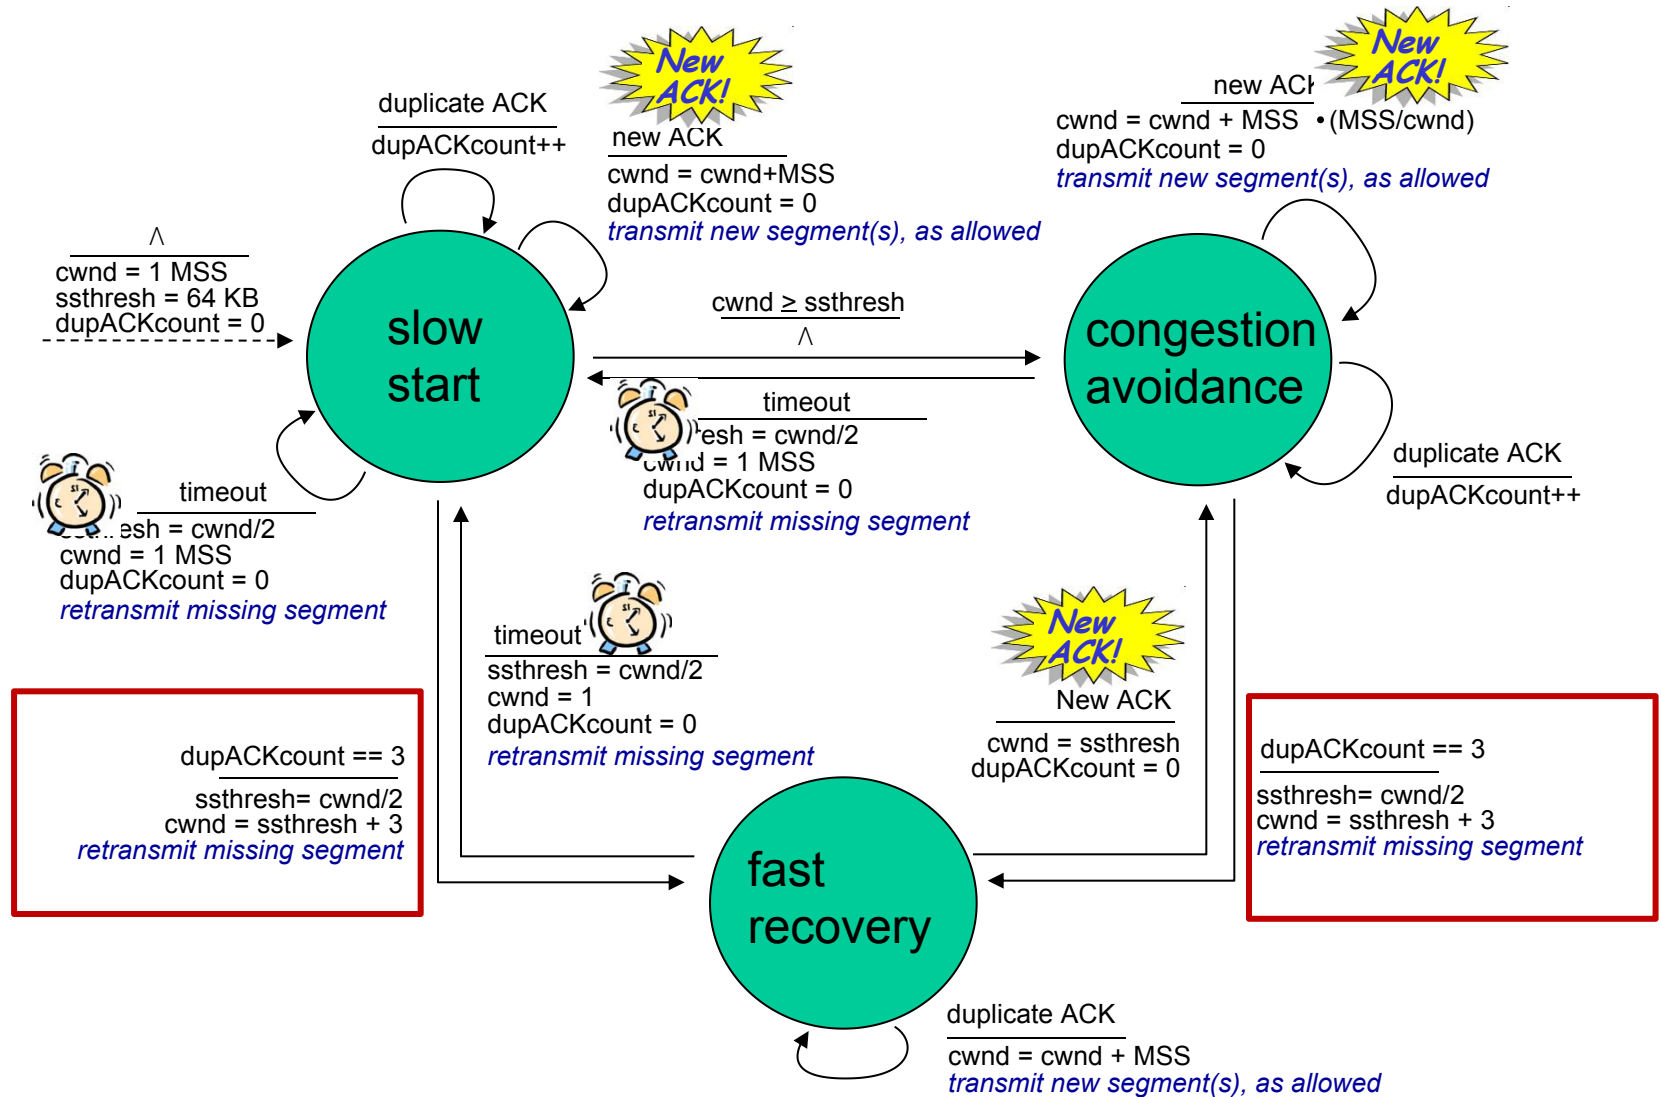




# Summary: TCP Congestion Control



# Summary: TCP Congestion Control




# TCP congestion control: additive increase multiplicative decrease

**Approach:** sender increases transmission rate (window size), **probing** for usable bandwidth, until loss occurs

- *additive increase:* increase **cwnd** by 1 MSS every RTT until loss detected
- *multiplicative decrease:* cut **cwnd** in half after loss


AIMD: probing for bandwidth



# TCP throughput

- ❖ Avg. TCP throughput as function of window size, RTT?
  - ignore slow start, assume always data to send
- ❖  $W$ : window size (measured in bytes) where loss occurs
  - avg. window size (# in-flight bytes) is  $\frac{3}{4} W$
  - avg. throughput is  $\frac{3}{4}W$  per RTT

$$\text{avg TCP thruput} = \frac{3}{4} \frac{W}{\text{RTT}} \text{ bytes/sec}$$



# TCP Futures: TCP over “long, fat pipes”

$$\text{avg TCP thruput} = \frac{3}{4} \frac{W}{\text{RTT}} \text{ bytes/sec}$$

- ❖ example: 1500 byte segments, 100ms RTT, want 10 Gbps throughput
- ❖ requires  $W = 83,333$  in-flight segments
- ❖ throughput in terms of segment loss probability,  $L$  [Mathis 1997]:


$$\text{TCP throughput} = \frac{1.22 \cdot \text{MSS}}{\text{RTT} \sqrt{L}}$$

→ to achieve 10 Gbps throughput, need a loss rate of  $L = 2 \cdot 10^{-10}$  – *a very small loss rate!*

- ❖ new versions of TCP for high-speed

# TCP Fairness


**Fairness goal:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $R/K$



# Why is TCP fair?

two competing sessions:

- ❖ additive increase gives slope of 1, as throughput increases
- ❖ multiplicative decrease decreases throughput proportionally



# Fairness (more)

## *Fairness and UDP*

- ❖ multimedia apps often do not use TCP
  - do not want rate throttled by congestion control
- ❖ use UDP:
  - send audio/video at constant rate, tolerate packet loss
- ❖ UDP sources to crowd out TCP traffic

## *Fairness, parallel TCP connections*


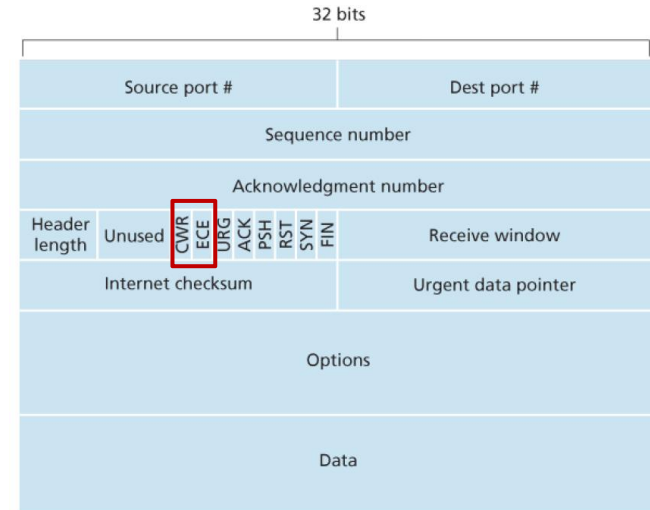
- ❖ application can open multiple parallel connections between two hosts
- ❖ web browsers do this
- ❖ e.g., link of rate  $R$  with 9 existing connections:
  - new app asks for 1 TCP, gets rate  $R/10$
  - new app asks for 11 TCPs, gets more than  $R/2$



# Explicit Congestion Notification (ECN)

*network-assisted congestion control:*

- two bits in IP header (ToS field) marked *by network router* to indicate congestion
- congestion indication carried to receiving host
- receiver sets ECE bit on receiver-to-sender ACK segment to notify sender of congestion



# Chapter 3: summary

- ❖ principles behind transport layer services:
    - multiplexing, demultiplexing
    - reliable data transfer
    - flow control
    - congestion control
  - ❖ instantiation, implementation in the Internet
    - UDP
    - TCP
- next:
- ❖ leaving the network “edge” (application, transport layers)
  - ❖ into the network “core”