# Lecture 8
Demand Paging

Prof. Yinqian Zhang

Fall 2024

Demand Paging Mechanisms

How To Go Beyond Physical Memory?

- How to support large address space?
    - 64-bit machine supports up to 4EB address space
    - Applications may use more space than available in physical memory

- Solution: stash away portions of address spaces that aren't currently in use
    - in the next-level of storage (e.g., hard disk drive)
   - slower but much larger
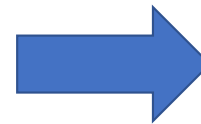
An Abstraction of Address Space

- Who is responsible for moving data?

- Application: **memory overlays**
  - Application in charge of moving data between memory and disk
  - e.g., calling a function needs to make sure the code is in memory!

- OS: **demand paging**
  - OS configures page table entries
  - Virtual page maps to physical memory or files in disk
  - Process sees an abstraction of address space
  - OS determines where the data is stored

Swap Space

- Swap space is a partition or a file stored on the disk
  - OS swaps pages out of memory to it
  - OS swaps pages from it into memory

- Swap space conceptually divided into page-sized units
  - OS maintains a disk address of each page-sized unit

disk

swap space

each page-sized unit in
swap space can be
addressed by the OS

Swap Space Example

- 4-page physical memory and an 8-page swap space
  - Proc 0 has three virtual pages
  - Proc 1 has four virtual pages
  - Proc 2 and Proc 3 each has two virtual pages

	PFN 0	PFN 1	PFN 2	PFN 3
Physical Memory	Proc 0 [VPN 0]	Proc 1 [VPN 2]	Proc 1 [VPN 3]	Proc 2 [VPN 0]

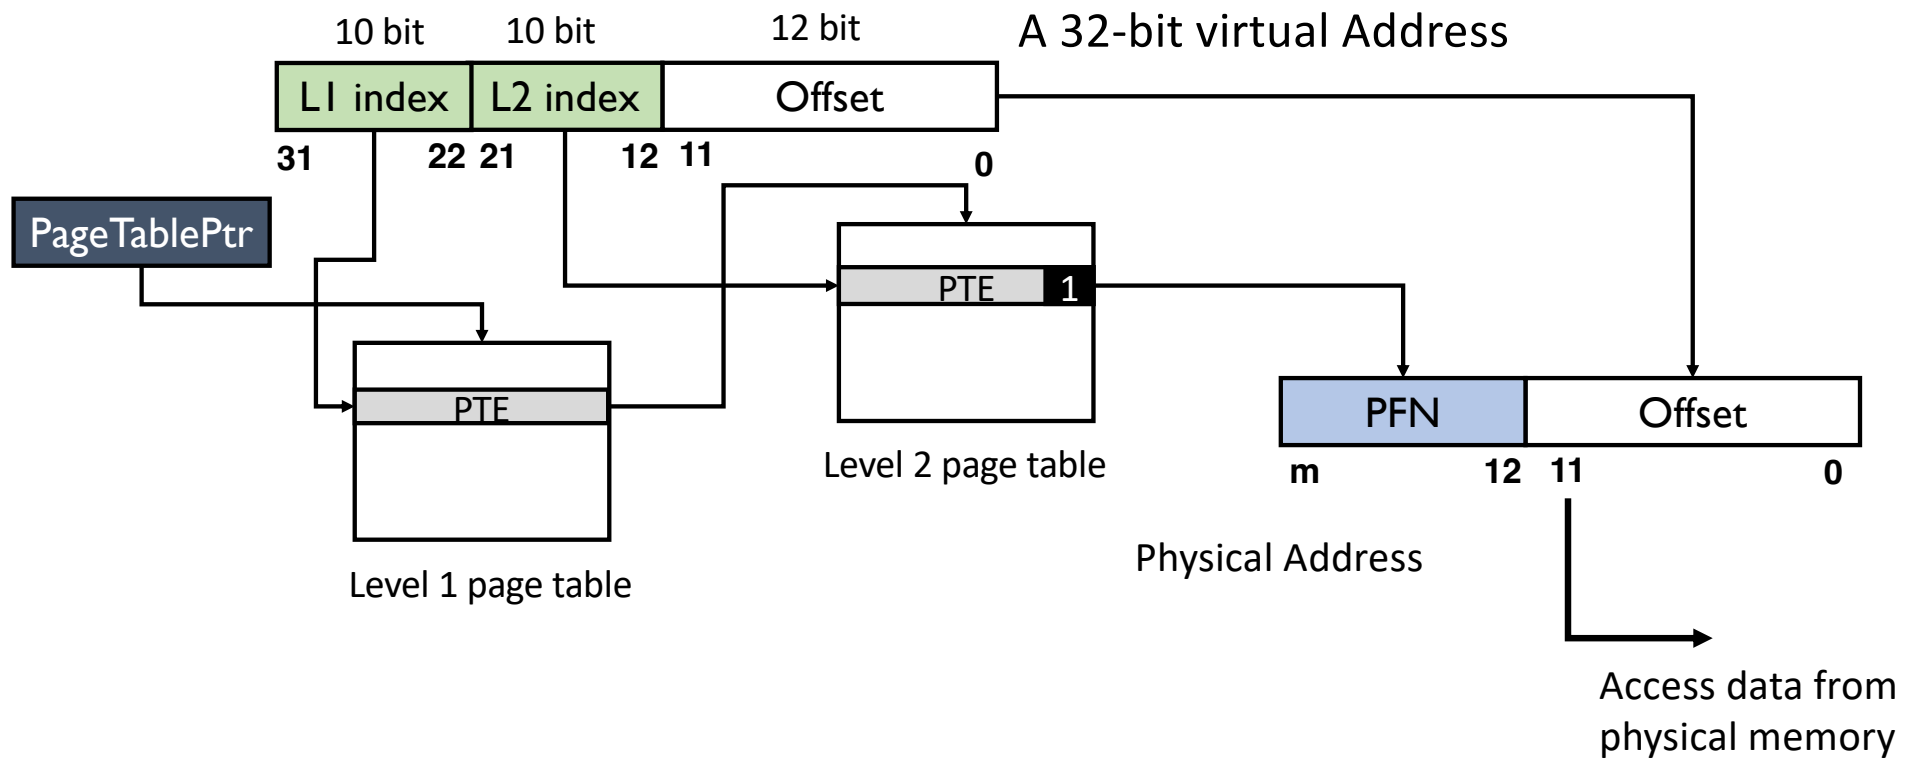|	Block 0	Block 1	Block 2	Block 3	Block 4	Block 5	Block 6	Block 7
Swap Space	Proc 0 [VPN 1]	Proc 0 [VPN 2]	[Free]	Proc 1 [VPN 0]	Proc 1 [VPN 1]	Proc 3 [VPN 0]	Proc 2 [VPN 1]	Proc 3 [VPN 1]

Demand Paging

- Load pages from disk to memory only as they are needed
  - Pages are loaded "on demand"
- Data transferred in the unit of pages
- Two possible on-disk locations
  - Swap space:
    - created by OS for temporary storage of pages on disk
    - e.g., pages for stack and heap
  - Program binary files:
    - The code pages from this binary are only loaded into memory when they are executed
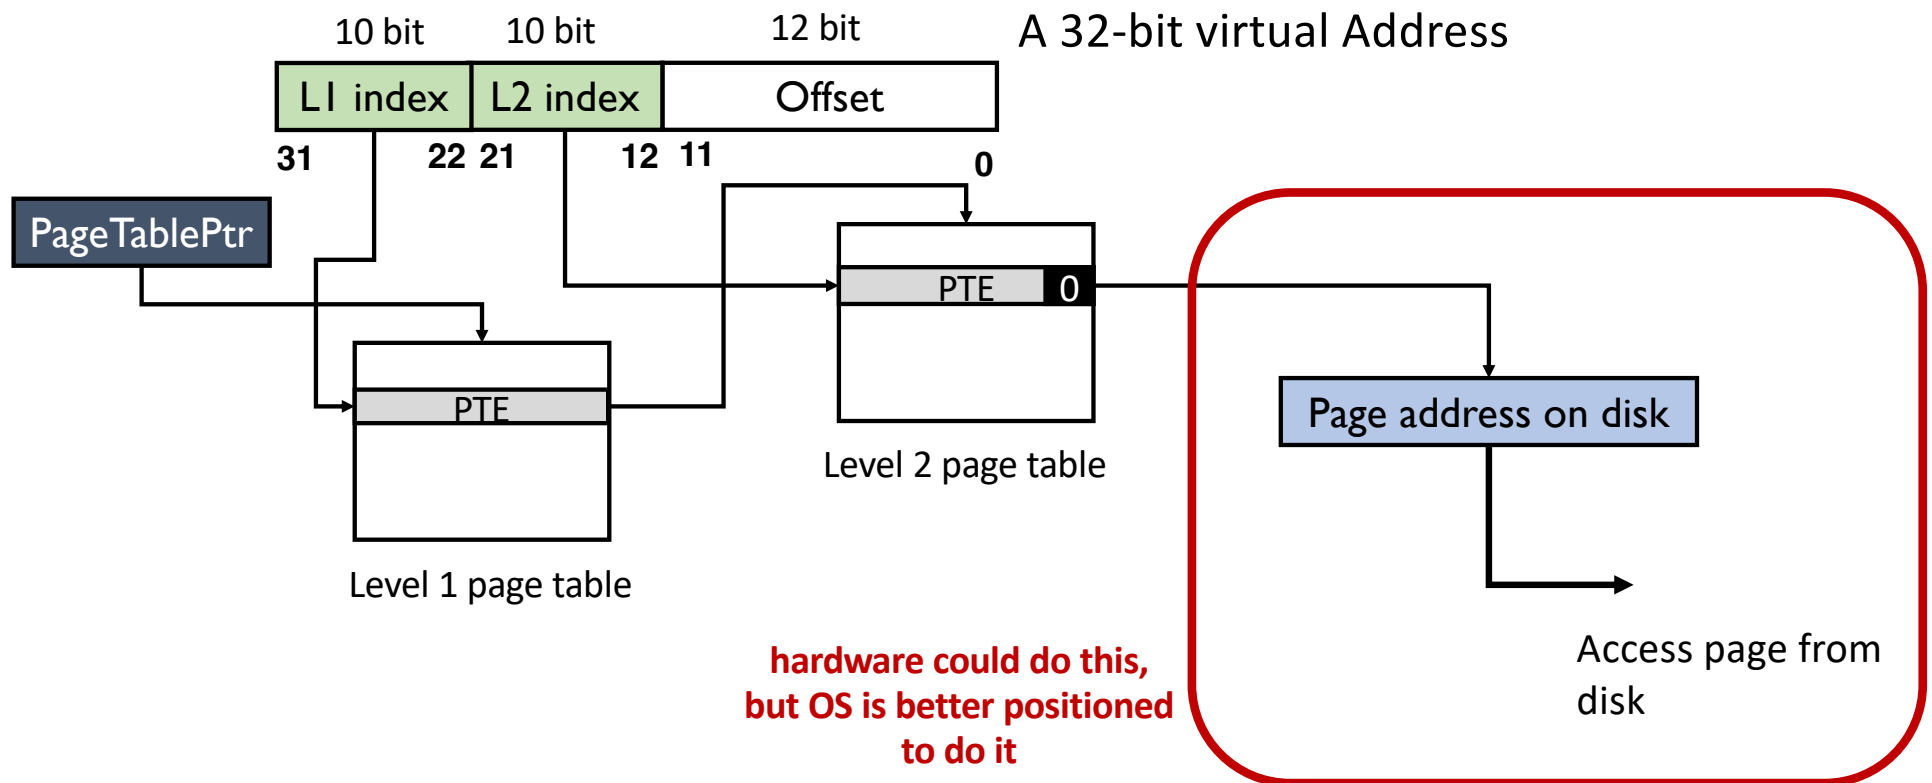   - Read-only, thus never write back

Physical Memory as a Cache

- Physical memory can be regarded as a cache of on-disk swap space
- Block size of the cache?
  - 1 page (4KB)
- Cache organization (direct-mapped, set-associative, fully-associative)?
  - Fully associative: any disk page maps to any page frame
- What is page replacement policy?
  - LRU, Random, FIFO
- What happens on a miss?
  - Go to lower level to fill page (i.e. disk)
- What happens on a write, write-through or write back?
  - write-back: changes are written back to disk when page is evicted

Present Bit

10 bit   10 bit   12 bit   A 32-bit virtual Address

L1 index	L2 index	Offset

31   22 21   12 11   0

PageTablePtr

PTE

Level 1 page table

PTE **1**

Level 2 page table

PFN	Offset

m   12 11   0

Physical Address

Access data from physical memory

9

Present Bit

A 32-bit virtual Address

10 bit	10 bit	12 bit
L1 index	L2 index	Offset

31   22 21   12 11   0

PageTablePtr

PTE

Level 1 page table

PTE | 0

Level 2 page table

Page address on disk

Access page from disk

hardware could do this, but OS is better positioned to do it

Page Faults

- Present bit = 0 raises a page fault exception
  - OS gets involved in address translation

- Page fault handler
  - (1) Find free page frame in physical memory
  - (2) Fetch page from disk and store it in physical memory

- After page fault
  - Return from page fault exception
  - CPU re-execute the instruction that accesses the virtual memory
  - No more page fault since present bit is set this time
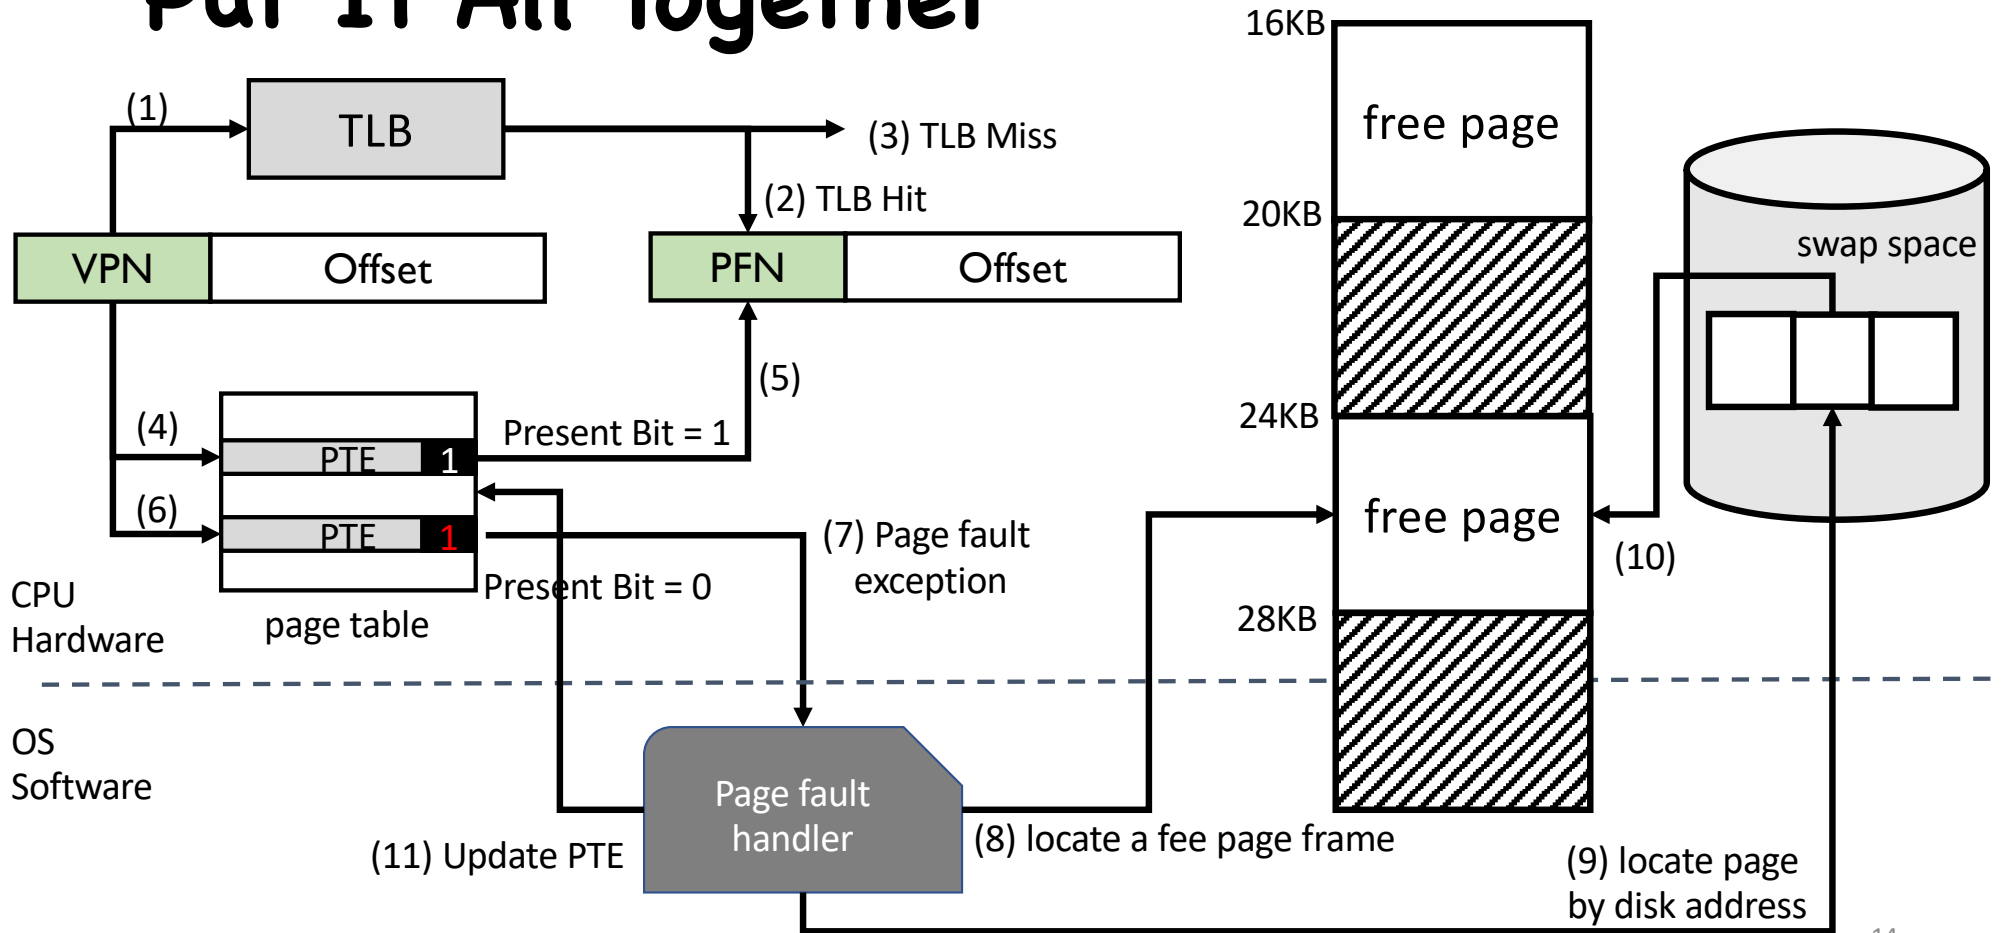 - TLB entry loaded from PTE

Page Faults (Cont'd)

- (1) Find free page frame in physical memory
  - Find one free page frame from a free-page list
  - If no free page, trigger **page replacement**

- **Page Replacement**
  - find a page frame to be replaced
    - Page replacement policy decides which one to replace
  - If page frame to be replaced is dirty, write it back to disk
  - Update all PTEs pointing to the page frame
  - Invalidate all TLB entries for these PTEs

Page Faults (Cont'd)

- (2) Fetch page from disk
  - Determine the faulting virtual address from register
  - Locate the disk address of the page in PTE (where PFN should be stored)
    - It is a very natural choice to make use of the space in PTE
  - Issues a request to disk to fetch the page into memory
  - Wait ...... (could be a very long time, context switch!)
  - When I/O completes, update page table entry: PFN, present bit

Put It All Together

TLB

(1)

(3) TLB Miss

(2) TLB Hit

VPN | Offset

PFN | Offset

(5)

(4)

PTE 1

Present Bit = 1

(6)

PTE 1

Present Bit = 0

page table

(7) Page fault
exception

CPU
Hardware

OS
Software

Page fault
handler

(11) Update PTE

(8) locate a fee page frame

16KB

free page

20KB

24KB

free page

(10)

28KB

swap space

(9) locate page
by disk address

14

When to Trigger Page Replacement

- Proactive page replacement usually leads to better performance
  - Page replacement even though no one needs free page frames (yet)
  - Always reserve some free page frames in the system

- Swap daemon
  - background process for reclaiming page frames
  - Low watermark: a threshold to trigger swap deamon
  - High watermark: a threshold to stop reclaiming page frames

Page Replacement Policy

Effective Access Time

- EAT = Hit Rate x Hit Time + Miss Rate x Miss Penalty
- Example:
  - Memory access time = 200 nanoseconds
  - Average page-fault service time = 8 milliseconds
  - Suppose p = Probability of miss, 1–p = Probably of hit
  - Then, we can compute EAT as follows:

$$EAT = (1-p) \times 200ns + p \times 8\ ms$$
$$= (1-p) \times 200ns + p \times 8{,}000{,}000ns$$

Effective Access Time (Cont'd)

$$\begin{aligned} \text{EAT} &= (1-p) \times \ 200\text{ns} + p \times 8 \text{ ms} \\ &= (1-p) \times 200\text{ns} + p \times 8{,}000{,}000\text{ns} \end{aligned}$$

- If one access out of 1,000 causes a page fault, then EAT is about 8.2 μs:
  - This is a slowdown by a factor of 40!

- What if we want slowdown by less than 10%?
  - 200ns x 1.1 < EAT $\Rightarrow$ p < 2.5 x 10^{-6}
  - This is about 1 page fault in 400,000!

Types of Cache Misses: Three Cs

- Compulsory Misses:
  - Cold-start miss: pages that have never been fetched into memory before
  - Prefetching: loading them into memory before needed

- Capacity Misses:
  - Not enough memory: must somehow increase available memory size
  - One option: Increase amount of DRAM (not quick fix!)
  - Another option: If multiple processes in memory: adjust percentage of memory allocated to each one!

- Conflict Misses:
  - fully-associative cache (OS page cache) does not have conflict misses

Page Replacement Policies

- Optimal (also called MIN):
  - Replace page that will not be used for the longest time
  - Lead to minimum page faults in theory

- FIFO (First In, First Out)
  - Throw out oldest page first
  - May throw out heavily used pages instead of infrequently used

- RANDOM:
  - Pick random page for every replacement
  - Pretty unpredictable – makes it hard to make real-time guarantees

Replacement Policies (Con't)

- Least Recently Used (LRU):
  - Replace page that has not been used for the longest time
  - **Temporal locality of program**
    - If a page has not been used for a while, it is unlikely to be used in the near future
- Least Frequently Used (LFU)
  - Replace page that has not been accessed many times
  - **Spatial locality of program**
    - if a page has been accessed many times, perhaps it should not be replaced as it clearly has some value.

Example: Optimal (MIN)

- Suppose we have 3 page frames, 4 virtual pages, and following reference string:
  - A B C A B D A D B C B

	Ref:	A	B	C	A	B	D	A	D	B	C	B
	Page:											
1		A									C	
2			B									
3				C			D					

- MIN: 5 faults
  - Where will D be brought in? Look for page not referenced farthest in future

Example: FIFO

- Suppose we have 3 page frames, 4 virtual pages, and following reference string:
  - A B C A B D A D B C B

Ref:	A	B	C	A	B	D	A	D	B	C	B
Page:											
1	A					D				C	
2		B					A				
3			C						B		

- FIFO: 7 faults
- When referencing D, replacing A is bad choice, since need A again right away

Example: LRU

- Suppose we have 3 page frames, 4 virtual pages, and following reference string:
  - A B C A B D A D B C B

Ref: Page:	A	B	C	A	B	D	A	D	B	C	B
1	A									C	
2		B									
3			C			D					

- LRU performs the same as Optimal

Is LRU Always Close to Optimal?

- Consider the following reference string: A B C D A B C D A B C D
- LRU performs as follows (the same as FIFO):

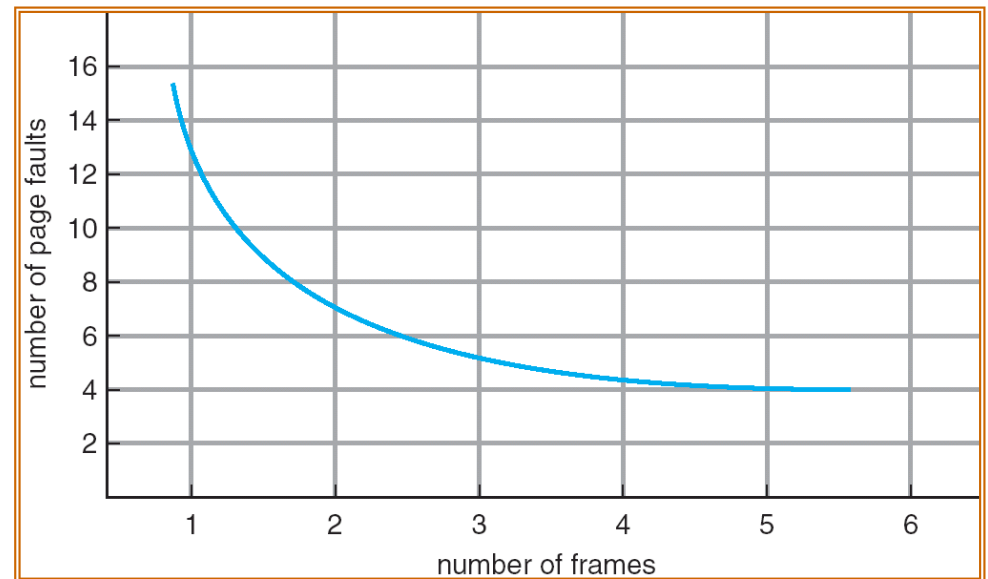| Ref: Page:	A	B	C	D	A	B	C	D	A	B	C	D
1	A			D			C			B		
2		B			A			D			C	
3			C			B			A			D

Is LRU Always Close to Optimal? (Cont'd)

- Consider the following: A B C D A B C D A B C D
- MIN performs better:

Ref: Page:	A	B	C	D	A	B	C	D	A	B	C	D
1	A									B		
2		B					C					
3			C	D								

Bélády's Anomaly

- One desirable property: When you add memory the miss rate drops
  - Yes for LRU and MIN
  - Not necessarily for FIFO!

- Bélády's anomaly
  - For FIFO, more page frames may lead to more page faults!

Bélády's Anomaly Example

- Page replacement with 3 page frames

| Ref: Page:	A	B	C	D	A	B	E	A	B	C	D	E
1	A			D			E					
2		B			A					C		
3			C			B					D	

- Page replacement with 4 page frames

Ref: Page:	A	B	C	D	A	B	E	A	B	C	D	E
1	A						E				D	
2		B						A				E
3			C						B			
4				D						C		

Page Fault Curve

- Page fault curve for FIFO on reference string
  - 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

How do you plot
a chart like this?
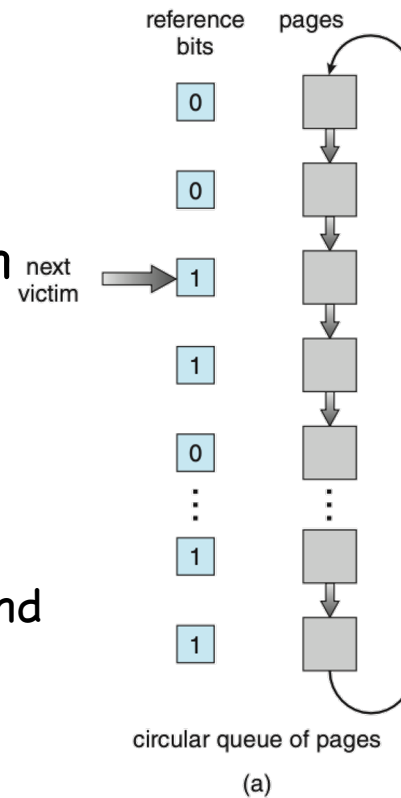
LRU Implementation

- Hardware support is necessary
  - Update a data structure in OS upon every memory access
  - E.g., a timestamp counter for each page frame

- Overhead
  - One additional memory write for each memory access
    - TLB hit does not save the extra memory access

  - Scan the entire memory to find the LRU one
    - 4GB physical memory has 1 million page frames
    - sorting is time consuming

LRU Approximation with Reference Bit

- Reference bit
  - One reference bit per page frame
  - All bits are cleared to 0 initially
  - The first time a page is referenced, the reference bit is set **by CPU**
    - Can be integrated with page table walk
  - The order of page accesses approximated by two clusters: **used** and **unused** pages

- Examples:
  - Clock algorithm (also called second-chance algorithm)
  - Enhanced clock algorithm with dirty bits

Clock Algorithm

- Arrange physical pages in a circular list
- CPU sets reference bit to 1 upon first access
- OS maintains a pointer
  - When a replacement occur, check reference bit of the current page
  - If 1: the page has been accessed recently, clear the bit (set to 0) and move to the next page
  - If 0: the page has not been accessed recently, good candidate for replacement, stop

reference bits   pages

0
0
next victim → 1
1
0
⋮   ⋮
1
1

circular queue of pages

(a)

reference bits   pages

0
0
0
0
→ 0
⋮   ⋮
1
1

circular queue of pages

(b)

Clock Algorithm with Dirty Bit

- Enhance clock algorithm with a dirty bit
  - dirty bit = 1: the page has recently been modified
- CPU sets dirty bit to 1 upon **write** access
- When a replacement occurs, OS checks (ref bit, dirty bit), and selects a candidate page in decreasing order
  - (0, 0) neither recently used nor modified — best page to replace
  - (0, 1) not recently used but modified — not quite as good, because the page will need to be written out before replacement
  - (1, 0) recently used but clean — probably will be used again soon
  - (1, 1) recently used and modified — probably will be used again soon, and the page will be need to be written out to secondary storage before it can be replaced

LRU Approximation with Reference Bit and Counter

- Each physical page frame is associated with one reference bit and a counter
  - Reference bit indicate recent access
    - set by CPU hardware, cleared by OS
  - Counter records history of accesses
    - Maintained by OS

- Examples
  - Additional-reference-bits algorithm
  - N[th]-chance clock algorithm

Additional-reference-bits Algorithm

- 8-bit history register associated with each page frame
- Timer interrupt every 100ms
  - reference bit shifts to highest bit in the history counter
  - other bits shift right and discard the lowest bit
  - 00000000 unused page in 800ms
- Compare history counter as unsigned integer
  - Larger value more recently used
  - 11000100 > 01110111
- Approximate LRU with more bits and more frequent interrupts

reference bit	history counter
1	00100101
0	10010010
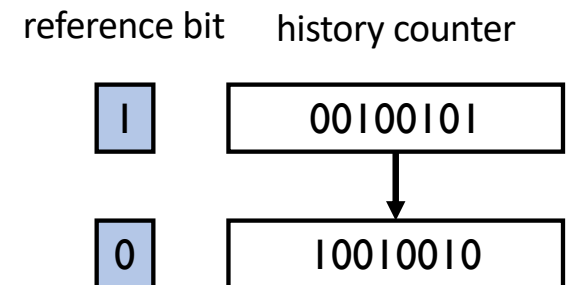
Nth-chance Clock Algorithm

- All page frames arranged in a circular list and each page frame is associated with a reference bit and a counter
- CPU hardware sets reference bit upon memory accesses
- OS checks the reference bit of the page pointed to by the clock hand
  - 1 → clear reference bit and the counter
  - 0 → increment counter; if count=N, replace page

- How do we pick N?
  - Large N? Better approximation to LRU
    - If N ~ 1K, really good approximation
  - Small N? More efficient
    - Otherwise might have to look a long way to find free page

Page Frame Allocation

Allocation of Page Frames

- How do we allocate memory among different processes?
  - Does every process get the same fraction of memory? Different fractions?
  - Should we completely swap some processes out of memory?
- Minimum number of pages per process
  - Depends on the computer architecture
  - How many pages would one instruction use at one time
    - x86 only allows data movement between memory and register and no indirect reference
    - needs at least one instruction page, one data page, and some page table pages
- Maximum number of pages per process
  - Depends on available physical memory

Global versus Local Allocation

- Global replacement
  - Process selects replacement frame from all page frames
  - One process can take a frame from another process

- Local replacement
  - Each process selects from only its own set of allocated frames

Allocation Algorithms

- Equal allocation:
  - Every process gets same amount of memory
  - Example: 100 frames, 5 processes $\rightarrow$ process gets 20 frames
- Proportional allocation
  - Number of page frames proportional to the size of process
    s_i = size of process p_i *and m* = total number of frame
    a_i = allocation for p_i = $m \times \dfrac{s_i}{\sum s_j}$
- Priority Allocation:
  - Number of page frames proportional to the priority of process
  - Possible behavior: If process p_i generates a page fault, select for replacement a frame from a process with lower priority number

Thrashing

- The memory demands of the set of running processes simply exceeds the available physical memory
- Early OS
  - Working set: the pages used actively of a process
  - Reduce the # of process so their working set fit into memory
- Modern OS
  - Out-of-memory killer when memory is oversubscribed
  - May need a reboot

Thank you!